

From Ledgers, to Options, to Latency

A History of Trading, Mathematical Finance, and the Quant Profession

Paul Alexander Bilokon

Mathematical Finance Section, Department of Mathematics, Imperial College London

paul.bilokon@imperial.ac.uk

August 14, 2026

Abstract

The quantitative finance profession is often narrated as a late-twentieth-century consequence of option-pricing theory, electronic markets, and the migration of mathematicians and physicists to Wall Street. That account is correct but incomplete. The modern quant is the latest expression of a much longer co-evolution among five systems: institutions of exchange, mathematical representation, information technology, organisational specialisation, and regulation. This paper develops that thesis from the prehistory of money and markets to machine learning and large language models. It treats merchant accounting, bills of exchange, bourses, joint-stock companies, organised derivatives markets, telegraphy, the stock ticker, and the formation of specialised trading desks as genuine antecedents of quantitative work rather than as decorative prehistory. It then follows the mathematical line from Bachelier through axiomatic probability, stochastic calculus, portfolio theory, asset pricing, no-arbitrage valuation, volatility modelling, market microstructure, optimal execution, and modern computational statistics. Parallel to this intellectual history, it traces the professional line from mathematically sophisticated merchants and actuaries to portfolio “quantifiers,” derivatives quants, systematic hedge funds, electronic market makers, quantitative developers, risk quants, data scientists, and AI-assisted researchers.

The paper also argues that the profession cannot be understood by following theory alone. Each wave of quantification depended on a matching substrate of market data, computing, communication, and institutional demand. Conversely, mathematical models reshaped the markets in which they were deployed, while crises and regulation repeatedly redirected quantitative labour toward new problems such as value-at-risk, counterparty exposure, xVA, capital, liquidity, execution quality, and model risk. The result is a history in which the quant is not a single occupation but a family of roles occupying the interface between formal models and executable decisions.

Keywords: quantitative finance; quant; trading; financial history; mathematical finance; derivatives; volatility; market microstructure; algorithmic trading; high-frequency trading; risk management; machine learning; artificial intelligence.

1 Introduction: what is a quant, and when does the history begin?

The word *quant* is recent; the activities from which the occupation was assembled are not. A useful history therefore has to distinguish at least three questions. First, when did people begin to use explicit calculation to make decisions under financial uncertainty? Second, when did institutions create specialist occupations whose principal product was a model, estimate, hedge, forecast, or

trading rule? Third, when did those occupations acquire the social identity now denoted by “quantitative analyst,” “quantitative researcher,” “strategist,” “quant developer,” or simply “quant”? These questions have different answers.

The modern word is best treated as the name of an occupational family rather than a single job description. One branch prices and hedges derivatives; another constructs portfolios and systematic signals; another studies execution and market microstructure; another builds risk and capital models; another develops the software and computing substrate on which all of the above run. The boundaries are porous. A portfolio manager may be a researcher, a researcher may write production code, and a quant developer may contribute materially to the mathematics of a model. What unifies the family is a workflow in which formal representation—mathematical, statistical, computational, or algorithmic—is expected to affect a financial decision.

The organising claim is that the quant profession emerged from the co-evolution of five forces:

- F1. Markets and institutions:** standardised claims, exchanges, banks, clearing, trading desks, asset managers, and hedge funds created repeated decisions for which quantitative methods could be economically valuable.
- F2. Mathematical representation:** probability, statistics, optimisation, stochastic processes, numerical analysis, control, and later machine learning supplied increasingly expressive languages for those decisions.
- F3. Information technology:** accounting systems, print, postal networks, telegraphy, the ticker, telephone, electronic feeds, databases, networks, and specialised hardware progressively compressed the delay between observation, inference, and action.
- F4. Organisational specialisation:** the merchant, broker, jobber, statistician, actuary, economist, operations researcher, programmer, trader, and risk manager were gradually recombined into specialist quantitative roles.
- F5. Law and regulation:** the permissions and constraints of markets—from exchange rules to capital requirements and execution obligations—continually changed the objective function the quant was asked to solve.

This framing avoids two opposite mistakes. The first is presentism: reading a seventeenth-century merchant as though he were a derivatives strategist with a laptop. The second is discontinuity: pretending that quantitative finance appeared ex nihilo with the publication of the Black–Scholes–Merton papers in 1973. The meaningful continuity is functional. A Medici clerk maintaining multi-city accounts, a nineteenth-century arbitrageur paying for faster information, an options dealer reasoning about premiums and fluctuations, a 1970s derivatives theorist constructing a replicating portfolio, and a 2020s electronic market maker optimizing queues and latency occupy very different technological worlds, yet all transform information into financial action through a formal or semi-formal procedure.

Claims of “the first” are consequently used sparingly. Historical records rarely line up with modern occupational categories. A dedicated room in a merchant bank can be recognisably desk-like without satisfying every feature of a modern trading desk; an early systematic investor can be a proto-quant without having used that word. Where an absolute priority is contestable, the paper prefers formulations such as “an early antecedent,” “among the earliest,” or “a recognisably modern form.”

The selection and interpretation offered here are necessarily incomplete, and the author is acutely conscious of the limits of any attempt to reconstruct so broad a history, particularly where practitioner contributions, proprietary technologies, institutional developments, and less well documented intellectual lineages are concerned. Important people, institutions, methods, and episodes will inevitably have been omitted, and some attributions, dates, emphases, or interpretations may benefit from correction or qualification. The author would therefore be especially grateful to readers for comments, suggestions, corrections, additional sources, and proposed additions. Such contributions will be welcomed in the spirit in which this history is intended: not as a definitive catalogue or final allocation of credit, but as an evolving attempt to record as faithfully and fairly as possible the many strands from which the quantitative-finance profession emerged.

A primary motivation for compiling this history is pedagogical. It is written in particular for my past, present, and future students at Imperial College London, including those on the MSc Mathematics and Finance, as well as for students, young researchers, and aspiring quantitative practitioners elsewhere in the world. Quantitative finance can appear from the outside as an intimidating collection of models, technologies, markets, institutions, and specialised professional cultures; one purpose of this document is to show how these pieces came to fit together, and to give newcomers a map of the intellectual and professional territory they are entering. If it helps even a few of them understand where their discipline came from, discover an overlooked idea or predecessor, or find their own place within its continuing history, it will have served an important part of its purpose. To all of them, embarking on what can be a demanding, strange, and immensely rewarding profession: good luck.

A remarkable number of people built this profession, often in places where the historical record barely noticed them. I have tried to reconstruct some of it; please help me get it right.

2 The long prehistory of exchange and calculation

2.1 Subsistence, reciprocity, barter, and tribute

Exchange predates markets, and markets predate finance. The earliest economic arrangements were embedded in kinship, subsistence, reciprocity, and political authority rather than organised around anonymous spot transactions. Anthropological and historical accounts warn against a simple evolutionary ladder in which barter mechanically gives rise to money and then to markets [159, 319]. Barter certainly occurred, particularly between groups and in periods of monetary disruption, but systems of obligation and accounting could precede widespread coin use. This distinction matters for a history of quantification because an obligation recorded in a ledger is already a numerical state variable even when no coin changes hands.

Tributary and palace economies add another ingredient: centralised measurement. Grain, labour, land, and tax obligations had to be counted, assessed, and reconciled. The practical mathematics of administration was not financial mathematics in the modern sense, but it established a durable relationship between economic power and record-keeping. A quantitative profession becomes possible only after transactions leave sufficiently standard traces to be compared across time, counterparties, and locations.

2.2 Money of account and money of exchange

It is useful to distinguish money as a unit of account from money as a physical medium of exchange. Tallying technologies, commodity accounts, clay tokens, written ledgers, coinage, paper claims, and bank deposits belong to a long sequence in which the informational representation of value is at least as important as its physical carrier [98, 155, 159]. Once obligations can be represented numerically, they can be netted, discounted, transferred, compared, and eventually modelled.

This observation foreshadows a recurring theme. Financial innovation often consists not in inventing a new underlying economic exposure but in creating a more tractable representation of it. A bill of exchange transforms geographically separated trade obligations into a negotiable instrument. A share transforms a claim on an enterprise into a transferable security. A futures contract standardises a forward commitment. A stochastic differential equation transforms a family of uncertain future prices into a manipulable mathematical object. An order book transforms dispersed intentions to trade into an observable data structure.

2.3 Oikonomia, scholastic economics, guilds, and moral constraints

Classical Greek *oikonomia* concerned household management and the appropriate acquisition and use of resources. Xenophon and Aristotle treated economic activity within a normative and political framework rather than as an autonomous science [17, 388]. Medieval scholastic thought likewise analysed trade, the just price, lending, and usury under moral and theological constraints; Thomas Aquinas is a canonical reference [16]. The point for the present history is not that scholasticism contained modern asset pricing, but that market rules and moral rules were inseparable. That remains true in a transformed form: today’s quantitative decisions are embedded in market-conduct rules, capital constraints, suitability requirements, exchange protocols, and model governance.

Guild systems added professional regulation. They controlled entry, training, standards, and permissible competition. Modern quants do not form a guild in the medieval sense, yet the profession developed its own informal credential system: advanced degrees, specialist textbooks, programming competence, internal model-validation standards, interviews dominated by mathematical problem solving, and role-specific communities of practice.

2.4 Mercantilism, political economy, and capitalism

The early modern expansion of long-distance trade increased the economic value of calculation. Mercantilist writers analysed trade balances, bullion, state power, navigation, and commercial policy. Antonio Serra’s 1613 treatise is an early systematic contribution to political economy [339]. James Steuart’s *Inquiry into the Principles of Political Oeconomy* (1767) explicitly attempted to organise commercial and monetary questions into a science [356]. Adam Smith’s *Wealth of Nations* (1776) supplied a radically influential account of specialisation, exchange, prices, and self-interest [349].

The development of capitalism was not a single event but a gradual recombination of private property, wage labour, joint-stock organisation, credit, tradable claims, and expanding markets [65, 319]. For the eventual quant profession, two features were decisive. First, financial claims became sufficiently numerous and standardised to support secondary markets. Second, the scale of intermediation became large enough that small improvements in pricing, hedging, execution, or capital allocation could be worth employing specialists to obtain.

3 Money, credit, exchanges, and the first financial infrastructures

3.1 Merchant banking and the bill of exchange

Long-distance commerce creates a difficult computational problem: values must be compared across currencies, locations, dates, credit qualities, and legal jurisdictions. Medieval and Renaissance merchant banking responded with account books, correspondent networks, and bills of exchange. The Medici Bank is especially well documented. Its branches and partnerships handled deposits, transfers, exchange transactions, and commercial finance across European centres; the surviving evidence shows the importance of bookkeeping and foreign-exchange calculation to the bank's operation [103, 104].

It would be anachronistic to call a Medici employee a “quant.” Nevertheless, merchant banking reveals three components that later become quantitative finance: a distributed information network, a state representation of positions and obligations, and an economic incentive to compare prices across locations. Spatial arbitrage is among the oldest recognisably quantitative trading ideas because the logic can be expressed as an inequality: after fees, transport, credit, and timing are accounted for, equivalent claims should not command persistently different prices.

3.2 Bourses: Antwerp, Hamburg, and London

The word *bourse* came to denote a place in which merchants and financiers met to transact. Antwerp's sixteenth-century bourse is important because it was a purpose-built commercial exchange before the age of continuously traded corporate shares. Hamburg developed a merchant exchange in the same broad European network, while London's Royal Exchange, founded in the sixteenth century, explicitly drew on continental precedents. Such institutions did more than provide buildings. They concentrated information, regularised meeting times, made counterparties easier to find, and created social mechanisms for enforcing norms.

Concentration is quantitatively consequential. A market with more participants can support tighter spreads, greater specialisation, and more informative prices. The later electronic limit order book is in one sense an extreme abstraction of the bourse: a common venue replaced by a common protocol and shared state.

3.3 The VOC and the Amsterdam securities market

The establishment of the Dutch East India Company (VOC) in 1602 and the subsequent secondary trading of its shares mark a foundational episode in the history of securities markets. The VOC combined long-lived corporate capital with transferable ownership claims; Amsterdam developed a deep ecosystem of shares, public debt, credit, and derivative-like transactions [146, 296, 315]. The significance lies less in any single priority claim than in the conjunction of durable securities, active secondary trading, specialised intermediaries, and financial innovation.

Joseph de la Vega's *Confusion de Confusiones* (1688) is an extraordinary contemporary account of the Amsterdam market [102]. It describes behaviour that remains recognisable: rumours, leverage, optionality, squeezes, crowd psychology, and strategic interaction. The vocabulary changes, but the central microstructural problem is stable: traders act on different information under uncertainty while attempting to infer one another's intentions from prices and transactions.

3.4 From merchant houses to trading desks

A modern trading desk is not merely a room with people who trade. It has an organisational boundary, a defined inventory or mandate, information systems, decision rights, profit-and-loss attribution, risk limits, execution and settlement processes, and usually specialist support. Those features appeared gradually. Merchant banks such as Barings developed distinct lines of business in government securities, bills, foreign exchange, commodities, and international finance [395]. By the nineteenth century, large financial institutions increasingly separated origination, dealing, brokerage, settlement, and accounting.

Accordingly, it is more defensible to describe Renaissance merchant banks and eighteenth- and nineteenth-century houses as *antecedents* of the modern trading desk than to assign a single first date. Organisational specialisation is the important variable. Once a group has a repeated market-making or risk-taking mandate and its own position book, the economic case for dedicated pricing and risk methods becomes much stronger.

4 Information speed: from couriers to electronic feeds

4.1 Speed as a trading technology

The idea that information arriving earlier can be monetised is centuries old. The basic inequality is simple. If a trader receives a price-relevant signal at time t and competitors receive it at $t + \Delta$, the signal has value only while prices have not fully adjusted. The economic return to faster communication therefore depends on both the size of the informational advantage and the market's response time. Modern latency arbitrage is a technologically extreme form of an old structure.

Historical markets repeatedly rewarded investment in faster news: private couriers, ships, carrier pigeons, optical signals, telegraph lines, dedicated telephone circuits, microwave links, and co-located servers. The medium changed by orders of magnitude, while the objective remained the reduction of Δ . Telegraphy was a qualitative discontinuity because it separated communication speed from physical transport [195, 353]. The stock ticker then converted recent transactions into a machine-readable stream for its time, changing both the scale and rhythm of price observation.

4.2 The ticker and the datafication of markets

For quantitative work, speed is only half the story. The other half is data regularity. A handwritten price copied once per day affords one class of method; a time-stamped sequence of trades affords another; a full-depth order book with microsecond events affords another again. Each increase in market-data granularity enlarges the space of strategies and the burden of computation.

This is why the history of market data belongs inside the history of mathematical finance. Statistical arbitrage, volatility estimation, optimal execution, and order-book modelling do not merely require theories. They require observations at the relevant frequency and a storage-and-compute system capable of processing them before their economic value decays.

5 Derivatives before modern mathematical finance

5.1 Forward commitments, options, and organised futures

Contracts contingent on future prices long predate modern derivatives exchanges. Forward purchase commitments naturally arise in trade because production, transport, and payment are separated in time. Optional contracts arise when parties are willing to exchange a premium for asymmetric future rights. By the seventeenth century, the Amsterdam market supported transactions recognisably related to puts, calls, forwards, and repos [102, 315].

The nineteenth century brought increasing standardisation. Organised commodity exchanges, including the Chicago Board of Trade, helped transform bespoke forward commitments into exchange-traded futures with standard contract terms, margin arrangements, and centralised price discovery. Standardisation is a prerequisite for scalable quantitative modelling: a model can be reused across thousands of contracts only if contracts can be parameterised rather than individually narrated.

5.2 Castelli’s 1877 theory of options

Charles Castelli’s *The Theory of “Options” in Stocks and Shares* (1877) is particularly instructive because it describes calls, puts, combinations, premiums, expiration conventions, and hedging logic in a language immediately recognisable to a modern derivatives practitioner [82]. Castelli understood that the option premium should be larger when the underlying exhibits more violent fluctuations. That is not yet an option-pricing model in the modern probabilistic sense, but it identifies the central comparative-static relation between option value and volatility.

The gap between Castelli and Black–Scholes–Merton is therefore not the discovery that options have nonlinear payoffs. Traders already knew that. The mathematical revolution was to show how, under explicit assumptions, a dynamically rebalanced position in the underlying and cash could replicate a contingent claim. The price could then be derived from no-arbitrage rather than from an expected payoff under the physical distribution alone.

6 From probability to Bachelier

6.1 Probability becomes a language of prices

Quantitative finance required a mathematical language in which uncertain future prices could be objects of calculation. Probability theory developed through problems in games, insurance, astronomy, statistics, and physics before it acquired its modern measure-theoretic foundation. Brownian motion entered mathematics through the convergence of physical and probabilistic ideas, notably in the work of Einstein and Wiener [119, 381]. Kolmogorov’s 1933 axiomatization provided a durable formal foundation, and Itô’s stochastic calculus supplied the differential machinery later used throughout continuous-time finance [111, 213, 240].

6.2 Bachelier’s 1900 thesis

Louis Bachelier’s 1900 thesis, *Théorie de la spéculation*, is the canonical starting point for mathematical finance [30]. In modern notation, its simplest price model is an arithmetic Brownian motion,

$$dS_t = \mu dt + \sigma_B dW_t, \tag{1}$$

where W_t is a Wiener process and σ_B controls the scale of price fluctuations. Bachelier also treated option-like claims and diffusion equations decades before stochastic calculus had been formalised in its modern form.

The model permits negative prices, which is awkward for many securities, but that criticism can obscure the historical achievement. Bachelier established a new kind of object: a market price represented by a stochastic process. Once prices are processes, a finance problem can be attacked by probability, partial differential equations, asymptotic arguments, simulation, and statistical inference. The intellectual possibility of the modern quant is already visible.

6.3 Lognormal prices and the road to continuous-time finance

A common alternative assumes geometric Brownian motion,

$$dS_t = \mu S_t dt + \sigma S_t dW_t, \quad (2)$$

with solution

$$S_t = S_0 \exp \left[\left(\mu - \frac{1}{2} \sigma^2 \right) t + \sigma W_t \right]. \quad (3)$$

This ensures positivity and makes log returns Gaussian with constant variance rate. The assumption later became central to the canonical Black–Scholes–Merton model, but empirical work would repeatedly show that actual return distributions exhibit heavy tails, volatility clustering, asymmetry, jumps, and changing market regimes.

7 Political Economy in the Cold War: Shared Mathematics, Rival Institutions, and the Rise of Quantitative Finance

7.1 Political economy as a problem of calculation

The economic confrontation of the Cold War is often described as a contest between planning and markets, or between state ownership and private ownership. That description is institutionally correct but mathematically incomplete. A striking feature of twentieth-century political economy is that economists and mathematicians working on opposite sides of the ideological divide repeatedly encountered the same formal problems: how to allocate scarce resources subject to constraints; how to represent opportunity cost; how to aggregate dispersed information; how to make decisions sequentially under uncertainty; how to coordinate many agents without solving every local problem in one place; and how to make a large optimization problem computationally tractable. The resulting mathematical languages—linear and convex programming, duality, general equilibrium, stochastic processes, optimal control, dynamic programming, game theory, information theory, and numerical algorithms—were not intrinsically socialist or capitalist. They were institutional technologies whose economic meaning depended on where objectives were specified, who possessed information, how constraints were enforced, and how the resulting multipliers or prices were allowed to affect behaviour [11, 226, 241, 289].

This observation matters for the history of quantitative finance. Modern quantitative finance is usually narrated as a sequence running from Bachelier through Markowitz, Black–Scholes–Merton, martingale pricing and contemporary computational finance. That genealogy is correct, but it hides a broader political-economic setting. During the Cold War, the West was simultaneously developing a mathematical theory of decentralized allocation and a practical infrastructure of securities markets,

while the Soviet Union was developing mathematical programming, cybernetics and optimization in an attempt to make planning more rational, adaptive and computationally scalable. The two programmes differed fundamentally in property rights, authority and the locus of decision-making, but they often shared mathematical machinery. The 1975 award of the economics prize to Leonid Kantorovich and Tjalling Koopmans “for their contributions to the theory of optimum allocation of resources” is an unusually vivid institutional recognition of that convergence [228, 303].

The argument of this section is therefore not that quantitative finance “proved” the superiority of decentralized markets, nor that Soviet mathematical economics secretly converged on capitalism. It is narrower and more interesting. First, the Cold War calculation debate accelerated the development of formal tools for allocation, information and control. Second, several of those tools were explicitly capable of both centralized and decentralized interpretations. Third, financial economics supplied a theory of decentralized allocation under uncertainty in which securities prices, portfolios and trading strategies became computational objects. Finally, quantitative finance helped make decentralized capital markets more scalable by supplying common rules for valuation, hedging, risk decomposition and arbitrage. Those achievements were conditional, however, on assumptions about information, liquidity, incentives, market completeness and institutional enforcement, and later work made the limitations explicit [165, 269].

7.2 The socialist calculation debate: prices, knowledge, and computability

The prehistory of the Cold War debate lies in the socialist calculation controversy. Ludwig von Mises argued that without market prices for capital goods generated through exchange, a socialist commonwealth would lack the economic calculation needed to compare alternative uses of resources [290]. Friedrich Hayek shifted the emphasis from mere arithmetic to dispersed knowledge. His central claim was that economically relevant information is fragmentary, local and often tacit; a price system is valuable because it allows agents to react to information they do not themselves possess in full [188]. In this view, the market is not simply a device for solving a known optimization problem. It is also an information-processing institution.

Oskar Lange and Abba Lerner responded by observing that competitive-equilibrium conditions could, at least in principle, be imitated by a socialist system. Lange’s trial-and-error mechanism asked a central planning board to announce accounting prices, allow enterprises to choose inputs and outputs subject to marginal rules, and adjust prices in response to shortages and surpluses [245, 246, 254]. The debate therefore moved quickly from ownership alone to mechanisms: what information must be transmitted, in what form, and by whom? Later retrospectives rightly emphasize that the disagreement cannot be reduced to whether equations have solutions. The central questions concern discovery, incentives, communication, institutional commitment and adaptation [205, 206, 313].

This is precisely the point at which mathematical economics becomes politically ambiguous in a productive sense. A system of equations describing a Pareto-efficient allocation does not specify whether the allocation is discovered through competitive bidding, computed by a ministry, negotiated among firms, or generated by a hybrid mechanism. General equilibrium, mathematical programming and mechanism design separate the formal allocation problem from the institution that implements it. The Cold War debate can thus be read as a debate over implementations of related mathematical objects.

7.3 Kantorovich, Novozhilov, Nemchinov, and the Soviet optimization programme

Leonid Kantorovich’s 1939 work on production planning is the clearest Soviet example. Faced with practical production-allocation problems, he formulated what would later be recognized as linear programming and introduced *resolving multipliers*, quantities that behaved mathematically like dual variables or shadow prices [226–228]. In modern notation, consider the primal allocation problem

$$\max_{x \geq 0} c^\top x \quad \text{subject to} \quad Ax \leq b, \quad (4)$$

with dual

$$\min_{y \geq 0} b^\top y \quad \text{subject to} \quad A^\top y \geq c. \quad (5)$$

The components of y measure the marginal value of relaxing resource constraints. In a planning interpretation, they are internal valuations useful for comparing production plans. In a competitive interpretation, closely related multipliers can be read as scarcity prices. The mathematical duality is exact; the institutional interpretation is not.

Kantorovich was careful to present his multipliers as devices for improving socialist calculation rather than as an argument for *laissez-faire* pricing. Yet the language of scarcity valuation created a bridge between planning and price theory. Viktor Novozhilov pushed further toward opportunity-cost reasoning, emphasizing the economic cost of scarce resources and the need for valuations that reflected alternatives forgone. Vasily Nemchinov promoted mathematical methods, input–output analysis and economic cybernetics, and helped institutionalize mathematical economics in Soviet higher education. Nikolai Fedorenko, as the founding director of the Central Economic-Mathematical Institute (CEMI), became a central figure in the programme known as the System for the Optimal Functioning of the Socialist Economy (SOFE) [122, 135, 236, 248, 300, 305].

SOFE is particularly important for the argument here because it was not simply a project to make central planning more centralized. It combined global objectives with attempts at *indirect centralization*: enterprises could be given more operational discretion while prices, profitability measures, contracts, credit and shadow valuations were used as instruments to align local decisions with system-wide objectives [123, 127]. This resembles, at the level of formal decomposition, a problem that also appears in decentralized optimization: a global objective is broken into local subproblems whose solutions are coordinated by messages, multipliers or prices. The political difference lies in who defines the objective, who owns the resources, what messages are permissible, and whether agents are free to reject the prescribed objective.

The Soviet programme also developed an institutional base. After the post-Stalin rehabilitation of cybernetics, mathematical economics programmes appeared at Leningrad University and Moscow State University; CEMI was founded in 1963; and *Economics and Mathematical Methods* began publication in 1965 [236]. These developments made mathematical reasoning a partially autonomous language within a political economy whose official doctrine remained suspicious of any implication that decentralized prices possessed independent authority.

7.4 Dantzig, Koopmans, Leontief, and the Western operations-research state

The Western counterpart was not born in a purely *laissez-faire* environment. George Dantzig developed the linear-programming model and simplex method in connection with U.S. Air Force planning problems, while operations research, systems analysis and game theory expanded rapidly through military and government institutions during and after the Second World War [96, 97,

141, 289]. Tjalling Koopmans and the Cowles Commission developed *activity analysis*, bringing production, transportation, substitution, convexity and computational questions into a common mathematical framework [241]. Wassily Leontief’s input–output analysis supplied another matrix-based representation of interdependence that could be used in capitalist forecasting, wartime mobilization, development planning or socialist accounting [252, 253].

The parallel is historically instructive. Soviet mathematical economists were trying to replace crude administrative norms with optimization and economically meaningful valuations. Western operations researchers were trying to replace ad hoc military and industrial planning with optimization and computation. In both settings, the immediate object was frequently a large planning problem. The ideological divergence appeared not in the algebra of $Ax \leq b$, but in the surrounding theory of institutions.

Koopmans made the connection between programming and decentralized markets especially explicit. Activity analysis showed that under convexity and appropriate conditions, supporting prices can decentralize efficient production decisions: each productive unit can optimize locally at common prices while the aggregate allocation satisfies system-wide feasibility [241, 242]. The same convex duality therefore supports two readings. A planner may use dual variables to evaluate scarce resources; a market may use prices to coordinate separate maximization problems. The 1975 Kantorovich–Koopmans prize recognized a mathematical theory whose institutional realizations could differ sharply [303].

7.5 The common language of duality

The shared machinery can be stated abstractly. Let aggregate welfare or output be separable across units,

$$\max_{x_1, \dots, x_n} \sum_{i=1}^n f_i(x_i) \quad \text{subject to} \quad \sum_{i=1}^n A_i x_i \leq b. \quad (6)$$

Introduce a multiplier $p \geq 0$ and form

$$\mathcal{L}(x_1, \dots, x_n, p) = \sum_{i=1}^n f_i(x_i) + p^\top \left(b - \sum_{i=1}^n A_i x_i \right). \quad (7)$$

For a fixed p , the maximization decomposes into n local problems,

$$\max_{x_i} (f_i(x_i) - p^\top A_i x_i). \quad (8)$$

The vector p is simultaneously a Lagrange multiplier, a scarcity valuation and a candidate coordinating message. This decomposition is mathematically indifferent to whether p is computed by a central algorithm, adjusted through an administrative process, or discovered through decentralized exchange.

This fact helps explain why the same vocabulary repeatedly crossed ideological boundaries. “Shadow prices”, “opportunity costs”, “decomposition”, “feedback”, “optimality” and “equilibrium” can describe quite different institutions. It also anticipates a central fact of quantitative finance: the no-arbitrage price of a contingent claim can be represented as a dual object—a linear pricing functional or expectation under an equivalent martingale measure—while individual trading desks make local decisions subject to those market prices.

7.6 Hayek, Hurwicz, and the mathematics of decentralization

Hayek’s critique of central planning was fundamentally informational: the planner does not possess, and may not even be able to articulate, the local knowledge embedded in decentralized decisions [188]. Leonid Hurwicz transformed this broad insight into a mathematical research programme. He represented resource-allocation processes as systems in which agents exchange messages and asked what information a mechanism must transmit to implement efficient outcomes [22, 23, 205, 206]. Kenneth Arrow and Hurwicz studied decentralized computation and stability, while later mechanism design added strategic incentives to the informational problem.

This work complicates any simple ideological mapping. Hurwicz had experienced both Soviet and Western institutions and developed a theory that could compare allocation mechanisms without assuming in advance that “market” or “planning” named a single mechanism. In this framework, a competitive market is one possible message-processing system among others. Its attraction lies in informational economy under certain conditions: agents need not transmit their complete utility functions or technologies to a central authority. Prices can act as low-dimensional coordinating signals. But whether a proposed mechanism is incentive-compatible, stable, robust to private information or computationally feasible becomes a mathematical question rather than a slogan.

The same perspective also clarifies the Soviet reformers’ interest in prices. The issue was not necessarily whether to abolish planning, but whether a high-dimensional command system could be decomposed so that local units responded to economically meaningful signals. In both traditions, decentralization became a problem of architecture.

7.7 Cybernetics and OGAS: an alternative information architecture

Viktor Glushkov’s proposed All-State Automated System for the Gathering and Processing of Information for the Accounting, Planning, and Governance of the National Economy (OGAS) represented the most ambitious Soviet attempt to attack the information problem technologically. Glushkov envisaged a nationwide hierarchy of computing centres linking enterprises and regional nodes to central facilities, with the aim of accelerating accounting, planning and control [149, 235, 314, 329]. It is tempting to describe OGAS as the exact opposite of a market: a giant computer intended to replace prices. That is too crude. Glushkov’s architecture itself was distributed and hierarchical, and Soviet economic cybernetics increasingly incorporated feedback, local calculation and delegated decision-making. What remained centralized was the ultimate political authority over objectives and the permissible structure of exchange.

The comparison with modern financial networks is revealing. Both systems confront distributed information and require common protocols, rapid communication, accounting identities and feedback. Yet modern securities markets do not ask a central machine to solve a single social objective. Exchanges, dealers, investors, clearing houses and risk systems solve heterogeneous local problems and interact through prices, orders, collateral constraints and contracts. The market therefore resembles a distributed computation whose global state is emergent rather than explicitly optimized by one authority. Quantitative finance supplies some of the algorithms run by the nodes of that distributed system.

7.8 Optimal control on both sides: Pontryagin and Bellman

Another striking example of shared machinery is optimal control. In the Soviet Union, Lev Pontryagin, Vladimir Boltyanskii, Revaz Gamkrelidze and Evgenii Mishchenko developed the

maximum principle for dynamic optimization [320]. In the United States, Richard Bellman developed dynamic programming and the principle of optimality [44]. The two approaches are mathematically related but computationally and conceptually distinct: one emphasizes necessary conditions expressed through costate variables and Hamiltonians, the other recursive value functions.

Both were shaped by engineering, military and planning problems, and both later became standard tools in economics and finance. A social planner can use optimal control to choose an investment path. A firm can use it to manage inventories. A portfolio manager can use it to choose consumption and investment. A market maker can use it to control inventory. The Hamilton–Jacobi–Bellman equation of continuous-time finance is therefore another example of political neutrality at the level of mathematics: the control problem acquires its political-economic meaning from the objective function, admissible controls and institutional constraints.

7.9 Probability, measure change, and Soviet mathematical contributions to financial machinery

The shared mathematical substrate extends beyond optimization. Andrei Kolmogorov’s axiomatization of probability placed stochastic reasoning on a measure-theoretic foundation [239]. Soviet probability theory became one of the strongest mathematical traditions in the world. Igor Girsanov’s 1960 theorem on absolutely continuous changes of measure later became central to risk-neutral pricing: under suitable conditions, a diffusion with one drift under a physical probability measure can be represented with another drift under an equivalent measure [152]. The theorem was not created for finance, but it became one of the key pieces of mathematical machinery by which modern finance separates expected returns under the empirical measure from arbitrage-free valuation under a pricing measure.

Leonid Khachiyan’s 1979 proof that linear programming is polynomial-time solvable by an ellipsoid method provides a later Cold War example of the same transnational mathematical economy [230]. An algorithm developed in the Soviet mathematical tradition immediately became part of the global theory of optimization. Here again, the computational method does not encode a political regime. It can solve a production plan, a portfolio problem, a network-flow problem or a decentralized subproblem.

One should therefore distinguish *political participants* in the planning-versus-market debate from *mathematical contributors* whose work supplied the common technical language. Kantorovich, Novozhilov, Nemchinov, Fedorenko, Glushkov, Mises, Hayek, Lange, Lerner, Koopmans, Arrow and Hurwicz belong directly to the former group. Kolmogorov, Girsanov, Pontryagin, Bellman, Khachiyan and many others belong principally to the latter, even though their mathematics profoundly affected what planners and market theorists could formulate.

7.10 From general equilibrium to securities as decentralized state-contingent allocation

Western general-equilibrium theory supplied a particularly direct bridge from political economy to finance. Arrow and Debreu formalized the existence of competitive equilibrium under convexity and other assumptions [21, 105]. Arrow’s work on securities showed that state-contingent claims can support efficient risk allocation without requiring every commodity in every future state to be traded separately at the initial date [20]. A security is thus an information-and-contract technology: it transfers purchasing power across states of the world.

Roy Radner extended equilibrium analysis to a sequence of markets in which agents trade over

time as information is revealed [327]. This is much closer to the institutional world of actual finance. Agents do not solve a social planner’s stochastic program. They form expectations, choose portfolios and trade securities at quoted prices. Under appropriate conditions, the resulting competitive allocation can possess welfare properties that are mathematically related to the solution of a planner’s problem.

The relationship is easiest to see through the first-order conditions. Suppose agent i maximizes expected utility

$$\max_{c_i(\omega)} \sum_{\omega \in \Omega} \pi(\omega) u_i(c_i(\omega)) \quad (9)$$

subject to a budget constraint expressed in state prices $q(\omega)$. The multiplier on the budget constraint links marginal utility across states to $q(\omega)$. In a complete-market equilibrium, the same state-price vector coordinates all agents. What a planner would express as multipliers on state-contingent resource constraints can therefore reappear as market prices for contingent claims. The formal resemblance to linear-programming duality is not accidental.

7.11 Portfolio theory: decentralizing risk rather than output

Harry Markowitz’s portfolio theory moved optimization from the allocation of physical production resources to the allocation of financial wealth under risk [277, 278]. In the simplest mean–variance form,

$$\min_w \frac{1}{2} w^\top \Sigma w \quad \text{subject to} \quad \mu^\top w \geq \bar{r}, \quad \mathbf{1}^\top w = 1. \quad (10)$$

This is a quadratic program. Its lineage runs directly through the same optimization culture as production planning and operations research. Yet its institutional interpretation is decentralized: each investor chooses a portfolio according to private objectives and beliefs, taking security prices and estimated risk parameters as inputs.

Tobin’s separation results and the capital-asset-pricing work of Sharpe, Lintner and Mossin then connected individual portfolio optimization to market equilibrium [256, 292, 340, 369]. The market portfolio and beta became reduced-form coordinates for the price of systematic risk. Modigliani and Miller likewise used arbitrage reasoning to establish conditions under which corporate financing choices do or do not affect firm value [291]. These theories did not merely describe investment practice; they offered decentralized consistency conditions. If many agents optimize and can trade freely, relative prices cannot be arbitrary.

Quantitative finance thus extended the political-economic role of prices. A commodity price coordinates production and consumption. A security price coordinates intertemporal and state-contingent transfers. A covariance matrix represents the joint uncertainty faced by private agents. A beta or stochastic discount factor compresses parts of that uncertainty into a price relation. The market becomes a distributed mechanism for allocating not only goods, but risk.

7.12 Information in prices: efficient markets and their limit

The efficient-markets literature gave an especially Cold-War-compatible interpretation of decentralized finance: competitive trading aggregates dispersed information into prices. Fama’s empirical and theoretical work systematized the efficient-markets hypothesis [133]. Yet the strongest version contains a paradox that echoes the older calculation debate. If prices already reveal all costly information, why should anyone pay to acquire information? Grossman and Stiglitz showed that perfectly informationally efficient markets are impossible when information acquisition is costly [165].

Markets need informed traders to make prices informative, but informed traders need compensation for producing information.

This result is important because it blocks a triumphalist reading of the Cold War history. A decentralized market does not costlessly solve an omniscient planner’s problem. It has its own informational frictions, rents, strategic behaviour and institutional requirements. Hayek’s insight that knowledge is dispersed survives, but the market’s aggregation of that knowledge is itself an endogenous and costly process. Quantitative finance helps structure the process; it does not abolish the underlying economic problem.

7.13 Black–Scholes–Merton: local arbitrage as a coordinating principle

The Black–Scholes–Merton framework brought a new kind of decentralization to financial pricing [59, 283]. Consider an idealized risky asset

$$dS_t = \mu S_t dt + \sigma S_t dW_t. \quad (11)$$

A derivative with value $V(t, S_t)$ can, under the classical assumptions, be locally replicated by trading the underlying and a riskless asset. The no-arbitrage argument removes the asset’s expected return μ from the valuation equation. This was conceptually profound: a derivative’s price need not be set by a central authority or by specifying every investor’s utility function. Local trading relations constrain the admissible price.

The economic force is arbitrage. If two portfolios generate the same future cash flows but have different prices, decentralized traders have an incentive to buy the cheaper and sell the dearer. Their trades push prices toward consistency. The mathematical model therefore converts a local profit-seeking rule into a system-wide pricing restriction. In that limited sense, arbitrage is a decentralized enforcement algorithm.

This interpretation should not be exaggerated. Real markets have transaction costs, funding constraints, discrete trading, model uncertainty, position limits and jumps. Nonetheless, the Black–Scholes–Merton framework provided a common numerical language in which exchanges, dealers, asset managers and risk controllers could quote, hedge and compare contingent claims. It increased the interoperability of decentralized markets.

7.14 Harrison, Kreps, and Pliska: the martingale formulation

J. Michael Harrison, David Kreps and Stanley Pliska supplied the general mathematical architecture behind this idea. Harrison and Kreps linked absence of arbitrage in securities markets to martingale measures, while Harrison and Pliska developed a continuous-trading model using semimartingales and stochastic integrals [177, 179, 180]. In modern language, an arbitrage-free market admits, under suitable technical conditions, a positive linear pricing rule; after normalization this can often be represented by an equivalent probability measure Q under which discounted traded prices are martingales.

Schematically,

$$V_t = B_t \mathbb{E}^Q \left[\frac{X_T}{B_T} \middle| \mathcal{F}_t \right], \quad (12)$$

where B_t is a numeraire and X_T the terminal payoff. Girsanov’s theorem is one of the central tools for constructing the required change of measure in diffusion models [152].

Equation (12) is another dual representation. Instead of solving every investor’s utility-maximization problem, one characterizes prices through a linear functional consistent with traded assets and no

arbitrage. When markets are complete, the pricing measure is unique; when they are incomplete, there are generally multiple compatible pricing measures and preferences or additional criteria re-enter the problem. The theorem therefore identifies exactly what decentralized trading can determine by arbitrage alone and what it cannot.

7.15 Quantitative finance as infrastructure for decentralized markets

It is useful to distinguish several ways in which quantitative finance supported the scaling of decentralized markets during the later Cold War and after it.

Common valuation conventions. Models supplied common reference prices, Greeks, yield curves, volatilities and risk measures. Agreement on a model was not required in every detail; the existence of a shared quantitative vocabulary reduced the dimensionality of negotiation. Traders could disagree about implied volatility while agreeing on how a quoted volatility maps into a price.

Replication and hedging. Replication decomposed complex contingent claims into positions in simpler traded instruments. This allowed specialized institutions to create, warehouse and redistribute risks rather than requiring each final investor to understand or bear the original exposure in its raw form.

Arbitrage relations. No-arbitrage restrictions tied prices across instruments and venues. Such restrictions are local enough to be enforced by self-interested traders, yet global enough to generate coherent term structures and relative-value systems when markets are sufficiently liquid.

Portfolio aggregation and risk budgeting. Markowitz-style covariance methods, factor models and later risk systems allowed large decentralized portfolios to be treated as structured objects. This made delegation easier: a fund, desk or trader could operate locally subject to risk limits expressed in common quantitative units.

Market completion and state contingency. Derivatives expanded the set of tradable state-contingent payoffs. In Arrow–Debreu language, they can move an economy closer to spanning relevant states of nature, although real markets remain incomplete and the creation of new claims also creates counterparty, liquidity and systemic risks.

Automation. Once prices, portfolios and constraints were represented mathematically, they became amenable to software. Quantitative finance thus converted parts of market coordination into machine-executable procedures: pricing libraries, optimization engines, electronic market making, margin systems and automated hedging. In this respect, there is an ironic parallel with Soviet cybernetics. Both sought to increase the scale of economic coordination through computation; one placed computation inside a command hierarchy, the other increasingly distributed it among competing market participants and infrastructures.

Donald MacKenzie’s notion of financial models as an “engine, not a camera” is useful here. Models can become constitutive of markets when participants quote, hedge, regulate and build systems around them [269]. The Black–Scholes model did not merely describe an already mature options market; standardized quantitative valuation helped the market itself become more liquid, comparable and organizationally tractable.

7.16 What was genuinely shared—and what was not

The strongest version of the “shared machinery” thesis can be summarized in Table 1. The table also indicates why mathematical similarity should not be mistaken for institutional equivalence.

Mathematical object	Planning / Soviet interpretation	Decentralized-market / finance interpretation
Linear/convex program	Feasible production plan maximizing an administratively specified criterion	Firm, investor, exchange or intermediary solves a local objective under market prices and constraints
Dual variable / multiplier	Resolving multiplier; objectively conditioned valuation; internal scarcity signal	Competitive price, shadow funding cost, state price, Lagrange multiplier in portfolio/risk optimization
Decomposition	Divide a national plan into enterprise or sector subproblems	Separate agents optimize locally and coordinate through prices, contracts, collateral and order flow
Feedback	Plan revision, cybernetic control, administrative correction	Price adjustment, arbitrage, market making, inventory control, margin feedback
Optimal control	State chooses dynamic investment/production path	Investor, firm or dealer chooses dynamic portfolio, execution or hedging policy
Probability model	Forecasting and stochastic control	Asset pricing, portfolio choice, risk measurement and derivative valuation
Change of measure	General theorem of stochastic-process transformation	Risk-neutral valuation and martingale pricing
Computer network	OGAS-style information collection and planning hierarchy	Exchanges, market-data networks, clearing systems and distributed trading infrastructure

Table 1: Shared mathematical machinery and divergent institutional interpretations.

The crucial differences remain outside the mathematics. Soviet planning assigned the state a privileged role in defining objectives, ownership and admissible enterprise behaviour. Competitive financial markets permit heterogeneous objectives and use property rights, contracts and bankruptcy to discipline local decisions. The market does not maximize a single explicit social objective. Conversely, a planner may be able to impose distributional, strategic or industrial objectives that cannot be represented by competitive prices without additional institutions.

The informational assumptions differ as well. A linear programme such as (4) assumes that the coefficients, constraints and objective have been specified. Hayek’s objection begins precisely where that assumption fails. Mechanism design asks what must be communicated to construct an outcome when preferences and technologies are private. Quantitative finance likewise distinguishes between parameters assumed known by a model and information that must be inferred from prices or data.

7.17 The political ambiguity of mathematical efficiency

The history also cautions against identifying mathematical optimality with political legitimacy. An allocation can be efficient relative to an objective that citizens reject. A market can be arbitrage-free and still generate extreme inequality, externalities or systemic fragility. A plan can solve an optimization problem exactly and still encode distorted priorities. Convex duality tells us how to

support an optimum; it does not tell us whose optimum ought to be supported.

This point was visible on both sides of the Cold War. Soviet mathematical economists repeatedly faced the problem of choosing an “optimality criterion” for the economy: output, consumption, growth, welfare, or some politically selected combination. Western welfare economics confronted social-choice impossibility, distribution and incomplete information. Arrow’s social-choice theorem itself showed the difficulty of aggregating individual rankings into a coherent social ordering under plausible axioms [18]. The shared machinery therefore sharpened rather than eliminated political questions.

7.18 A qualified political-economic interpretation of quantitative finance

Within these limits, quantitative finance can be understood as one of the technologies that made decentralized market coordination more capable. It did not create private property, enforce contracts, provide monetary stability or guarantee competition. Nor did it eliminate the need for regulation. Its contribution was more specific: it made heterogeneous financial claims calculable in a common language; it connected prices across states, maturities and instruments; it supplied local decision rules for hedging and portfolio construction; and it made risk transferable among specialized agents.

In a centralized planning system, a large part of the coordination problem is presented explicitly: the planner seeks an allocation satisfying global constraints. In a decentralized financial system, the global problem is not normally written down. Instead, many agents solve local optimization problems. Prices, arbitrage relations and balance-sheet constraints couple those problems. Quantitative finance supplies models for those couplings.

One can express the contrast schematically. The planner asks

$$\text{choose } x \text{ to maximize } W(x) \text{ subject to system constraints;} \quad (13)$$

whereas the market theorist asks whether there exists a price system p such that

$$x_i \in \arg \max_{x \in X_i} U_i(x; p), \quad i = 1, \dots, n, \quad (14)$$

and the individually chosen allocations clear markets. Financial economics adds time and uncertainty, so that p becomes a system of spot prices, forward prices, state prices and stochastic pricing kernels. The difference is not that one side uses mathematics and the other does not. It is where the coordination problem is solved.

7.19 Failures and counterexamples: why decentralization still needs institutions

A balanced account must also note that the mathematical theory of decentralized markets contains its own failure modes. Incomplete markets prevent full state-contingent risk sharing. Asymmetric information may produce adverse selection and moral hazard. Market power invalidates price-taking assumptions. Transaction costs and illiquidity break exact replication. Leverage and collateral create nonlinear feedback. Grossman–Stiglitz informational inefficiency means that information aggregation is costly. Financial crises show that individually rational risk management can generate collectively destabilizing behaviour.

These limitations do not reinstate the central-planning solution. They show instead that “market versus planner” is too coarse a dichotomy. Real economies use layered mechanisms: markets, firms, governments, clearing houses, central banks, exchanges, regulators, internal transfer prices, auctions,

risk limits and algorithmic controls. The mature lesson of mechanism design and quantitative finance is that institutional architecture matters as much as optimization.

The Soviet experience supplies a parallel warning. A sufficiently large computer network could not by itself solve questions of incentives, data quality, bureaucratic rivalry or objective selection. Histories of OGAS emphasize that organizational and political conflicts were as consequential as computational constraints [149, 314]. The failure of a central network and the imperfections of decentralized financial markets are therefore not mirror-image proofs of opposing ideologies; both demonstrate that computation is embedded in institutions.

7.20 The Cold War legacy

The most revealing symbolic event may remain the 1975 Kantorovich–Koopmans prize. A Soviet mathematician of planning and a Western economist of activity analysis were honoured jointly for optimum resource allocation. Their methods demonstrated that efficiency, shadow valuation and decomposition belong to a common mathematical structure even when the surrounding political economies differ.

From this perspective, the rise of quantitative finance during the second half of the Cold War was part of a broader transformation in which economies became increasingly describable, and increasingly operable, through formal models and computers. The same period produced linear programming, dynamic programming, optimal control, general equilibrium under uncertainty, portfolio theory, state-contingent securities, rational-expectations models, option pricing and martingale finance. Some of these developments were nurtured by military planning, some by socialist economic reform, some by universities and securities markets. Their eventual combination in modern finance should not obscure their mixed genealogy.

The political-economic significance of quantitative finance is therefore best stated cautiously. It supplied a rigorous and executable theory of how decentralized agents can coordinate some decisions under uncertainty through prices and tradable claims. Arbitrage can enforce local consistency without a central price setter; securities can redistribute risk across agents and states; portfolio and control theory can delegate decisions while preserving quantitative constraints; and market prices can compress dispersed information. These are genuine mechanisms of decentralization.

But the mathematics is not an ideological certificate. The same dual variables that coordinate a competitive equilibrium can guide a planner. The same optimal-control equations can govern a state investment programme or an algorithmic trading strategy. The same stochastic-process theory can support industrial forecasting or derivative pricing. The deepest Cold War lesson is therefore not that one side possessed mathematics and the other did not. It is that *the locus of objectives, information, authority and adjustment determines the political meaning of shared mathematical machinery*.

That observation provides a natural bridge to the history of the quant profession. The quant is a descendant not only of Bachelier and Black–Scholes, but also of the twentieth-century science of allocation and control. The profession inherited optimization from planners and operations researchers, probability from mathematical schools that crossed national systems, equilibrium from mathematical economists, control from military and engineering research, and computational discipline from both state and corporate laboratories. Quantitative finance then recombined those tools around a distinctive institutional object: a decentralized network of tradable claims whose prices are continuously revised by heterogeneous agents.

8 Portfolio theory, equilibrium, and informational efficiency

8.1 Markowitz and optimisation

Harry Markowitz’s 1952 portfolio-selection framework transformed diversification from an aphorism into an optimisation problem [277]. For portfolio weights w , expected returns μ , and covariance matrix Σ , the mean–variance problem can be written

$$\min_w w^\top \Sigma w \quad \text{subject to} \quad w^\top \mu = \mu_*, \quad \mathbf{1}^\top w = 1. \quad (15)$$

The formula is deceptively compact. Operationally it requires estimates, matrix computation, constraints, data cleaning, and decisions about how a theoretical optimum is mapped into trades. In other words, portfolio theory created not only an academic research programme but a demand for a quantitative production process.

Tobin extended the analysis of risk and liquidity; Sharpe, Lintner, and Mossin developed equilibrium asset-pricing models relating expected return to systematic risk [256, 292, 340, 369]. Ross later proposed arbitrage pricing theory [334]. These models helped establish a new professional vocabulary of beta, factor exposure, covariance, benchmark, active risk, and optimisation.

8.2 Random prices and efficient markets

The statistical study of price changes grew alongside portfolio theory. Fama documented return behaviour and later systematised the efficient-markets literature [132, 133]. Samuelson proved that properly anticipated prices should fluctuate randomly under an appropriate model of information and valuation [337]. Yet informational efficiency is not synonymous with a claim that no one can profitably gather information. Grossman and Stiglitz showed the tension in the idea of perfectly informationally efficient markets: if information were costlessly and fully reflected in prices, no one would pay to acquire it [164].

This tension is constitutive of the quant profession. Quantitative research is economically valuable precisely because markets are neither trivially predictable nor perfectly unexploitable. The practitioner lives in the gap between an obvious inefficiency, which competition quickly destroys, and a perfectly efficient market, which would make research pointless.

9 The Black–Scholes–Merton revolution

9.1 Replication and no-arbitrage pricing

The 1973 papers of Fischer Black and Myron Scholes and of Robert C. Merton changed the relationship between mathematics and derivatives markets [58, 284]. Consider an option value $V(S, t)$ when the underlying follows Equation (2). Applying Itô’s lemma gives

$$dV = \left(V_t + \mu S V_S + \frac{1}{2} \sigma^2 S^2 V_{SS} \right) dt + \sigma S V_S dW_t. \quad (16)$$

A dynamically rebalanced portfolio holding V_S units of the underlying removes the diffusion term. Under idealised assumptions the remaining locally riskless position must earn the risk-free rate r , yielding the Black–Scholes–Merton PDE

$$V_t + \frac{1}{2} \sigma^2 S^2 V_{SS} + r S V_S - r V = 0. \quad (17)$$

For a European call with strike K and maturity T , this gives

$$C = S_0 N(d_1) - K e^{-rT} N(d_2), \quad (18)$$

$$d_1 = \frac{\ln(S_0/K) + (r + \sigma^2/2)T}{\sigma\sqrt{T}}, \quad d_2 = d_1 - \sigma\sqrt{T}. \quad (19)$$

The deepest idea is not the closed form. It is replication. Once a contingent claim can be reproduced by dynamically trading more elementary claims, no-arbitrage links their prices. This shift from expected-payoff pricing to replication pricing created a practical programme: estimate model inputs, compute values and Greeks, hedge dynamically, and analyse the residual risks when assumptions fail.

9.2 From formula to profession

A theory creates a profession when it can be embedded in a repeated industrial workflow. Listed options markets expanded in the 1970s; banks and securities firms increasingly needed models for pricing, hedging, scenario analysis, and risk control. The intellectual ecosystem widened through binomial methods [90], martingale pricing [175, 178], and term-structure models [91, 190, 194, 201, 378]. Textbooks eventually made stochastic calculus and derivatives pricing portable across institutions [41, 203, 383].

The derivatives quant became a recognisable specialist because the implementation burden was large. A desk required yield curves, calendars, conventions, interpolation, implied volatilities, numerical solvers, calibration, Greeks, scenario engines, and interfaces to trading and risk systems. The elegant PDE sat inside a far less elegant production stack. The profession formed in that interface between abstract model and operational detail.

10 The emergence of the modern quant profession

10.1 Thorp, Kelly, and systematic risk-taking

Edward O. Thorp is often described as one of the earliest modern quants because he combined mathematics, computation, statistical edge estimation, explicit bankroll management, and real-money risk-taking across games and markets. *Beat the Dealer* and, with Sheen Kassouf, *Beat the Market* made mechanical, quantified strategies central to the narrative [366, 368]. The Kelly criterion supplied a formal connection between edge, risk, and long-run growth [229].

Thorp matters historically not because no earlier trader calculated, but because the entire modern quantitative workflow is visible: formulate an exploitable structure, estimate parameters, code or mechanise the calculation, size bets according to risk, execute repeatedly, and update when the environment changes. This is closer to a contemporary systematic fund than to traditional discretionary security analysis.

10.2 Institutional portfolio quantification

The 1960s and 1970s also saw quantitative methods migrate into institutional asset management. Portfolio theory, factor models, index construction, and computing made it possible to manage large portfolios by rules rather than solely by security-by-security judgement. The historical literature on modern portfolio theory emphasises the interaction among academics and practitioners around

institutions such as Wells Fargo’s Management Sciences group [47, 48]. The important shift was organisational: mathematics moved from an advisory input to a production technology for managing portfolios.

This buy-side lineage is sometimes overshadowed by derivatives history. Yet the two branches generated different archetypes of quant. The derivatives quant asks for an arbitrage-consistent price and hedge. The portfolio quant asks for a forecast, covariance model, optimiser, and portfolio construction rule. Both require statistics and computation, but their objective functions, error tolerances, and validation regimes differ.

10.3 Simons and the systematic hedge-fund model

Jim Simons and Renaissance Technologies became the emblem of a third lineage: the research-driven systematic hedge fund. Its public history illustrates the industrialisation of statistical trading—large data sets, teams of scientists, signal combination, transaction-cost modelling, automated execution, and relentless empirical testing [396]. The precise strategies remained proprietary, but the organisational model was highly influential.

The modern systematic fund differs from a single “black box.” It is better understood as a research factory. Hypotheses are generated, data are cleaned, features are constructed, tests are run, false discoveries must be controlled, portfolios are built under constraints, costs are estimated, orders are generated, execution is monitored, and live performance feeds back into research. The central professional asset is not one equation but an iterative scientific and engineering process.

10.4 Physicists, mathematicians, engineers, and the derivatives boom

The expansion of derivatives in the 1980s and 1990s drew mathematicians, physicists, engineers, computer scientists, and statisticians into banks. Derman’s memoir provides a first-person account of this migration and of the emergence of “quant” as a workplace identity [107]. Mehrling’s intellectual biography of Fischer Black and MacKenzie’s sociology of financial modelling illuminate the same period from different directions [269, 281].

The profession’s prestige during this period was tied to a specific institutional proposition: difficult mathematics could be turned into tradable prices and risk numbers. But the proposition was always conditional. A model that prices accurately but cannot be calibrated, coded, run within latency constraints, reconciled with market conventions, or explained to risk control is not a production model. The enduring skill of the quant is therefore translation across domains.

10.5 The term “quant” and occupational identity

Dating a word is easier than dating an occupation, but even the word’s diffusion is gradual. “Quantitative analyst” and “quantifier” appeared before “quant” became ubiquitous. By the 1990s, journalistic and practitioner accounts used “quant” for mathematically trained finance specialists, while textbooks and career markets normalised the label. It is safer to describe this as a process of occupational consolidation than as an invention at one moment.

By the early twenty-first century, the umbrella had widened to include at least pricing quants, quantitative traders, systematic researchers, quantitative developers, desk strategists, risk quants, model validators, execution researchers, and quantitative portfolio managers. Later, data scientists and machine-learning researchers further blurred the boundary between “quant” and general computational research.

11 Volatility: from a parameter to a stochastic object

11.1 The empirical challenge to constant volatility

The constant- σ assumption in geometric Brownian motion is analytically productive but empirically fragile. Mandelbrot documented heavy-tailed price changes and challenged Gaussian descriptions of returns [276]. Black discussed volatility asymmetry and what became known as the leverage effect [54]. A broad empirical literature established stylised facts including heavy tails, volatility clustering, aggregational effects, and dependence in nonlinear functions of returns [87].

This produced a characteristic quant dynamic: a simple model creates a useful benchmark; systematic residual patterns then create a research programme. Model failure is not merely an embarrassment. It is often the source of the next generation of quantitative work.

11.2 ARCH and GARCH

Engle’s ARCH model made conditional variance depend on past shocks [126]. In a simple ARCH(q) specification,

$$r_t = \sigma_t \varepsilon_t, \quad \sigma_t^2 = \alpha_0 + \sum_{i=1}^q \alpha_i r_{t-i}^2, \quad (20)$$

with standardized innovations ε_t . Bollerslev’s GARCH extension added lagged conditional variances [61]; the ubiquitous GARCH(1,1) form is

$$\sigma_t^2 = \omega + \alpha r_{t-1}^2 + \beta \sigma_{t-1}^2. \quad (21)$$

Nelson’s EGARCH model allowed log variance dynamics and asymmetric responses [299]. The family remains important because it offers a computationally cheap way to forecast conditional scale, often competitive with much more elaborate models [173].

11.3 Stochastic volatility

A distinct family treats volatility itself as latent and stochastic. Clark’s subordinated-process model is an important precursor [86]; Taylor developed discrete-time stochastic-volatility models [360–362]. In a canonical state-space form,

$$r_t = \exp(h_t/2) \varepsilon_t, \quad (22)$$

$$h_{t+1} = \mu + \phi(h_t - \mu) + \eta_t. \quad (23)$$

Inference becomes a latent-state problem. That in turn creates demand for Kalman-style approximations, simulation, Markov chain Monte Carlo, and particle filtering [157, 216, 232, 318, 338].

Continuous-time stochastic volatility supplied another route. Heston’s 1993 model,

$$dS_t = \mu S_t dt + \sqrt{v_t} S_t dW_t^{(S)}, \quad (24)$$

$$dv_t = \kappa(\theta - v_t) dt + \xi \sqrt{v_t} dW_t^{(v)}, \quad (25)$$

$$d\langle W^{(S)}, W^{(v)} \rangle_t = \rho dt, \quad (26)$$

combined mean-reverting variance with a correlation parameter capable of producing skew [191].

11.4 Smiles, local volatility, jumps, and roughness

The volatility smile observed after the 1987 crash made it untenable to treat a single constant implied volatility as a literal description of the world. Local-volatility and implied-tree approaches fit a cross-section of option prices by allowing volatility to depend on state and time [108, 116, 335]. Jump-diffusion models introduced discontinuous price moves [285], while stochastic-volatility and jump models offered competing explanations of smiles and dynamics [34].

Later work argued that volatility itself exhibits rough sample-path behaviour. Rough-volatility models use fractional or Volterra-type dynamics with a low Hurst parameter and have generated a substantial modern literature [42, 144]. The historical arc is instructive: volatility began as one coefficient in a diffusion, became a time series to forecast, then a latent stochastic state, a surface indexed by strike and maturity, and eventually an object with its own fine-scale regularity.

12 Market microstructure: the mathematics of how prices are made

Asset-pricing theory often begins with a price process as given. Market microstructure asks how that price is produced by trading rules, information, inventory, and strategic interaction. The distinction became increasingly important as markets automated.

12.1 Spreads, information, and inventory

Early microstructure models decomposed the bid–ask spread into economically meaningful components. The dealer may require compensation for inventory risk, order-processing costs, or adverse selection. Roll derived an implicit spread estimator from return autocovariances [332]. Glosten and Milgrom modelled a specialist facing traders with heterogeneous information [154]; Kyle modelled strategic informed trading and market depth in a continuous auction [244]. Grossman and Miller analysed liquidity as the capacity of intermediaries to absorb temporary order imbalances [163].

These models changed the quant’s object of study. The problem was no longer only “what is the fair value?” but “what will happen when I try to trade?” A signal with positive forecast return can be worthless after spread, impact, delay, and information leakage. This is the mathematical reason execution became a first-class quantitative discipline.

12.2 The limit order book as state space

In an electronic limit order book, standing orders form a high-dimensional state. Prices, queues, cancellations, order arrivals, hidden liquidity, and venue rules all affect fill probabilities and short-term price dynamics. Standard references systematise this institutional and empirical terrain [1, 174, 183, 307].

The move from daily prices to event-level order data brought finance close to fields such as point processes, stochastic control, queueing theory, online learning, and distributed systems. It also moved model risk closer to software risk: a conceptually correct strategy implemented with stale state, incorrect sequence handling, or poor clock synchronisation can lose money for reasons that never appear in the mathematical paper.

13 Electronic markets and the compression of time

13.1 From quotation screens to automated matching

Electronic trading did not arrive in one step. Early systems first automated quotation dissemination or institutional communication; later systems automated order routing, matching, and execution. NASDAQ, Instinet, electronic communication networks, exchange automation, and futures platforms each contributed to a transition in which market state became natively digital rather than a digital representation of activity that fundamentally occurred on a floor.

The digital transition changed the economics of research. Once both market data and order submission are machine interfaces, the loop

observe $\rightarrow$ estimate $\rightarrow$ decide $\rightarrow$ submit $\rightarrow$ observe

can be automated. The smallest economically meaningful time interval then contracts from minutes to seconds, milliseconds, microseconds, and in some settings below. The same contraction amplifies the cost of bad engineering.

13.2 Regulation and fragmentation

Electronic markets also became more fragmented. Competing venues, alternative trading systems, dark pools, internalisation, and smart order routing created a quantitative routing problem. In the United States, market-structure rules and Regulation NMS shaped the architecture of equity trading; in Europe, MiFID and later MiFID II/MiFIR promoted competition and imposed extensive transparency and reporting obligations [128–130, 375].

Regulation therefore changed the demand for quants in two ways. First, it altered trading opportunities and execution mechanics. Second, it created measurement obligations: transaction-cost analysis, best-execution evidence, market-abuse surveillance, capital models, and reporting systems. In modern finance, compliance itself can be a computational problem.

14 Algorithmic execution and high-frequency trading

14.1 Optimal execution

A large order creates a control problem. Trading quickly reduces exposure to future price risk but increases impact; trading slowly reduces instantaneous impact but leaves the investor exposed and may leak information. Bertsimas and Lo formulated execution as a dynamic optimisation problem [51]. Almgren and Chriss made the risk–impact trade-off canonical [7]. In schematic form, an execution objective can be written

$$\min_v \mathbb{E}[C(v)] + \lambda \text{Var}[C(v)], \quad (27)$$

where v is a trading schedule, C is implementation cost, and λ controls risk aversion. Later work modelled transient impact and order-book resilience [306].

The practical descendants are execution algorithms such as VWAP, TWAP, participation strategies, implementation-shortfall algorithms, and adaptive routers. Their mathematics interacts strongly with market microstructure because “impact” is not a universal physical constant. It depends on venue, participation, information content, order type, queue position, and regime.

14.2 Automated market making

Electronic market makers continuously solve a different control problem: where to quote, how much inventory to tolerate, when to cancel, and how to balance spread capture against adverse selection. Avellaneda and Stoikov provided an influential stochastic-control formulation [28]. Modern implementations enrich the state with order-book imbalance, fill probabilities, toxicity estimates, correlated instruments, fees, rebates, and latency.

14.3 The speed race

As electronic competition intensified, latency became an endogenous strategic variable. Empirical research documents the role of low-latency activity in modern markets [69, 184, 282]. Market-design research asks whether continuous-time priority creates socially wasteful races for tiny speed advantages and studies alternatives such as frequent batch auctions [72, 309]. MacKenzie’s sociological history shows how fibre routes, microwave towers, co-location, and hardware became part of market structure itself [270].

The 6 May 2010 Flash Crash became a defining episode in the debate over automated markets. Regulatory and academic analyses showed the interaction of a large automated sell programme, liquidity withdrawal, and rapid trading in a stressed electronic market [234, 376]. It is too simple to treat the event as “algorithms went wrong.” The deeper lesson is systemic: individually rational automated policies can interact nonlinearly when many agents respond to the same deteriorating state.

15 Systematic investing and statistical arbitrage

15.1 From factors to relative value

Systematic investing generalises the Markowitz insight: investment can be represented as an algorithm mapping data into positions. The map may be economically interpretable—value, momentum, carry, quality, term structure—or statistically learned. Relative-value strategies exploit relationships among instruments rather than the outright direction of a single asset. Pairs trading is a classic example and has been studied empirically over long samples [142]. Principal-component and factor residual methods provide more scalable constructions [26].

A stylised statistical-arbitrage pipeline is

raw data $\rightarrow$ clean data $\rightarrow$ features $\rightarrow$ forecast $\rightarrow$ portfolio $\rightarrow$ orders $\rightarrow$ execution $\rightarrow$ attribution.

Every arrow is a potential source of overfitting or implementation failure. The research problem is therefore not “find a predictive regression” but build a chain whose live distribution remains close enough to the backtest distribution after costs, constraints, and feedback.

15.2 The limits of arbitrage

Arbitrage in textbook models is often riskless; practical relative-value trading rarely is. Capital is finite, financing terms change, counterparties impose limits, and convergence horizons are uncertain. Shleifer and Vishny formalised limits to arbitrage arising from delegated capital and performance-sensitive investors [345]. The LTCM crisis supplied a vivid institutional example of how highly

sophisticated relative-value positions can become vulnerable to crowdedness, leverage, liquidity shocks, and adverse funding dynamics [115, 267].

This distinction is crucial to professional quantitative practice. A model can correctly identify a long-run relation and still produce a disastrous trade if it ignores path dependence, margin, liquidity, or capacity. Risk is not merely the variance of terminal payoff; it is the set of states through which the strategy must survive.

16 Machine learning and the expansion of the model class

16.1 Why machine learning fit the quant workflow

Machine learning entered finance not because finance suddenly discovered prediction, but because data, computing, and model classes expanded simultaneously. Tree ensembles, regularisation, boosting, neural networks, and representation learning made it possible to model nonlinear interactions without specifying each interaction *ex ante*. Standard statistical-learning references provided the methodological foundation [66, 156, 185].

For asset pricing, machine-learning methods have been used to model complex relations between characteristics and expected returns [166]. Recurrent networks and LSTMs have been applied to market prediction and volatility forecasting [138, 145, 196, 231, 260]. Deep learning has also been used to identify common structure in order-book dynamics across instruments [348].

16.2 The finance-specific difficulty: nonstationarity

The central obstacle is not lack of expressive power. It is that financial data are generated by adaptive agents under changing rules, technologies, and incentives. Predictors can decay because they are discovered, crowded, regulated, or arbitrated away. Structural breaks in volatility are one example [259]. More generally, the distribution of (X, Y) is not guaranteed to remain fixed after deployment.

This makes finance a hostile environment for naive machine learning. Cross-validation must respect time; leakage is pervasive; multiple testing can create compelling false discoveries; costs are state-dependent; labels are noisy; and strategy deployment can change the market. The modern quant increasingly resembles an ML engineer, but a finance-aware quant must remain sceptical of static predictive benchmarks.

16.3 Alternative data and the widening boundary of research

As conventional price and accounting data became widely available, firms invested in alternative data: text, web activity, geospatial observations, transaction proxies, shipping, supply chains, and other nontraditional sources. This widened the profession toward data engineering, natural-language processing, entity resolution, privacy governance, and causal inference. It also increased the value of infrastructure. The bottleneck often moved from model invention to data rights, cleaning, provenance, timestamp integrity, and reproducibility.

17 Crises, risk measurement, and the regulatory quant

17.1 Crises as regime changes for quantitative work

Financial crises repeatedly redirect quantitative labour. A crisis changes correlations, liquidity, volatility, defaults, policy, market rules, and institutional priorities simultaneously. Historical surveys emphasise the recurrence of leverage, maturity mismatch, asset-price booms, runs, and fire sales across very different monetary systems [233, 326, 328]. The implication is not that every crisis is the same, but that quantitative models calibrated to tranquil data can fail when the financial mechanism itself changes.

The 1929 crash and Great Depression were followed in the United States by major institutional reforms including the Banking Act of 1933 [373]. The collapse of Bretton Woods transformed exchange-rate risk and helped enlarge markets in currency derivatives; the history of the system and its demise is treated in Bordo and Eichengreen [62], Steil [354]. The 1987 crash intensified attention to portfolio insurance, liquidity, and volatility smiles. The 1992 sterling crisis demonstrated the interaction between macroeconomic policy commitments and speculative pressure. LTCM in 1998 exposed the interaction of leverage, liquidity, and model-based relative value.

17.2 VaR and coherent risk measures

The growth of large trading books created demand for an aggregate risk number. Value-at-risk became the dominant language of market-risk reporting in the 1990s, supported by systems such as RiskMetrics [222, 224]. For loss L , a confidence level α , and distribution function F_L , one definition is

$$\text{VaR}_\alpha(L) = \inf\{\ell : F_L(\ell) \geq \alpha\}. \quad (28)$$

VaR is intuitive but not generally subadditive. The theory of coherent risk measures formalised desirable axioms such as monotonicity, translation invariance, positive homogeneity, and subadditivity [24]. Expected shortfall,

$$\text{ES}_\alpha(L) = \mathbb{E}[L \mid L \geq \text{VaR}_\alpha(L)] \quad (29)$$

under suitable regularity, focuses on tail severity and became central to modern market-risk regulation; spectral risk measures generalise the idea [2].

17.3 Basel and the institutionalisation of quantitative risk

The Basel Accords made quantitative risk measurement part of the regulatory architecture of international banking. Basel I established common minimum capital concepts; Basel II expanded risk sensitivity and formalised a three-pillar framework; post-crisis Basel III reforms increased the emphasis on capital quality, leverage, liquidity, counterparty risk, and systemic resilience [37–39]. The Fundamental Review of the Trading Book revised market-risk rules, including the trading/banking-book boundary, liquidity horizons, model approval, and the use of expected shortfall in the internal-models framework [40].

These rules created a large class of “regulatory quants.” Their objective differs from a front-office pricing quant’s. They must map portfolios to regulatory categories, calculate stressed and capital measures, document assumptions, pass backtests and attribution tests, and maintain governance acceptable to supervisors. The model is simultaneously a statistical object, a legal object, and an operational system.

17.4 The global financial crisis and model critique

The 2007–2008 crisis exposed failures in mortgage underwriting, securitisation incentives, wholesale funding, ratings, leverage, liquidity management, and regulatory architecture. It cannot be reduced to a failure of one equation. Brunnermeier and Gorton provide complementary accounts of the liquidity and shadow-banking mechanisms [71, 158]; the U.S. Financial Crisis Inquiry Commission supplies the official investigative record [137]. Post-crisis U.S. legislation included the Dodd–Frank Act [374].

For the quant profession, the crisis changed what counted as a complete valuation. A derivatives price could no longer be treated as independent of counterparty default, collateral, funding, capital, and margin. This helped produce the family of valuation adjustments collectively known as xVA.

18 xVA, funding, and the post-crisis enlargement of valuation

Suppose a textbook model produces a “clean” value V^{clean} . In a bilateral, collateralised, regulated world, a desk may instead track a decomposition schematically written

$$V^{\text{all-in}} = V^{\text{clean}} - \text{CVA} + \text{DVA} - \text{FVA} - \text{KVA} - \text{MVA} + \dots, \quad (30)$$

where the signs and exact definitions depend on conventions and modelling choices. CVA reflects counterparty credit, DVA own credit, FVA funding, KVA capital, and MVA initial-margin funding. The mathematical literature developed PDE, BSDE, simulation, and exposure-based approaches to these effects [74, 93, 161].

The historical importance of xVA is organisational as much as mathematical. Valuation became a negotiation among trading, treasury, credit, collateral, capital, and regulation. A “price” was increasingly understood as institution-dependent because two firms with different funding, netting, collateral, and capital positions could rationally attach different all-in economics to the same contractual payoff. This is a profound departure from the clean representative-agent intuition often associated with textbook derivatives pricing.

19 Computing as a constitutive part of quantitative finance

19.1 From hand calculation to numerical production

Quantitative finance is impossible to separate from computing. Even simple models become industrially useful only when they can be evaluated across large books, scenarios, calibrations, and markets. Early electronic computers supported portfolio optimisation, actuarial work, and simulation. Mainframes gave way to departmental systems and workstations; spreadsheets democratised model construction; compiled languages supported pricing and risk libraries; specialised time-series databases made tick data tractable; clusters, GPUs, and cloud systems expanded the scale of simulation and machine learning.

Each computing regime favoured different mathematical compromises. A closed-form approximation is valuable when CPU is scarce or latency is strict. A recombining lattice is attractive when memory and interpretability matter. Monte Carlo becomes powerful when parallel compute is abundant. Adjoint methods become valuable when a desk needs sensitivities to thousands of inputs. Machine learning becomes practical when training data, accelerators, and software ecosystems exist together.

19.2 High-performance and low-latency computing

High-performance computing and low-latency computing overlap but optimise different things. HPC seeks throughput: more paths, scenarios, instruments, or training examples per unit time. Low-latency systems seek response time and jitter: the tail of the time-to-decision distribution may matter more than average throughput. The engineering tools accordingly differ. Vectorisation, GPUs, and distributed batches dominate many risk and ML workloads; kernel bypass, lock-free structures, hardware timestamping, FPGAs, co-location, and careful network design matter in latency-sensitive trading.

This bifurcation created a distinct occupational branch: the quantitative developer. The role is not merely “programmer for quants.” In many firms, the quantitative developer owns the translation from mathematical object to production semantics: numerical stability, data structures, concurrency, testing, deployment, and performance. The boundary between research and engineering is especially thin in electronic trading.

19.3 Reproducibility and model governance

As systems grew, reproducibility became a risk-control problem. A backtest that cannot be reconstructed from versioned data, code, parameters, and environment is not merely inconvenient; it is difficult to validate. Modern research platforms therefore increasingly borrow practices from software engineering and scientific computing: version control, deterministic pipelines where feasible, containerised environments, automated tests, data lineage, experiment tracking, and independent validation.

This professionalises scepticism. The mature quant organisation assumes that code will contain bugs, data will be revised, dependencies will change, and attractive results may be statistical accidents. Governance is not antithetical to quantitative innovation; it is the machinery by which an institution decides which innovations can safely become positions.

20 Blockchain, decentralised finance, and programmable market infrastructure

Bitcoin’s 2008 design combined a distributed ledger, cryptographic proof, and an incentive mechanism to create a digitally scarce bearer-like asset without a central operator [295]. Ethereum generalised the idea toward a programmable state machine supporting smart contracts [386]. Decentralised finance then recreated many familiar financial functions—exchange, collateralised lending, leverage, derivatives, and market making—inside programmable protocols [181, 380].

For quants, this environment is historically interesting because market design, settlement, and software are unusually inseparable. Automated market makers replace or supplement limit order books with deterministic pricing curves; transaction ordering and block production create new forms of execution risk; smart-contract state is directly observable yet can be adversarially manipulated; and protocol rules can change through governance. The quant becomes partly a mechanism designer and security engineer.

A substantial strand of research associated with William J. Knottenbelt and his collaborators at Imperial College London illustrates this convergence particularly well. Early work by McGinn et al. [279] treated the Bitcoin ledger as a rich empirical object, developing high-resolution visualisations of transaction, mempool, and network activity and showing how the transparency of a public ledger

permits behavioural and structural patterns to be studied directly. The same programme addressed the security of programmable financial infrastructure: Harz and Knottenbelt [182] surveyed smart-contract languages and formal verification techniques, while Stewart et al. [357] analysed the threat that quantum computation poses to Bitcoin’s elliptic-curve cryptography and proposed a migration mechanism toward quantum-resistant outputs.

Interoperability created another layer of financial engineering. Zamyatin et al. [393] introduced XCLAIM, a framework for creating trust-minimised cryptocurrency-backed assets across otherwise independent blockchains. This line of work was subsequently placed in a broader framework by Zamyatin et al. [394], who systematised mechanisms for communication across distributed ledgers. Such work anticipates a world in which the relevant “market infrastructure” is not one exchange, clearing house, or ledger, but a collection of interacting state machines whose economic guarantees depend on cryptography, collateral, consensus, and cross-chain communication.

Knottenbelt and collaborators also studied the explicitly financial mechanisms built on top of this infrastructure. Gudgeon et al. [167] analysed DeFi protocols for loanable funds such as Compound, Aave, and dYdX, examining their endogenous interest-rate mechanisms, liquidity, cross-protocol dependence, and market efficiency. Werner et al. [380] then provided a systematic account of DeFi primitives, protocol types, and security, importantly distinguishing technical security from *economic security*. The distinction is crucial for quantitative finance: a smart contract may execute exactly as programmed while the economic mechanism implemented by that contract nevertheless admits manipulation, destabilising feedback, adverse selection, liquidation cascades, or strategic exploitation.

The underlying blockchain protocol itself can exhibit analogous feedback effects. Ilie et al. [207] showed how a poorly designed proof-of-work difficulty-adjustment rule can generate oscillatory transaction throughput and strategic coin-hopping, and derived an alternative adjustment mechanism based on a negative exponential filter. This is a striking example of a familiar control-theoretic problem reappearing inside monetary infrastructure: protocol parameters alter incentives, incentives alter participant behaviour, and participant behaviour feeds back into the dynamics the protocol was intended to stabilise.

The same programme has increasingly connected blockchain engineering to formal verification, regulation, and quantitative risk. Xiong et al. [389] developed a taxonomy and empirical analysis of market misconduct in DeFi, emphasising that programmability can create novel forms of manipulation rather than merely automate familiar ones. More recent work by Jones et al. [221] uses copula-based methods to study the transmission of volatility and activity from stablecoins to broader cryptocurrency markets, treating stablecoin liquidity as a measurable component of systemic crypto-market risk. At the protocol level, Jones and Knottenbelt [220] applies automated theorem proving to formal verification of blockchain consensus, extending the earlier concern with smart-contract correctness toward the correctness of the consensus machinery itself.

At the same time, crypto markets reiterate older themes. Leverage produces liquidations; maturity and liquidity mismatches produce runs; intermediaries fail; informational advantages are monetised; and regulation follows after institutional scale. New infrastructure does not abolish the recurring economic logic documented in histories of financial crises. What changes is that part of the market’s institutional structure has become executable code, making software semantics, consensus design, cryptographic security, mechanism design, and quantitative risk management components of a single financial system.

21 Large language models and the next boundary of the quant role

Large language models (LLMs) add a new kind of interface between unstructured information and quantitative workflow. Finance-specific models such as BloombergGPT and open projects such as FinGPT demonstrate the adaptation of large language models to financial corpora and tasks [387, 390]. Early empirical research has explored whether general-purpose LLMs can extract return-relevant sentiment from news [266]. Broader work by the Bank for International Settlements surveys potential effects of AI across financial intermediation [5].

The strongest near-term impact may be less romantic than autonomous “AI traders.” LLMs can accelerate tasks at the edges of formal models: reading documentation, generating and reviewing code, mapping contracts to data, extracting entities from text, writing tests, comparing model assumptions, producing research summaries, and translating natural-language questions into analytical workflows. In that sense they reduce the cost of moving between representations.

But finance exposes LLM weaknesses sharply. A plausible but false numerical claim is dangerous; stale regulatory knowledge can be consequential; generated code can compile while implementing the wrong convention; a summary can erase a tail-risk assumption. The appropriate architecture therefore treats language models as fallible components surrounded by deterministic tools, retrieval, tests, permissions, provenance, and human accountability. The quant profession may become more productive through AI, but its comparative advantage shifts toward problem formulation, verification, and systems judgement.

22 People, schools, and lineages in modern quantitative finance

A history of quantitative finance written only as a succession of equations misses an important fact: the field has been built by overlapping communities of mathematicians, economists, physicists, computer scientists, traders, risk managers, software engineers, educators, and institution-builders. The same model may have an academic origin, a bank implementation, a software embodiment, a conference life, and a later reinterpretation in a different subfield. Conversely, individuals who never attached their names to a canonical model may have had enormous influence through books, software, research groups, trading firms, conferences, or the professional labour market.

The following sections therefore adopt a prosopographical approach. The objective is not to rank individuals, nor to imply that the people named below exhaust the field, but to show how different technical and institutional lineages intersect. Where a priority claim is ambiguous, the emphasis is on a person’s demonstrable contribution rather than on assertions of being “first.” For several practitioners whose work was proprietary, the public record is necessarily thinner than for academics; such cases are explicitly described as practitioner history rather than reconstructed as a fictitious publication record.

22.1 No-arbitrage, martingales, stochastic calculus, and control

The modern mathematical-finance language of contingent claims was placed on a rigorous footing by the martingale approach. J. Michael Harrison and David Kreps connected arbitrage, equilibrium, and martingale measures, while Harrison and Stanley Pliska developed the stochastic-integration formulation of continuous-time trading and the idea that suitably discounted security prices should become martingales under an equivalent probability measure [175, 176, 178]. In historical terms,

Harrison–Pliska is a hinge: the Black–Scholes–Merton insight becomes not merely an option formula but an abstract theory of attainable claims, complete markets, and state-price systems.

Steven Shreve did much to turn this framework into a teachable canon. His two-volume treatment connects Brownian motion, martingales, stochastic calculus, risk-neutral valuation, change of measure, and term-structure modelling in a form that trained several generations of quants [346, 347]. Bernt Øksendal’s text played a similarly important role in diffusing stochastic differential equations beyond specialist probability departments [308]. Albert Shiryaev’s work on probability, optimal stopping, sequential analysis, and mathematical finance supplied a deeper probabilistic lineage, particularly for problems in which stopping rules and incomplete information matter [343, 344].

Nicole El Karoui’s influence is difficult to reduce to a single model. Her school helped establish backward stochastic differential equations, stochastic control, hedging in incomplete markets, and mathematically rigorous derivatives modelling as central parts of the French mathematical-finance tradition [121]. Huy  n Pham systematised stochastic control and viscosity-solution methods for finance and developed influential work on optimal investment, hedging, and later machine-learning approaches to control [316]. Jak  sa Cvitan  i   connected martingale methods, constraints, contracting, and portfolio choice; his books and papers with collaborators helped make constrained optimisation and incomplete markets standard quant topics [94, 95]. Marek Musiela, both through the Musiela parametrisation of the HJM framework and through his work with Marek Rutkowski, made the functional-analytic and term-structure formulation of arbitrage pricing more usable [293, 294].

Mark H. A. Davis represents another crucial line: stochastic control and piecewise-deterministic Markov processes brought into finance from control theory rather than from equilibrium economics. His work ranges from general stochastic-control foundations to option pricing under transaction costs, credit risk, and mathematical models whose state evolves deterministically between random events [99–101]. This control-theoretic viewpoint became increasingly important as quantitative finance shifted from pricing a terminal payoff to choosing a sequence of actions under costs, constraints, and partial information.

Panos Parpas and Ber    R  stem belong naturally in this optimisation lineage. R  stem’s work on robust decision making, min–max optimisation, and financial risk anticipated the increasing importance of parameter uncertainty and adversarial formulations, while Parpas has developed large-scale stochastic, robust, and decomposition methods relevant to portfolio and risk problems [310, 336]. Their work is a reminder that a substantial portion of “quant finance” is really operations research under financial semantics.

Dimitri Bertsekas and Warren Powell are not primarily finance researchers, yet their impact on the modern quant toolkit is profound. Bertsekas’ dynamic programming and optimal-control programme gives the classical conceptual backbone behind value functions, Bellman recursion, approximate dynamic programming, and reinforcement learning [49, 50]. Powell’s unified treatment of sequential decisions under uncertainty similarly connects stochastic optimisation, approximate dynamic programming, simulation, and reinforcement learning [321, 322]. Their inclusion is important precisely because modern financial RL inherits far more from control and operations research than the label “AI” sometimes suggests.

22.2 Term structures, rates, and the architecture of fixed-income modelling

Oldřich Va    ek’s 1977 short-rate model is one of the earliest tractable equilibrium-style term-structure models in continuous time and remains a canonical example of mean reversion in rates [378]. John Hull and Alan White extended the short-rate approach into a flexible family capable of fitting an initial yield curve, making the Hull–White model a mainstay of rates desks and teaching

[201]. Hull’s broader impact is even larger: *Options, Futures, and Other Derivatives* became one of the profession’s most widely used bridges between theory and desk practice, and his work with White ranges from volatility and interest rates to credit and valuation adjustments [202, 204].

The Black–Derman–Toy model, associated with Fischer Black, Emanuel Derman, and William Toy, represents the tree-based, arbitrage-free short-rate tradition that accompanied the growth of callable and mortgage-related interest-rate products [60]. Piotr Karasiński’s work with Black produced the Black–Karasiński lognormal short-rate model, another influential desk model designed to combine mean reversion, positivity, and calibration convenience [56].

The Heath–Jarrow–Morton framework shifted the state variable from a short rate to the full forward curve and showed how absence of arbitrage constrains drift once volatility is specified [189]. The Musiela parametrisation then recasts maturity as time-to-maturity, turning the forward curve into an infinite-dimensional state evolving in function space [293]. This line ultimately fed both rigorous stochastic analysis and industrial yield-curve engines.

Damiano Brigo and Fabio Mercurio helped consolidate the post-HJM world for practitioners, particularly through affine models, LIBOR-market models, calibration, and the systematic comparison of interest-rate model families [67]. Vladimir Piterbarg’s work on funding, collateral, rates, and multi-curve pricing became central after the financial crisis, when the single-curve idealisation broke down and valuation had to account explicitly for collateral and funding conventions [14, 317]. Alexander Lipton has repeatedly occupied the boundary between theory and implementation, with contributions across rates, FX, credit, numerical methods, and later machine learning and quantum computing; his texts and industrial papers exemplify the broad “quant architect” role rather than a single-model identity [257, 258].

John Hull, Alan White, Brigo, Mercurio, Musiela, Karasiński, Piterbarg, and Lipton therefore illustrate a general pattern in quant history: successful models are not only stochastic processes. They are calibration conventions, numerical algorithms, data requirements, hedge representations, software interfaces, and organisational practices.

22.3 Volatility, smiles, jumps, roughness, and the geometry of option markets

Steven Heston’s 1993 stochastic-volatility model is one of the key bridges between the first generation of constant-volatility derivatives theory and the modern implied-volatility surface. By coupling the asset price to a square-root variance process while retaining a semi-closed-form characteristic-function solution, the model made stochastic volatility computationally practical for option pricing [191]. Its descendants remain ubiquitous, and “Heston” has become almost a generic adjective in calibration and hybrid-model discussions.

Lorenzo Bergomi pushed the state variable one step further: rather than model only spot variance, forward-variance models treat the variance term structure itself as the natural object. The Bergomi approach, later synthesised in his monograph, provides a conceptual route from stochastic volatility to stochastic-local and rough-volatility modelling and focuses attention on the dynamics of the entire volatility surface [45, 46]. Jim Gatheral’s work is closely connected to this development: his treatment of volatility surfaces, no-arbitrage parameterisations, and later rough volatility helped turn smile modelling into a coherent field rather than a collection of desk recipes [143, 144].

The SABR model of Patrick Hagan, Deep Kumar, Andrew Lesniewski, and Diana Woodward is another decisive practitioner model. Its asymptotic implied-volatility formula gave rates and FX desks a parsimonious, interpretable way to fit smiles while retaining a direct connection to stochastic volatility and elasticity [169]. The “SABR crew” is historically important not only because the

model endured, but because it exemplifies a style of quant work in which asymptotic mathematics is chosen to match the calibration and quoting language of traders.

Dilip Madan’s work on variance-gamma and, with Carr and Chang, on finite-variation Lévy models, helped establish jump processes as practical alternatives to pure diffusion [272, 273]. Peter Carr contributed across option replication, static hedging, transforms, variance swaps, volatility derivatives, and jump-process modelling; the Carr–Madan Fourier-pricing method became especially influential in computational derivatives [79, 80]. His career illustrates how a quant can reshape both mathematical technique and desk conventions.

Rama Cont developed a body of work spanning stylised facts, Lévy processes, volatility, functional Itô calculus, market microstructure, systemic risk, and model uncertainty. His 2001 survey of empirical properties remains a standard point of departure for the argument that finance models should be disciplined by robust empirical regularities rather than only tractability [87, 88]. Mathieu Rosenbaum, Jean-Philippe Bouchaud, and collaborators connected high-frequency order flow, Hawkes processes, rough volatility, and market impact, helping blur the older boundary between “derivatives modelling” and “microstructure” [31, 63, 120].

Antoine (Jack) Jacquier’s work sits at several intersections: large deviations and asymptotic methods for implied volatility, stochastic and rough volatility, numerical and machine-learning approaches, and more recently quantum computing [140, 215]. Blanka Horvath has likewise worked across rough volatility, generative modelling, signatures, machine learning, and data-driven volatility, representing the younger generation for whom stochastic analysis and ML are not rival paradigms but parts of one modelling toolkit [197, 198]. Milena Vuletić’s VolGAN work, recognised by Risk.net’s 2025 Rising Star in Quant Finance award, is a particularly clear example of this convergence: a learning-based market generator is assessed by its ability to reproduce arbitrage-free and realistic volatility-surface dynamics rather than by predictive accuracy alone [89, 330].

Elisa Alòs and collaborators developed Malliavin-calculus methods for stochastic-volatility models, implied-volatility expansions, and sensitivity analysis. Malliavin calculus is especially useful when Greeks and conditional expectations must be represented in forms suitable for Monte Carlo computation; the Alòs line connects abstract stochastic analysis with practical asymptotics and simulation [9, 10]. Marc Yor’s vast work on exponential functionals of Brownian motion, Bessel processes, local times, and related stochastic identities provided many of the probabilistic tools later exploited in Asian-option and stochastic-volatility calculations [391, 392].

Nassim Nicholas Taleb and Espen Gaarder Haug belong to a different but influential volatility tradition: practitioner critique. Taleb’s writing and trading-oriented work emphasise convexity, fat tails, model error, and the danger of treating calibrated distributions as physical laws [358, 359]. Haug has repeatedly excavated pre-Black–Scholes option-pricing history and challenged overly linear narratives of derivatives theory, while also producing practitioner formula collections that became desk references [186, 187]. Their role is partly technical and partly historiographical: they remind the profession that a model’s social authority is itself a financial variable.

Didier Sornette and Jean-Philippe Bouchaud represent the statistical-physics or econophysics current. Sornette’s work on crashes, critical phenomena, herding, and log-periodic structures sought signatures of collective instability [352]; Bouchaud’s work on anomalous diffusion, order flow, impact, and agent-based market phenomena became deeply influential in modern microstructure [64]. Raphael Douady similarly connected dynamical-systems thinking, crisis indicators, and quantitative risk, adding another route by which nonlinear science entered finance.

22.4 Credit, counterparty risk, xVA, and the post-crisis quant

Credit derivatives and counterparty risk created a distinct quant ecology. Robert Jarrow’s work with Stuart Turnbull on reduced-form credit-risk modelling provided a tractable intensity-based alternative to structural default models [219]. Darrell Duffie’s work spans asset pricing, term structures, credit risk, OTC markets, and market design; with Singleton he gave a powerful reduced-form framework for defaultable claims [112, 113]. Hélyette Geman contributed across commodities, interest rates, numeraire methods, and insurance-linked products, and her work helped normalise measure-change techniques in practitioner finance [147, 148].

Damiano Brigo’s later work with Morini and Pallavicini helped codify counterparty-credit and funding adjustments as a coherent subject, while Jon Gregory’s practitioner books made CVA and xVA accessible to the wider industry [68, 160]. Harvey Stein’s work at the practitioner–academic boundary is especially important here. Across model risk, counterparty and credit analytics, numerical methods, and professional organisations, Stein exemplifies the senior bank quant whose influence is distributed through methodology, implementation, seminars, and standards rather than through a single eponymous model. His role in the IAQF/Thalesians New York seminar series also illustrates how practitioner networks helped circulate technical ideas outside formal journals [355, 365].

Youssef Elouerkhaoui developed dynamic and correlation models for credit, including work on pricing and hedging in dynamic credit models, and later led broad sell-side quantitative organisations spanning credit, CVA, funding, capital, and markets modelling [124, 125]. Alexander Sokol and Alexander Denev pushed the practical treatment of counterparty and credit risk through modelling, stress testing, and data-driven risk methods [106, 351]. Alexandre Antonov, Serguei Issakov, and related practitioner researchers contributed to industrial derivatives modelling, especially in rates, volatility, and numerical implementation; much of this tradition appeared in bank notes, conference presentations, and specialist books rather than only in journals.

Luca Capriotti’s work on adjoint algorithmic differentiation and Monte Carlo Greeks became especially important once xVA and large risk systems made the cost of sensitivities a first-order engineering problem [78]. Paul Glasserman’s Monte Carlo work provides one of the definitive mathematical and computational foundations for simulation-based derivatives and risk [153]. Peter Jaeckel’s practitioner work on Monte Carlo, calibration, implied volatility, and numerical robustness similarly shaped the implementation culture of derivatives desks [217, 218]. Daniel Duffy’s books and training on finite differences, C++, and computational finance helped translate numerical PDE methods into production-oriented practice [114].

Andrew Kalotay’s work on fixed-income analytics, term-structure models, mortgage and callable-bond valuation represents another path by which industrial fixed-income problems shaped quantitative technique. Joe (Jon) Gregory’s xVA work, Harvey Stein’s risk infrastructure, and the large bank quant teams led by figures such as Elouerkhaoui show a post-crisis transformation: the quant’s unit of analysis became not simply a trade but a portfolio, netting set, counterparty, funding curve, collateral agreement, legal entity, and capital regime.

22.5 Asset pricing, portfolio construction, systematic trading, and the quant fund

Kenneth Arrow’s state-contingent securities and general-equilibrium work supplied a conceptual foundation for modern state pricing long before “quant” became a job title [19]. Stephen Ross’s arbitrage pricing theory and work on options and agency broadened the equilibrium and factor-

model foundations of investment finance [333]. Jonathan Ingersoll’s textbook and research helped systematise modern financial theory, including portfolio choice, asset pricing, and contingent claims [211]. These economists belong in a quant history because portfolio models and derivative valuation inherited the language of state prices, spanning, factor exposures, and no-arbitrage from this broader theoretical tradition.

Victor Niederhoffer represents the empirical trader-scholar tradition: statistical regularities, contrarian hypotheses, and market experimentation were treated as objects for systematic investigation, though his later fund experience also became a case study in leverage and tail exposure [302]. Andrew Lo’s work on market efficiency, adaptive markets, hedge funds, and empirical finance brought econometrics and institutional realism together; his Adaptive Markets Hypothesis is explicitly evolutionary rather than equilibrium-static [261, 262]. Marcos López de Prado has become influential through work on financial machine learning, backtest overfitting, feature importance, portfolio methods, and the methodological discipline required when data are scarce relative to researcher degrees of freedom [33, 265].

Attilio Meucci developed a coherent practitioner framework for risk and portfolio management that links estimation, views, simulation, and allocation; his entropy-pooling approach provides a particularly flexible way to update distributions under subjective or stress views [287, 288]. Frank Fabozzi has had a vast editorial and pedagogical impact across fixed income, portfolio management, and quantitative investment, often acting as a bridge between specialised research and professional education [131]. Joseph Simonian represents the more recent applied-investment-science tradition in which machine learning, portfolio construction, and institutional asset management are treated as a single research programme.

Jim Simons and D. E. Shaw are central not because they published a canonical pricing equation, but because they created durable organisations in which mathematical, statistical, and computational research itself became the investment process. Renaissance Technologies and D. E. Shaw institutionalised the idea that data, research infrastructure, software, and portfolio construction could be integrated at firm scale. Alexander Gerko later represents a similar institution-building role in electronic and systematic trading: the salient historical innovation is not a paper but the creation of research-intensive trading organisations in which quants are not a support function to traders but the core production technology.

Igor Tulchinsky’s WorldQuant made the “alpha factory” idea explicit: large numbers of researchers generate, test, combine, and continuously monitor predictive signals. The model is organisationally different from the classical derivatives desk and helps explain why modern buy-side quants increasingly resemble experimental scientists embedded in a software platform [372]. Svetlana Borovkova’s work across commodities, energy, volatility, and later climate-risk analytics illustrates the broadening of quant research toward non-equity data and sustainability-related risk. Saeed Amen’s work on FX, alternative data, and systematic macro research, together with his role in community education, exemplifies the practitioner who combines research, data engineering, writing, and professional network building [12].

Miquel Noguer i Alonso has contributed to the financial-ML and education ecosystem, with work spanning AI, alternative data, and quantitative-investment curricula. Alec N. Schmidt has similarly contributed to the practitioner literature on machine learning and quantitative trading. Where the field’s earlier texts taught Itô calculus as the differentiating skill, these authors represent a period in which model validation, ML, data pipelines, and robust experimental design became equally central.

22.6 Market microstructure, impact, optimal execution, and electronic market making

Robert Almgren and Neil Chriss converted the practical problem of liquidating a large position into a transparent optimisation between expected impact cost and execution risk [8]. Their framework became foundational because it expresses the trader’s problem as control under endogenous cost rather than as frictionless portfolio selection. It also gave the profession a vocabulary—temporary impact, permanent impact, implementation shortfall, efficient execution frontier—that could be measured and debated.

Marco Avellaneda’s work spans derivatives, homogenisation, uncertain volatility, statistical arbitrage, and market making. The Avellaneda–Stoikov model made inventory-aware optimal quoting a standard entry point into quantitative market making [29]; his statistical-arbitrage work with Lee likewise helped formalise relative-value trading in a factor-residual framework [27]. Álvaro Cartea, Sebastian Jaimungal, and José Penalva subsequently developed a broad stochastic-control treatment of algorithmic and high-frequency trading, encompassing execution, market making, order types, signals, and limit-order-book state [81].

Olivier Guéant’s work on optimal market making and execution made explicit the connection between stochastic control, inventory risk, and limit-order placement [168]. Charles-Albert Lehalle connected optimal execution theory to empirical order-book behaviour, transaction-cost analysis, and actual trading algorithms; his work with Eyal Neuman on incorporating predictive signals into optimal trading is a good example of the post-Almgren evolution from deterministic schedules toward state-dependent control [249, 251]. Eyal Neuman’s broader work on crowding, transient impact, and interacting traders places execution in a game-theoretic setting rather than treating the market as a passive cost curve [301].

Jean-Philippe Bouchaud, Rama Cont, and Mathieu Rosenbaum supplied complementary empirical and stochastic-process foundations: long-memory order flow, propagator models, Hawkes processes, price-impact functions, and scaling limits. This literature changed the object being modelled from a smooth price diffusion to a coupled system of orders, queues, agents, and feedback. Charles-Albert Lehalle’s phrase “market microstructure in practice” captures the shift: the quant is now modelling the mechanism by which a transaction creates a price, not merely the resulting price series.

Johannes Muhle-Karbe has developed a rigorous modern theory of trading with frictions, liquidity, and price impact, including asymptotics and tractable formulations of nonlinear impact. David Itkin’s recent work on stochastic portfolio theory with price impact and nonlinear price-impact structures belongs to this same frontier [70, 212]. Risk.net selected Itkin as its 2026 Rising Star in Quant Finance for price-impact research, making the award a useful marker of how central market impact has become to current mathematical finance [331].

Petter Kolm and Gordon Ritter connect microstructure and portfolio problems to reinforcement learning. Their work on dynamic replication and hedging and their survey of RL in finance frame pricing, execution, allocation, and market making as sequential decision problems rather than isolated prediction tasks [237, 238]. This is conceptually close to the control tradition of Davis, Pham, Bertsekas, and Powell, but computationally oriented toward learning policies from data.

22.7 Machine learning, deep hedging, reinforcement learning, and data-driven finance

Matthew Dixon is among the researchers who helped consolidate machine learning as a serious quantitative-finance discipline rather than a collection of trading heuristics. His work spans

deep learning for market data, execution, forecasting, and the methodological relationship between statistical learning and financial decision problems; his textbook with Halperin and Bilokon explicitly places ML beside econometrics and stochastic control [110]. Igor Halperin’s QLBS programme is a particularly elegant bridge from Black–Scholes reasoning to reinforcement learning: option pricing and hedging are formulated as a Q-learning problem in which the price and hedge emerge from one sequential optimisation [170, 171].

Hans Buehler and collaborators’ “deep hedging” work made neural networks into policy approximators for hedging under market frictions and nonclassical risk objectives [73]. This matters historically because it shifts the role of a network from predicting a market variable to representing an action rule. Petter Kolm and Gordon Ritter similarly emphasise RL for hedging and portfolio decisions, while Ben Hambly, Renyuan Xu, and Huining Yang surveyed the rapidly expanding field of financial RL across execution, market making, pricing, portfolio choice, and robo-advice [172].

Eyal Neuman and Yufei Zhang have worked on learning and control in price-impact and propagator settings; Yufei Zhang’s broader work on continuous-time reinforcement learning connects stochastic control theory with policy-gradient and entropy-regularised methods. Ben Hambly provides another bridge between rigorous probability and ML/RL in finance. Abbas Edalat enters this story from computer science: his work on exact computation, domain theory, optimisation, and semantics represents the foundations-oriented side of computational mathematics, relevant to a field increasingly concerned with correctness as well as speed.

Blanka Horvath, Antoine Jacquier, Cristopher Salvi, and collaborators exemplify the newer synthesis of stochastic analysis, signatures, rough paths, generative modelling, and deep learning. Salvi’s Sig-SDE work with Imanol Pérez Arribas and Łukasz Szpruch treats path signatures as a bridge between stochastic-process models and generative learning [312]. Milena Vuletić’s volatility-generative work and Jochen Papenbrock’s applied ML research illustrate the same institutional shift: modelling teams increasingly combine classical quant validation standards with modern statistical learning.

Joerg Osterrieder has been influential in quantitative-finance education, applied research, and the professional diffusion of data science, fintech, and AI. Jochen Papenbrock has similarly operated across financial ML, applied research, and industry education. Thomas (Tom) Hadfield represents a particularly revealing career arc: an early research background in pure mathematics and noncommutative geometry, a long industrial quantitative-research career at G-Research, and a return to academia in quantitative finance with interests including the mathematics of LLMs and computer-assisted formalisation [208]. The trajectory is itself historical evidence of how the profession has widened: abstract mathematics, large-scale systematic research, language models, and formal proof can now lie on one quant career path.

22.8 Practitioner-authors, textbooks, and the transmission of tacit knowledge

Several individuals shaped the profession primarily by making specialised knowledge portable. John Hull’s textbooks created a common derivatives vocabulary across universities, banks, regulators, and professional qualifications. Paul Wilmott combined PDE-based derivatives modelling with an unusually broad publishing, training, magazine, and conference ecosystem; his books with Jeff Dewynne and Sam Howison helped establish the Oxford-style PDE treatment of derivatives among practitioners [384, 385]. Salih Neftci’s *Introduction to the Mathematics of Financial Derivatives* and *Principles of Financial Engineering* provided a complementary practitioner route into stochastic calculus and instrument construction [297, 298].

Emanuel Derman’s technical work on the Black–Derman–Toy model and local-volatility ideas

was accompanied by an influential account of the quant’s professional identity. *My Life as a Quant* helped make visible the migration of physicists into finance and the tensions between models as scientific representations and models as engineered tools [107, 109]. Paul Wilmott’s public persona played a different but comparable role: he created a culture in which model critique, numerical recipes, practitioner questions, and humour coexisted with serious mathematics.

Irene Aldridge helped codify high-frequency trading and electronic-market practice for a broad audience [6]. Yves Hilpisch did something similar for Python: his books and training helped normalise Python as a production-relevant language for finance rather than merely a scripting environment [192, 193]. Mark Joshi’s books on C++ design patterns and quantitative finance helped bridge mathematical modelling and production software engineering [223]. Marek Capiński’s texts helped make mathematical finance accessible to students crossing from pure mathematics into the subject.

Espen Haug’s formula compendia, Peter Jaeckel’s numerical texts, Daniel Duffy’s C++/PDE books, and Frank Fabozzi’s extensive publishing programme all belong to this knowledge-transmission infrastructure. A profession cannot scale by apprenticeship alone; these authors transformed tacit desk practices into portable curricula.

22.9 Named institutions and firm builders: Simons, Shaw, Gerko, and the research organisation

The emergence of the quant fund changed the organisational status of quantitative research. At a traditional sell-side derivatives desk, a quant could be an adviser to traders or a producer of pricing libraries. At Renaissance Technologies, D. E. Shaw, and later systematic firms, research and software became the business itself. Jim Simons’ institutional significance lies in creating a durable organisation in which mathematical research, data, computing, transaction-cost control, and portfolio aggregation were integrated. D. E. Shaw followed a similarly research-intensive path and demonstrated that computer science and distributed systems could be as important to systematic investing as classical financial economics.

Alexander Gerko and XTX Markets represent the later electronic-market-making version of the same institutional inversion: quantitative research, data, and technology are not “support” for a discretionary trading hierarchy but the core of the firm. The rise of such organisations contributed to the modern labour market in which a senior quant researcher may have more in common with a machine-learning scientist or systems engineer than with the spreadsheet-based “analyst” implied by the older job title.

22.10 Practitioner case studies: Bruno Iksil and Jérôme Kerviel

Bruno Iksil (the “London Whale”) and Jérôme Kerviel should not be placed in a list of model inventors. Their importance to a history of the quant profession is instead diagnostic. Both episodes exposed the gap that can open between formal risk systems and the actual socio-technical system of limits, marks, controls, incentives, data, management escalation, and trader behaviour.

The 2012 JPMorgan Chief Investment Office losses associated with Iksil demonstrated that a sophisticated institution can possess extensive VaR, stress, and risk infrastructure while still suffering from model changes, concentration, liquidity, valuation uncertainty, and governance failures. The lesson is not that quantitative risk is useless; it is that a risk number is an output of a model embedded in an organisation, and the organisation determines what positions enter, how parameters are estimated, which exceptions are escalated, and whether the metric constrains

behaviour [225, 377].

Kerviel’s losses at Société Générale similarly illuminate the operational boundary of quantitative trading. The central questions concerned unauthorised positions, fictitious hedges, control circumvention, and the relationship between front-office systems and supervisory processes. These events belong in quant history because they helped push model risk, independent validation, operational controls, and surveillance into the same conversation as pricing theory [350].

23 The computational substrate: numerical methods, languages, databases, and specialised hardware

Mathematical finance is often narrated as though the model exists first and is later “implemented.” In production, the distinction is artificial. What can be computed, calibrated, stored, queried, differentiated, and updated at the required latency determines which models are economically usable. The history of the quant profession is therefore inseparable from the history of numerical software, programming languages, data systems, and hardware.

23.1 Monte Carlo, PDEs, numerical analysis, and algorithmic differentiation

Paul Glasserman’s *Monte Carlo Methods in Financial Engineering* synthesised variance reduction, discretisation, Greeks, American-option methods, credit simulation, and rare-event problems into a computational canon [153]. Mike Giles’ multilevel Monte Carlo transformed the complexity analysis of simulation by combining many cheap low-resolution paths with fewer expensive high-resolution paths [150, 151]. Luca Capriotti’s work on adjoint algorithmic differentiation showed how sensitivities for large books could be computed at a cost much closer to one valuation than one valuation per risk factor [78]. Peter Jaeckel’s numerical work emphasised robustness, inversion, calibration, and practical algorithms in precisely the places where elegant formulas can fail on production edge cases.

Claudio Albanese represents the strand in which numerical finance was rethought around massively parallel hardware. His operator methods, calibration work, GPU pricing libraries, and later XVA simulation projects treated architecture and numerical method as a joint design problem [3, 4]. The point was not merely to port a CPU loop to a graphics card, but to reformulate simulation and linear-algebra workloads so that the hardware’s parallel structure was economically useful.

23.2 MATLAB, MathWorks, and the era of the numerical workbench

Before Python became the default research language at many firms, MATLAB provided a highly productive environment for matrix computation, optimisation, signal processing, statistics, and rapid prototyping. MathWorks’ finance toolboxes and industry interfaces helped create a generation of quants for whom a numerical workbench linked research, visualisation, calibration, and application deployment. Steve Wilcockson’s long role in financial-services technology at MathWorks is representative of the vendor-side specialists who translated quant requirements into product features, examples, interfaces, and training [92, 382].

The significance of MATLAB is cultural as much as technical. It normalised vectorised thinking, interactive numerical experimentation, and packaged numerical algorithms. Those habits later transferred easily to NumPy, pandas, SciPy, JAX, PyTorch, and the broader Python ecosystem.

23.3 Python and the open-source turn

Yves Hilpisch’s books, courses, and community activity helped make the case that Python could span research, backtesting, data handling, derivatives, and deployment [192]. Python’s eventual dominance was not the victory of one numerical language over another so much as the convergence of several ecosystems: scientific computing, data engineering, web services, machine learning, notebooks, package management, and cloud infrastructure. Quantitative research became easier to connect to general-purpose software engineering.

This open-source turn also weakened the old boundary between “quant code” and “technology code.” A research notebook could import industrial optimisation libraries, automatic differentiation, GPUs, language models, and database clients. At the same time, the ease of composition increased model-governance problems: dependency risk, version drift, hidden preprocessing, and non-deterministic ML pipelines became part of quantitative validation.

23.4 APL, A+, k, q, kdb+, and the array-language tradition

Arthur Whitney is one of the most important programming-language figures in financial technology. His APL background led to A/A+ at Morgan Stanley and then to k, kdb+, and q through Kx Systems. The essential idea is that concise array operations and column-oriented data structures are exceptionally well matched to market time series and event data. A+ was developed for numerically intensive financial applications at Morgan Stanley; k and q pushed the same array philosophy into an integrated language-and-database environment [77, 243].

Fintan Quill became a prominent educator and engineer in the kdb+ ecosystem, helping explain time-series and streaming analytics to users beyond the original bank community [325]. Jan Novotny, Paul Bilokon, Aris Galitos, and Frédéric Déléze’s *Machine Learning and Big Data with kdb+/q* explicitly connected the traditional q/kdb+ stack to machine learning, econometrics, text processing, and agent-based methods [304]. Alexander Unterrainer’s DefconQ work represents the later community-led diffusion of kdb+/q knowledge through consulting, blogging, meetups, and practical training.

The importance of q/kdb+ to quant history is larger than its market share. It represents a distinct epistemology of financial computing: data are ordered event streams; joins are temporal; vectorisation is natural; the database and language are one environment; and terseness is a productivity feature for expert users. That style influenced how entire electronic-trading organisations thought about market data.

23.5 Proprietary bank languages: Slang, A+, and internal query systems

Large banks also built proprietary domain-specific languages and data systems. Goldman Sachs’ SecDB/Slang ecosystem pioneered by Sam Wisnia and his colleagues is the best-known surviving example: a proprietary language and object environment became deeply intertwined with front-office pricing and risk [75]. Morgan Stanley’s A+ history is public, while other internal query languages and frameworks are much less documented. Nicholas Zinn’s and Ron Robbins’ work in Morgan Stanley technology and the “query language” lineage supplied in practitioner accounts belongs to this less visible layer of bank infrastructure.

Nicholas Zinn, Martin Zinkin, Jason Batt, William Osborn, Rob Smith, and other senior practitioners named in desk histories are important partly because proprietary platforms make conventional bibliometrics systematically understate their contribution. A production query language,

risk engine, or e-trading platform can affect thousands of daily decisions without generating a citable paper.

23.6 GPUs: Albanese, Giles, and massively parallel finance

GPUs entered quantitative finance naturally through Monte Carlo, finite differences, calibration, and linear algebra. A useful historical marker is the first Thalesian full-day “GPUs in Finance” workshop at Imperial College in September 2009, which included Claudio Albanese, Mike Giles, Paul Kelly, David Thomas, Thomas Bradley, Robert Tong, and Gernot Ziegler. The programme combined CUDA laboratories with Monte Carlo, finite-difference, random-number, and numerical-library work, making the workshop a snapshot of the moment when GPUs were moving from graphics hardware to serious scientific-finance infrastructure [363].

Giles’ work on GPU computation and multilevel Monte Carlo, Albanese’s CUDA pricing libraries, and wider work by Arnaud Doucet and collaborators on massively parallel simulation demonstrate a recurring pattern: hardware matters most when it changes which statistical or numerical method is affordable [247]. The GPU later became even more central with deep learning, making finance one of the domains in which scientific-computing and AI hardware lineages converged.

23.7 FPGAs and dataflow computing: Mencer, Gaydadjiev, and Maxeler

Field-programmable gate arrays offered a different proposition: deterministic pipelines, custom data paths, and very high throughput per watt. Oskar Mencer and Georgi Gaydadjiev, together with Tobias Becker and Stephen Weston, developed the Maxeler dataflow approach and applied it to computational finance [43]. The Maxeler model treats computation as a streaming dataflow graph compiled onto reconfigurable hardware rather than as an instruction stream executed by a conventional CPU.

For low-latency finance, FPGAs became associated with market-data decoding, feed handling, order gateways, and selected numerical kernels. For computational finance, the broader lesson is architectural: a pricing or risk algorithm can be redesigned around data movement and pipeline structure, not merely “accelerated.” Mencer and Gaydadjiev therefore belong in the same historical chapter as Albanese and Giles, but on a different branch: GPU throughput and FPGA/dataflow determinism are complementary answers to the limits of general-purpose processors.

24 Data terminals, platforms, communities, conferences, and the labour market

The profession was shaped by institutions that do not appear in stochastic-calculus textbooks. Data vendors standardised observation; software vendors standardised computation; conferences and seminar series created common vocabularies; specialist media and job platforms made the labour market legible; and professional societies created reputational networks. These institutions helped determine which ideas travelled from one desk or university to another.

24.1 Bloomberg, Refinitiv, and the standardisation of observable finance

Bloomberg and the Reuters–Thomson Reuters–Refinitiv lineage did far more than sell prices. Terminals, real-time feeds, symbology, reference data, news, messaging, analytics, and increasingly

APIs created a shared informational substrate across institutions. A quant model cannot be separated from identifiers, corporate actions, calendars, curves, fixings, timestamps, and data entitlements. Vendors standardised these mundane but essential objects and thereby reduced the cost of reproducing calculations across firms.

Bloomberg also developed in-terminal analytics and later public-facing research such as BloombergGPT, showing how a data vendor can become a model and AI platform [387]. Refinitiv, now within LSEG, similarly combined data, analytics, market infrastructure, and model-risk functions. The resulting ecosystems shaped what “market data” meant to quants and made vendor APIs a persistent layer between mathematical models and exchange reality.

A younger generation of data and analytics firms further illustrates the migration from data as a vendor product toward data and analytical infrastructure embedded directly in quantitative workflows. Paul Humphrey, drawing on senior roles spanning electronic broking, exchanges, banking, and market data, has led BMLL as it developed harmonised full-depth historical order-book data and analytics into a research substrate for market microstructure, execution analysis, and quantitative trading. Alexei Jiltsov represents an unusually direct bridge between academic financial economics and production markets: his research examined inflation risk premia, model uncertainty, heterogeneous beliefs, and option markets, while his subsequent career encompassed quantitative foreign exchange, systematic investment, and the co-founding of Tradefeedr, whose execution-data and analytics infrastructure addresses the fragmentation of trading information across institutions. Christina Qi provides a complementary trajectory from quantitative trading to market-data infrastructure: after co-founding the high-frequency trading firm Domeyard, she co-founded Databento, whose API-centred approach seeks to reduce the engineering and commercial friction involved in obtaining historical and real-time exchange data.

Turnleaf Analytics illustrates a related transition from the provision of raw or normalised data toward the industrialisation of quantitative inference itself. Co-founded by Alexander Denev and Saeed Amen, whose backgrounds span quantitative research, artificial intelligence, systematic trading, alternative data, and financial modelling, Turnleaf applies machine learning and alternative information sources to economic forecasting, particularly inflation. Its development is representative of a broader shift in which specialised quantitative firms increasingly package not merely observations but reproducible forecasting processes, transforming heterogeneous economic information into model-ready signals and probabilistic views. Taken together, BMLL, Tradefeedr, Databento, and Turnleaf illustrate the growing importance of an intermediate computational layer between primary data sources and final trading or investment decisions: acquisition, normalisation, enrichment, forecasting, analytics, and distribution themselves become specialised quantitative products.

24.2 The Thalesians and the seminar/community model

The Thalesians were established in 2008 and incorporated in 2009, with an early focus on quantitative-finance seminars, education, research, and later machine learning and high-performance computing. The seminar series began in London in January 2009 and expanded internationally; in 2011 an IAQF/Thalesians series began in New York with Harvey Stein playing a central organisational role [364, 365]. The 2009 GPU workshop is an example of the community’s role in bringing together academics, bank quants, numerical-software specialists, and hardware vendors at an early point in GPU adoption.

The Thalesians’ significance is best understood as connective tissue. Seminar communities shorten the diffusion delay between a paper, a desk implementation, and a new research question. They also create weak ties across banks that cannot share proprietary code but can share mathematical

problems, public benchmarks, numerical ideas, and hiring signals. Saeed Amen’s later *Trading Thalesians* and the organisation’s training and publishing activity extended that function into systematic trading and data science.

24.3 WBS Training and Neil Fowler

Neil Clive Fowler founded WBS Training in 2000 as a quantitative-research and risk-management training and conference company. Its long-running Quantitative Finance Conference, specialised workshops, Women in Quantitative Finance events, and later machine-learning programmes constitute another durable professional network [324]. WBS is historically useful because it documents the changing agenda of the profession: early derivatives and risk topics coexist over time with xVA, electronic trading, machine learning, climate risk, and AI.

The WBS “Legend in Quant Finance” awards and related events also show how the profession creates its own memory. Recognition of figures such as Vašíček and Hélyette Geman turns technical lineages into a shared professional canon, while Women in Quantitative Finance events respond to the composition and career structure of the field itself.

24.4 Informa: Global Derivatives and QuantMinds

Informa’s Global Derivatives conference, later rebranded as QuantMinds, is one of the most visible examples of a conference evolving with the profession. The organisation describes more than three decades of quant-finance activity, while participants have noted the shift from a programme dominated by derivatives pricing toward a broader mix of portfolio construction, machine learning, risk, data, liquidity, and quant development [13, 210]. The rebranding is historically revealing: “derivatives” ceased to be a sufficient umbrella for the profession that derivatives mathematics had helped create.

Conferences such as Global Derivatives/QuantMinds did more than disseminate finished results. They provided a venue where bank quants could present partially public versions of proprietary research, where academics learned which problems desks would fund, where software firms demonstrated new infrastructure, and where a labour market formed around reputation and repeated contact.

24.5 eFinancialCareers and Sarah Butcher: making the quant labour market visible

The rise of specialist online careers media changed how quant careers were understood. eFinancialCareers combined vacancies with reporting on pay, hiring, organisational changes, skills, and career transitions. Sarah Butcher, its long-serving global editor, became one of the most visible journalists covering the microstructure of the finance labour market [118]. Her influence is not a technical contribution to asset pricing; it is informational infrastructure for the profession itself.

This distinction matters. Compensation benchmarks, stories about which languages banks use, reports on hiring by hedge funds, and accounts of role changes affect career choices and therefore the supply of skills. eFinancialCareers helped make previously opaque internal categories—strats, quant developers, desk quants, model validators, electronic-trading researchers—legible to people outside a particular bank. It also documented proprietary languages such as Slang and A+, thereby preserving fragments of a technical history that firms themselves rarely publish.

24.6 Risk.net, awards, and contemporary reputation

Risk magazine and Risk.net have long occupied a hybrid position between trade press and technical publication. Model papers, practitioner debates, annual awards, “Quant of the Year,” and “Rising Star in Quant Finance” recognitions create a public record of what the industry considers technically significant. The 2025 Rising Star award to Milena Vuletić for VolGAN and the 2026 award to David Itkin for price-impact research are examples of awards functioning as snapshots of the field’s frontier: generative volatility models in one year, high-dimensional market impact in the next [330, 331].

24.7 Universities, practitioner programmes, and hybrid careers

Imperial, Oxford, NYU, Columbia, Carnegie Mellon, Princeton, Chicago, and other universities developed mathematical-finance, financial-engineering, computational-finance, and later data-science programmes that professionalised the quant pipeline. The traffic was two-way: academics moved into banks and funds, while practitioners returned as lecturers, research fellows, and supervisors. Tom Hadfield’s 2026 return to Imperial after sixteen years at G-Research is a particularly vivid example [208].

Joerg Osterrieder and Jochen Papenbrock represent another strand of this educational ecosystem, combining applied research, conferences, executive education, and financial data science. The profession increasingly learns through a mixture of degree programmes, vendor training, open-source notebooks, specialist conferences, online courses, and practitioner books rather than through one canonical academic curriculum.

25 Contemporary frontiers: reinforcement learning, formal verification, quantum computing, and model governance

25.1 Reinforcement learning as the reunion of control and data

The current enthusiasm for reinforcement learning is historically less discontinuous than it appears. The Bellman/control lineage of Davis, Bertsekas, Pham, and Powell already treats finance as sequential decision making. What changes with RL is that transition laws, value functions, or policies may be learned from data and simulation rather than fully specified analytically.

Igor Halperin’s QLBS programme deliberately uses Black–Scholes as a benchmark environment for this transition [170]. Petter Kolm and Gordon Ritter’s work extends RL to dynamic hedging and surveys financial applications more broadly [238]. Ben Hambly, Renyuan Xu, and Huining Yang’s survey places execution, market making, pricing, hedging, portfolio allocation, and robo-advice in one RL taxonomy [172]. These strands make Bertsekas and Powell historically relevant to quant finance even when their primary audience lies outside finance: they provide the general theory of approximate sequential optimisation on which many financial RL methods implicitly rely.

The most important conceptual correction is that RL is not “prediction plus trading.” In a proper formulation, the state must include inventories, costs, constraints, and information; actions affect future states; and the objective is a path-dependent functional of rewards and risks. That places RL closer to stochastic control and market microstructure than to ordinary supervised learning.

25.2 Formal verification and Lean: from plausible mathematics to machine-checked finance

Formal verification introduces a new possible stage in the history of quant methodology. Classical finance papers rely on human proof; production libraries rely on unit tests, code review, and model validation. A theorem prover such as Lean makes it possible, at least in principle, to specify financial definitions and prove properties of algorithms against a small trusted kernel.

Kevin Buzzard has been one of the leading advocates of formalising modern mathematics in Lean. Cristopher Salvi’s current research explicitly includes formal verification and Lean 4, alongside stochastic analysis and machine learning [209]. Robert Smith brings the complementary perspective of a senior bank quant and automated-trading research leader. Their 2026 Logos Research initiative, with Buzzard as scientific adviser, explicitly seeks to extend Lean-based formalisation beyond pure mathematics toward applied and high-stakes algorithmic reasoning [263].

For quantitative finance, the attraction is obvious: pricing identities, arbitrage arguments, risk aggregation rules, portfolio constraints, and parts of model implementation have specifications that can be stated mathematically. The harder problem is connecting those proofs to floating-point code, data assumptions, calibration procedures, and changing market conventions. Formal proof will therefore not eliminate model risk; it can isolate one component of it—logical and algorithmic correctness—more sharply than ordinary testing.

25.3 Quantum computing for financial mathematics

Antoine Jacquier and Oleksiy Kondratyev have developed one of the clearest practitioner-academic programmes on quantum computing for finance, including teaching and a book on quantum machine learning and optimisation. Their work covers quantum optimisation, variational circuits, quantum machine learning, and the realistic constraints of noisy intermediate-scale quantum hardware [214]. With Gordon Lee and Mugad Oumgari they also provided an overview of quantum computing for financial mathematics [215]. Jacquier’s related work includes quantum algorithms for finance PDEs and quantum generative models for volatility [25, 139].

This should be distinguished from older “quantum finance” that imports quantum-mechanical formalisms as models of markets. The Jacquier–Kondratyev line is about using quantum computers to solve classical financial problems: optimisation, Monte Carlo, linear systems, generative modelling, and related computational tasks. Whether a durable quantum advantage will emerge remains an empirical hardware question, but the research programme belongs in the history because it continues the same search seen with GPUs and FPGAs: match the financial algorithm to a new computational substrate.

25.4 Model risk, AI governance, and Sanja Hukovic

As models became more numerous and ML systems less transparent, model risk evolved into a profession of its own. Sanja Hukovic’s career in quantitative risk control and model-risk leadership, including work on risk standards and later AI/model governance, illustrates this institutionalisation [199]. The model validator is historically significant because the role formalises scepticism: an organisation creates a separate quantitative function whose task is to challenge assumptions, implementation, data, calibration, and use.

The arrival of generative AI intensifies the same problem. A deterministic pricing library can be regression-tested against benchmarks; an LLM may change outputs with context and sampling.

Model governance must therefore widen from parameter validation toward data lineage, evaluation design, prompt/tool control, monitoring, and formal constraints. In that sense, the Hukovic model-risk lineage and the Buzzard–Salvi–Smith formal-verification lineage are converging on a common question: how does a financial institution know that a computational decision process deserves trust?

25.5 Recent rising stars and the changing frontier

David Itkin and Milena Vuletić provide two useful recent markers. Itkin’s work applies stochastic portfolio theory and control ideas to markets with nonlinear price impact, a problem forced by the recognition that an investor is not outside the market but changes the prices at which the portfolio is implemented [212]. Vuletić’s work with Rama Cont uses generative learning for implied-volatility surfaces while preserving financial structure [89]. Their Risk.net Rising Star awards in 2026 and 2025, respectively, capture two contemporary themes: endogenous market impact and data-driven generative modelling [330, 331].

26 A prosopographical index of the modern quant ecosystem

Table 2 provides a structured but necessarily incomplete index of individuals selected for their representative contributions to the development of quantitative finance, including foundational theory, methodological innovation, practitioner influence, software and computational infrastructure, institution-building, education, and the evolution of professional practice. It is intended to illustrate the principal strands developed in this history rather than to constitute an exhaustive canon of the field. It is deliberately broader than a list of academic model inventors. Some entries are theorists, some practitioner-authors, some software and hardware pioneers, some firm builders, and some case studies in risk and governance. Where the public record is too sparse to support a specific technical attribution, the entry says so rather than inferring one.

Table 2: Selected people and their principal place in the history developed here.

Person	Strand	Contribution / historical role
John Hull	derivatives/rates	Derivatives pedagogy; Hull–White rates models; credit and xVA work.
Alan White	derivatives/rates	Hull–White term-structure modelling; derivatives and credit collaboration with Hull.
Dilip Madan	jumps/credit	Variance-gamma and Lévy-process finance; credit and option modelling.
Marcos López de Prado	ML/systematic	Financial ML, backtest overfitting, portfolio methods, research methodology.
Matthew Dixon	ML	Deep learning, market data, quantitative ML pedagogy and applications.
Igor Halperin	RL/credit	QLBS reinforcement-learning approach to option pricing/hedging; credit/ML work.
Blanka Horvath	volatility/ML	Rough volatility, deep calibration, generative and signature-based methods.

Person	Strand	Contribution / historical role
Saeed Amen	systematic/data/forecasting	Systematic FX research, alternative data and quantitative software; co-founder and CTO/Head of Data of Turnleaf Analytics, applying machine learning and alternative data to macroeconomic forecasting.
Antoine (Jack) Jacquier	volatility/quantum	Asymptotics, volatility, stochastic analysis, quantum computing for finance.
Rama Cont	stochastic analysis/microstructure	Stylised facts, Lévy models, functional Itô calculus, market microstructure, systemic risk.
Emanuel Derman	derivatives/profession	Black–Derman–Toy, local volatility, and influential writing on quant identity.
Nassim Nicholas Taleb	tails/convexity	Practitioner critique of Gaussian risk, convexity, tail exposure and model error.
Salih Neftci	pedagogy/engineering	Major textbooks connecting stochastic calculus with financial engineering.
Paul Wilmott	derivatives/pedagogy	PDE finance, textbooks, magazine/community, model critique and professional education.
Johannes Muhle-Karbe	frictions/impact	Rigorous trading with frictions, liquidity, price impact and stochastic control.
Alexander Lipton	cross-asset modelling	Rates, FX, credit, numerical methods, ML/quantum; practitioner-academic bridge.
Jon (Joe) Gregory	xVA	Practitioner synthesis of counterparty risk, CVA and xVA.
Harvey Stein	risk/credit/community	Model risk, credit/counterparty analytics, computational finance; IAQF/Thalesians community role.
Hans Buehler	ML/hedging	Deep hedging and neural policy methods for derivatives under frictions.
Alexander Gerko	institution building	Research-intensive electronic market making and systematic trading organisation.
Steven Shreve	stochastic finance	Canonical stochastic-calculus and mathematical-finance pedagogy.
Bernt Øksendal	stochastic calculus	Widely used SDE text that diffused Itô-calculus methods across applied fields.
Miquel Noguer i Alonso	AI/education	Applied AI/alternative-data and quantitative-investment education; normalisation of “Miquel Alonso.”
Irene Aldridge	HFT	Practitioner literature on high-frequency and electronic trading.
Marco Avellaneda	derivatives/microstructure	Uncertain volatility, statistical arbitrage, and Avellaneda–Stoikov market making.
Petter Kolm	portfolio/RL	Portfolio research; RL for hedging and finance; education at NYU.
Gordon Ritter	RL/systematic	RL for dynamic hedging and financial sequential decisions; systematic research.

Person	Strand	Contribution / historical role
Alec N. Schmidt	ML/practitioner	Practitioner writing on ML and quantitative trading.
Albert Shiryaev	probability/finance	Optimal stopping, sequential analysis, probability foundations, mathematical finance.
Damiano Brigo	rates/credit/xVA	Interest-rate models, credit derivatives, counterparty risk and xVA.
Fabio Mercurio	rates	Interest-rate derivatives, LIBOR market models and calibration; co-author with Brigo.
Jim Simons	systematic/institution	Renaissance Technologies and the research-firm model of systematic investing.
Svetlana Borovkova	commodities/risk	Commodities, energy, volatility, quantitative risk and climate-risk analytics.
Jonathan Ingersoll	financial economics	Asset-pricing and contingent-claims theory; influential graduate text.
Robert Jarrow	credit/rates	Term structure, derivatives and reduced-form credit-risk modelling.
Kenneth Arrow	state pricing/equilibrium	State-contingent claims, risk allocation and general equilibrium foundations.
Mathieu Rosenbaum	microstructure/rough volatility	Hawkes processes, high-frequency limits, market impact, rough volatility.
Marek Musiela	rates/martingales	Musiela parametrisation; martingale methods and interest-rate modelling.
Victor Niederhoffer	systematic trading	Statistical/empirical trading tradition and later tail-risk case study.
Stephen Ross	asset pricing	Arbitrage Pricing Theory and broad contributions to financial economics.
Oldřich Vašíček	rates	Mean-reverting short-rate model and modern fixed-income modelling lineage.
Peter Carr	derivatives/volatility	Static replication, transforms, volatility derivatives, variance swaps and jump models.
Jean-Philippe Bouchaud	econophysics/microstructure	Order flow, impact, anomalous diffusion, agent/statistical-physics approaches.
Charles-Albert Lehalle	execution/microstructure	Optimal trading, order books, transaction-cost analysis and market practice.
Didier Sornette	econophysics/crashes	Critical phenomena, bubbles, crash diagnostics and collective behaviour.
Darrell Duffie	asset pricing/credit/OTC	Dynamic asset pricing, credit risk, OTC-market structure and financial institutions.
Daniel Duffy	computational finance	C++, finite differences and practitioner numerical-finance education.
Nicole El Karoui	stochastic control	BSDEs, incomplete markets, stochastic control and the French mathematical-finance school.
Hélyette Geman	commodities/measure change	Numeraire methods, commodities, insurance and derivatives.
Espen Gaarder Haug	practitioner/history	Option formula compendia and revisionist history of option pricing.

Person	Strand	Contribution / historical role
Jeff Dewynne	PDE finance	Oxford PDE derivatives tradition; co-author with Wilmott and Howison.
Sam Howison	PDE/stochastic finance	Asymptotics, PDEs and mathematical finance; Oxford school.
Piotr Karasiński	rates	Black–Karasinski short-rate model and industrial derivatives research.
Andrew Lo	empirical/systematic	Adaptive Markets, hedge-fund analysis, econometrics and market efficiency.
Marek Capiński	pedagogy	Mathematical-finance textbooks and education.
Marc Yor	probability	Exponential Brownian functionals, Bessel processes and tools used in derivatives.
D. E. Shaw	systematic/institution	Quantitative-investment firm built around computer science and research infrastructure.
Igor Tulchinsky	systematic/institution	WorldQuant and the large-scale “alpha factory” research organisation.
Attilio Meucci	portfolio/risk	Quantitative risk/portfolio framework and entropy pooling.
Paul Glasserman	simulation/risk	Monte Carlo, rare events, derivatives and financial risk.
Peter Jaeckel	numerics	Monte Carlo, calibration and robust practitioner numerical methods.
Andrew Kalotay	fixed income	Callable-bond, mortgage and interest-rate analytics.
Mark Joshi	quant development	C++ design patterns and derivatives-software pedagogy.
Luca Capriotti	AAD/risk	Adjoint algorithmic differentiation and efficient Greeks/risk.
Mike Giles	computational finance	Multilevel Monte Carlo, GPU/scientific computing and numerical analysis.
Albert Hibbs	probability/prehistory	Early stochastic-process and physics lineage often cited in the prehistory of modern quantitative modelling.
Frank Fabozzi	fixed income/publishing	Extensive fixed-income, portfolio and quantitative-investment publishing.
Joseph Simonian	investment science	Applied quantitative investment, ML and institutional portfolio research.
Robert Almgren	execution	Mathematical market impact and optimal execution.
Neil Chriss	execution	Almgren–Chriss optimal execution and practitioner quant leadership.
Raphael Douady	nonlinear risk	Dynamical systems, instability and quantitative-risk research.
Jakša Cvitanić	stochastic control/finance	Constrained portfolios, incomplete markets, contracting and mathematical finance.
Vladimir Piterbarg	rates/xVA	Interest-rate derivatives, collateral/funding and post-crisis valuation.

Person	Strand	Contribution / historical role
Vladimir Lucic Eyal Neuman	practitioner quant execution/control	Practitioner quantitative-finance lineage. Price impact, optimal trading, interacting agents and mean-field/crowding models.
Yufei Zhang	control/RL	Continuous-time stochastic control and reinforcement-learning methods.
Ben Hambly	probability/RL	Mathematical finance and influential survey work on RL in finance.
Abbas Edalat	exact computation/foundations	Domain theory, exact computation, optimisation and correctness-oriented computing.
Huy��n Pham	stochastic control	Control, viscosity methods, optimal investment/hedging and later ML methods.
William Toy Patrick Hagan	rates volatility/rates	Black–Derman–Toy interest-rate model. SABR asymptotics and practitioner derivatives modelling.
Deep Kumar	volatility/rates	Co-author of SABR.
Andrew Lesniewski	volatility/rates	Co-author of SABR; rates and derivatives modelling.
Diana Woodward	volatility/rates	Co-author of SABR.
Bruno Iksil	governance case study	London Whale episode: concentration, model change, liquidity and risk governance.
J��r��me Kerviel	operational-risk case study	Unauthorised trading and control-circumvention case study.
Arnaud Doucet	SMC/computation	Particle filtering and sequential Monte Carlo with applications to finance and statistics.
��lvaro Cartea	algorithmic trading	Stochastic control of execution and market making; textbook synthesis.
Olivier Gu��ant Alexey Poyarkov	market making practitioner quant	Optimal execution and market-making control. Name supplied for practitioner-history coverage; public sources are insufficient here for a specific technical priority claim.
Yuri Bedny Igor Vereshchetin Alexandre Antonov	practitioner quant practitioner quant derivatives	Practitioner quantitative-finance lineage. Practitioner quantitative-finance lineage. Practitioner research in rates/volatility and quantitative modelling.
Serguei Issakov	derivatives	Practitioner quantitative modelling; rates/derivatives lineage.
Alexander Sokol	credit/xVA	Counterparty credit risk, simulation and quantitative risk.
Alexander Denev	risk/ML/forecasting	Probability, stress testing, credit/counterparty risk and data-driven modelling; co-founder and CEO/Head of R&D of Turnleaf Analytics, developing machine-learning-based macroeconomic and inflation forecasting.
Vincent Berthelemy Damien Couture	practitioner quant practitioner quant	Practitioner quantitative-finance lineage. Practitioner quantitative-finance lineage.

Person	Strand	Contribution / historical role
Mark H. A. Davis	control/credit	Stochastic control, PDMPs, transaction costs and credit risk.
Joerg Osterrieder	education/fintech	Applied quant research, data science, fintech and professional education.
Jochen Papenbrock	ML/industry	Applied ML/quant research and industry education.
Steve Wilcockson	software ecosystem	MathWorks financial-services bridge; quantitative research-to-production tooling.
Paul Humphrey	market data/infrastructure	BMLL leadership; industrialisation and distribution of harmonised full-depth order-book data and analytics for quantitative research, market microstructure and execution.
Alexei Jiltsov	FX/systematic/execution analytics	Academic work on inflation risk premia, heterogeneous beliefs and option markets; quantitative FX and systematic investment; co-founder of Blacktree and Tradefedr, developing common execution-data and analytics infrastructure across market participants.
Christina Qi	HFT/data infrastructure	Co-founder of Domeyard and Databento; trajectory from high-frequency trading to API-native historical and real-time market-data infrastructure for quantitative research and trading.
Yves Hilpisch	Python	Major role in Python-for-finance education and open-source quant culture.
Jan Novotny	q/kdb+/systematic	Co-author of the kdb+/q ML text; eFX/systematic and market-microstructure work.
Aris Galiotos	q/kdb+ engineering	Co-author of the kdb+/q ML text; large-scale kdb+ engineering and implementation.
Frédéric Déléze	q/kdb+/quant research	Co-author of the kdb+/q ML text; quantitative research and kdb+/q applications.
Alexander Barzykin	market making academic and practitioner.	
Alexander Unter-rainer	q/kdb+	Practitioner educator, DefconQ founder and kdb+/q community contributor.
Panos Parpas	optimisation	Stochastic/robust optimisation and decomposition methods relevant to finance.
Berç Rüstem	robust optimisation	Robust control/optimisation and financial risk.
Dimitri Bertsekas	control/RL foundations	Dynamic programming, optimal control, approximate DP and RL.
Warren Powell	sequential decisions	Approximate dynamic programming, stochastic optimisation and unified RL framework.
Elisa Alòs	Malliavin/volatility	Malliavin-calculus methods for volatility, Greeks and asymptotics.

Person	Strand	Contribution / historical role
Sarah Butcher	labour-market media	eFinancialCareers reporting made quant hiring, compensation, roles and technologies more legible.
Katya Gorbatyuk	practitioner/institutional	Practitioner and institutional quantitative-finance lineage.
Veronika Lunina	practitioner quant	Practitioner quantitative-finance lineage.
Youssef Elouerkhaoui	credit/quant leadership	Credit modelling and senior quantitative leadership across credit, CVA, funding, capital and markets.
Sam Wisnia	practitioner quant	Practitioner quantitative-finance lineage.
Martin Zinkin	e-trading leadership	Senior electronic-trading and quantitative leadership.
Jason Batt	e-trading/quant leadership	Senior bank quant/e-trading practitioner and organisation builder.
David Itkin	price impact/SPT	Stochastic portfolio theory with price impact; Risk.net 2026 Rising Star.
Milena Vuletić	volatility/ML	VolGAN and arbitrage-aware generative volatility modelling; Risk.net 2025 Rising Star.
Sanja Hukovic	model risk	Quantitative risk standards, model validation, model/AI risk governance.
Tom Hadfield	systematic/AI/formalisation	Co-Research quant leadership; later Imperial work on quantitative finance, LLM mathematics and formalisation.
Kevin Buzzard	formal verification	Lean/mathlib advocacy and machine-checked modern mathematics.
Robert Smith	credit/e-trading/formalisation	Senior credit/automated-trading quant leadership; Logos formal-verification initiative.
Cristopher Salvi	stochastic ML/formal verification	Signatures, generative SDEs, stochastic ML, Lean/formal verification.
Oleksiy Kondratyev	quantum/quant practice	Quantum machine learning and optimisation in finance; senior practitioner perspective.
William Osborn	e-trading/proprietary systems	Citi electronic-trading practitioner; associated in practitioner history with proprietary platform development.
Nicholas Zinn	risk/practitioner	Prime-brokerage risk and senior practitioner lineage.
Ron Robbins	bank technology	Morgan Stanley technology/query-language practitioner history.
Arthur Whitney	programming languages	A/A+, k, q and kdb+; array-language/database revolution in market data.
Fintan Quill	kdb+	KX engineering, time-series/database education and ecosystem development.
Oskar Mencer	FPGA/dataflow	Maxeler and dataflow computing.
Georgi Gaydadjiev	FPGA/dataflow	Reconfigurable/dataflow architecture; co-author of Maxeler finance work.

Person	Strand	Contribution / historical role
Claudio Albanese	GPUs/numerics	Operator methods, high-performance pricing, GPU/CUDA finance and XVA simulation.

The table intentionally preserves several entries with limited public documentation. They are not presented as equivalent to the well-documented academic lineages; rather, they flag an important agenda for oral history. Quantitative finance contains a large “dark literature” of internal notes, codebases, model libraries, query languages, and desk practices. An exhaustive history will ultimately need interviews and archival material in addition to published papers.

27 A taxonomy of quant roles

The term “quant” now spans enough occupations that historical discussion benefits from a taxonomy. Table 3 separates roles by their primary object, although real jobs frequently combine several columns.

Table 3: A functional taxonomy of quantitative finance roles.

Role family	Primary object	Typical methods	Characteristic failure mode
Pricing / derivatives	contingent claim	stochastic calculus, PDEs, trees, Monte Carlo, calibration	model misses smile, dynamics, funding, or exercise behaviour
Systematic research	forecast and portfolio	statistics, factors, optimisation, ML, backtesting	overfit signal, leakage, crowding, regime change
Quant trading	decision policy	stochastic control, forecasting, state estimation	poor live adaptation, hidden exposure, execution drag
Execution research	order trajectory	impact models, microstructure, control, causal analysis	signal consumed by spread, impact, or information leakage
Electronic market making	quotes and inventory	order-book models, control, queue estimation, low-latency statistics	adverse selection, inventory accumulation, latency race
Risk quant	loss distribution	simulation, stress testing, time series, dependence models	tail misspecification, aggregation error, false diversification
XVA / counterparty	exposure and funding	simulation, credit models, PDE/BSDE methods	wrong netting/collateral/funding assumptions
Model validation	model claim	benchmarking, independent implementation, statistical tests	false assurance from incomplete challenge

Role family		Primary object	Typical methods	Characteristic failure mode
Quant oper	devel-	production seman- tics	numerical meth- ods, C++/Python, databases, distributed systems	implementation diverges from research model
Data quant	/ ML	representation and prediction	feature engineering, ML, NLP, deep learning	data leakage, drift, opaque failure, weak provenance
Portfolio	quant	constrained alloca- tion	covariance/factor models, optimisation, transaction-cost models	unstable estimates, unin- tended concentration
Regulatory quant		capital / compli- ance measure	regulatory models, stress, backtesting, reporting	rule/model mismatch, gov- ernance or audit failure

The taxonomy clarifies why debates over whether “quants” are mathematicians, programmers, traders, or statisticians are unproductive. The profession is defined by the coupling. Different market structures and business models demand different points on the mathematics–statistics–software–trading simplex.

28 Five co-evolutionary forces: a synthesis

The history assembled above can be compressed into five interacting trajectories.

1. Tradability. The universe of standardised claims expanded: bills, shares, bonds, futures, options, swaps, securitisations, exchange-traded funds, crypto-assets, and bespoke but modellable contracts. A quant needs an object that can be parameterised and acted upon.

2. Formalisation. Financial reasoning moved from arithmetic and bookkeeping toward probability, optimisation, stochastic calculus, econometrics, control, and machine learning. Each formalism made some previously qualitative problem computationally explicit.

3. Observability and speed. Market information moved from periodic and local to continuous and global. The same transition created both richer statistical evidence and shorter half-lives for exploitable information.

4. Industrial organisation. Finance split into specialised desks, funds, risk functions, technology groups, and control functions. The economic value of specialisation justified a labour market for people who produce models rather than merely use them.

5. Constraint. Regulation, capital, collateral, liquidity, market design, and governance became endogenous parts of the modelling problem. The “best” strategy is never simply the maximum of an unconstrained expected-return function; it is the best feasible strategy inside an institution and a rule system.

These forces form feedback loops. A mathematical model can make a product easier to quote, which increases trading, which generates data, which attracts competition, which changes the market’s microstructure, which motivates new regulation, which creates new models. MacKenzie’s notion of performativity is relevant precisely because financial models do not merely describe markets; once institutionalised, they can alter the practices being described [269, 271].

29 Conclusion

The modern quant is neither a sudden invention of 1973 nor simply an ancient trader equipped with more mathematics. The profession arose when several long-running developments finally intersected: securities became standardised and liquid enough to model; probability and optimisation became expressive enough to represent the relevant risks; computers became fast enough to turn models into repeated decisions; financial institutions became specialised enough to hire dedicated model builders; and regulation made quantitative measurement part of the operating licence of large firms.

The Black–Scholes–Merton revolution remains a pivotal moment because it demonstrated with unusual clarity that abstract mathematics could become market infrastructure. Yet the subsequent history shows that no single model defines the profession. Volatility ceased to be constant; markets became electronic; execution became a control problem; risk became a capital problem; valuation became entangled with funding and counterparty state; and prediction absorbed machine learning. The quant role survived each transition because its deeper function was not attachment to one formalism. It was the disciplined conversion of uncertain information into decisions that can be priced, tested, executed, and controlled.

The latest transition, toward AI-assisted research and increasingly automated market infrastructure, makes that function more rather than less important. When model generation becomes cheaper, verification becomes scarcer. When code can be generated quickly, specification and testing become more valuable. When information is abundant, provenance and relevance dominate access. The history from ledgers to latency thus ends with a paradox: quantitative tools become more autonomous, while the highest-value human work shifts toward deciding what should be formalised, what can be trusted, and what failure would mean.

A Chronology of major milestones

Table 4: Selected milestones in the co-evolution of trading and quantitative finance. Dates are deliberately approximate where the underlying development was gradual.

Date	Milestone	Quantitative significance
Ancient world	ledgers, standard weights, coinage	durable numerical representation of obligations and exchange
Medieval period	bills of exchange and merchant banking	multi-currency, multi-date claims invite discounting and arbitrage calculations
1397–1494	Medici Bank	distributed accounting and foreign-exchange operations at institutional scale
1531	Antwerp bourse	purpose-built concentration of merchants and commercial information

Date	Milestone	Quantitative significance
1560s	Royal Exchange, London	institutionalisation of a permanent commercial meeting place
1602	VOC founded	long-lived, transferable joint-stock claims and active secondary trading
1600s	Amsterdam securities market	shares, public debt, credit, derivatives, and specialised dealing coexist
1688	de la Vega, <i>Confusion de Confusiones</i>	early systematic description of securities-market behaviour
1767	Steuart, <i>Political Oeconomy</i>	economic administration treated as a system susceptible to principles
1776	Smith, <i>Wealth of Nations</i>	foundational account of specialisation, prices, and market coordination
Late 1700s	large merchant banks	increasingly specialised dealing in bills, government securities, FX, and commodities
1800s	organised commodity futures	standardisation of forward risk and margining
1830s–1860s	commercial telegraphy	price-relevant information decouples from physical transport speed
1860s	stock ticker	market prices become a quasi-continuous machine-distributed data stream
1877	Castelli’s options treatise	explicit account of calls, puts, combinations, premiums, and volatility intuition
1900	Bachelier thesis	stochastic-process model of prices and mathematical option valuation
1933	Kolmogorov axioms	measure-theoretic probability provides modern formal foundation
1944	Itô stochastic integral	calculus for diffusion processes later central to derivative pricing
1952	Markowitz portfolio selection	portfolio construction becomes a quadratic optimisation problem
1956	Kelly criterion	information and long-run capital growth linked mathematically
1964–1966	CAPM formulations	expected return related to systematic risk in equilibrium
1965	Fama and Samuelson	statistical/random-price foundations of modern asset-pricing debate
1966–1967	Thorp; Thorp and Kasouf	systematic, computer-assisted risk-taking becomes a public archetype
Late 1960s–1970s	institutional quantitative portfolio management	factor models, indexing, and optimisation enter production asset management
1973	Black–Scholes and Merton	replication/no-arbitrage framework industrialises derivatives pricing and hedging
1973	CBOE era begins	listed options create repeated pricing and market-making demand
1970s–1980s	electronic quotation/trading systems	market state and order submission become increasingly digital

Date	Milestone	Quantitative significance
1979–1981	martingale pricing theory	arbitrage pricing placed on a general probabilistic foundation
1982	ARCH	conditional volatility becomes explicitly dynamic
1980s	derivatives desk expansion	banks recruit mathematicians, physicists, and programmers in growing numbers
1986	GARCH	parsimonious persistent volatility model becomes a workhorse
1987	global equity crash	volatility smile and liquidity risk become impossible to ignore
Late 1980s–1990s	Renaissance/systematic funds	empirical signal research becomes an industrial scientific process
1990s	VaR and Basel market-risk models	quantitative risk becomes embedded in management and regulation
1992	sterling crisis	macro policy, speculative capital, and FX risk interact visibly
1993–1994	Heston, local vol, implied trees	smile-consistent models proliferate
1998	LTCM crisis	leverage, liquidity, crowding, and model risk become central case study
Late 1990s	ECNs and electronic fragmentation	routing and execution become quantitative optimisation problems
2001	Almgren–Chriss execution framework	trading cost and timing risk formalised jointly
2004–2007	MiFID era	European venue competition and transparency requirements expand
2006–2008	electronic market making accelerates	queue position, latency, and microstructure become primary research objects
2007–2008	global financial crisis	counterparty, liquidity, funding, and systemic risks reshape quant priorities
2010	Flash Crash	interacting automated strategies and liquidity withdrawal become policy focus
2010s	xVA	valuation expands to counterparty, funding, collateral, capital, and margin
2010s	GPUs, cloud, large-scale data	simulation and ML workloads become much cheaper and more scalable
2014 on-ward	MiFID II/MiFIR and FRTB development	execution, reporting, market risk, and model governance become more data-intensive
2016–2018	rough volatility	fine-scale regularity of volatility becomes a new modelling frontier
2017 on-ward	deep learning in market data	representation learning applied to order books, returns, and volatility
2020s	alternative data at scale	research boundaries expand toward NLP, geospatial data, and data engineering
2023 on-ward	finance LLMs	language interfaces enter research, coding, information extraction, and model workflows

Date	Milestone	Quantitative significance
2026	increasingly hybrid quant roles	research, software, data, execution, and AI tooling converge further

B A genealogy of representative quantitative models

Table 5: Representative model families and the problem each made tractable.

Family	Canonical date	Core state / equation	Professional use
Bachelier diffusion	1900	$dS = \mu dt + \sigma dW$	stochastic representation of prices, early option valuation
Mean-variance	1952	$w^\top \Sigma w$	portfolio construction, diversification, active risk
CAPM / factor equilibrium	1960s	$\mathbb{E}[R_i] - r_f = \beta_i(\mathbb{E}[R_M] - r_f)$	expected-return benchmarking, cost of capital, factor exposure
Geometric Brownian motion	1960s–1970s	$dS = \mu S dt + \sigma S dW$	positive-price diffusion, basis for canonical option model
Black–Scholes–Merton	1973	replication PDE	option pricing, Greeks, dynamic hedging
Short-rate models	1977 onward	stochastic r_t	bond and interest-rate derivative pricing
Martingale pricing	1979–1981	equivalent martingale measure	general no-arbitrage valuation framework
ARCH/GARCH	1982–1986	recursive conditional variance	volatility forecasting, risk, econometrics
Stochastic volatility	1980s onward	latent variance state	volatility forecasting, option smiles, filtering
Heston	1993	square-root variance diffusion	semi-analytic smile-aware option pricing
Local volatility	1994	$\sigma(S, t)$	exact fit to an option surface under diffusion assumptions
VaR	1990s	loss quantile	aggregate market-risk reporting and capital
Statistical arbitrage	1990s onward	factor/residual dynamics	systematic relative-value trading
Optimal execution	1998 onward	impact + timing-risk control	schedule large trades under cost/risk trade-off
Electronic market making	2000s onward	order-book state + inventory	optimal quotes, inventory and adverse-selection control
Realized volatility	2000s	high-frequency return variation	nonparametric ex-post volatility and forecast inputs
Machine learning	2000s onward	learned nonlinear function	forecasting, feature selection, execution, NLP

Family	Canonical date	Core state / equation	Professional use
Rough volatility	2010s	fractional/Volterra variance	volatility term structure and fine-scale dynamics
xVA	2010s	exposure/funding/capital system	all-in derivatives valuation and desk economics
LLM-assisted finance	2020s	language model + tools/retrieval	text analysis, code, documentation, research interfaces

C Selected crises and dislocations

The table below focuses on episodes with durable relevance to trading, liquidity, risk modelling, or regulation. It is not intended as a complete catalogue of every national banking crisis; broader comparative histories are given by Kindleberger and Aliber [233], Quinn and Turner [326], Reinhart and Rogoff [328].

Table 6: Selected crises and their relevance to quantitative practice.

Date	Episode	Quantitative lesson
1630s	Dutch tulip boom	leverage, optional contracts, and narrative dynamics can precede modern securities theory
1720	South Sea and Mississippi bubbles	financial innovation, credit expansion, and political privilege interact
1763	Amsterdam credit crisis	interconnected bills and leverage propagate funding shocks
1790s	early U.S. securities panics	information speed and public-debt speculation already matter
1825	British banking panic	global capital flows and speculative lending produce systemic stress
1866	Overend Gurney crisis	wholesale funding and liquidity of a major intermediary become systemic
1873	international panic	railroad/real-estate finance and cross-border credit amplify downturn
1907	U.S. banking panic	lender-of-last-resort problem precedes Federal Reserve creation
1929–1933	crash and Great Depression	leverage, banking failure, macro feedback, and institutional reform
1960s–1970s	Bretton Woods strain and collapse	fixed-exchange-rate regime risk gives way to market FX volatility
1973–1974	oil and equity shock	inflation, rates, commodities, and equities move under a new macro regime
1987	global equity crash	dynamic strategies, liquidity, jumps, and smile risk challenge diffusion assumptions
1992	ERM / sterling crisis	policy bands can create one-sided speculative payoff structures

Date	Episode	Quantitative lesson
1994	bond-market sell-off	duration, convexity, leverage, and rapid rate repricing matter across portfolios
1997	Asian financial crisis	currency mismatch, leverage, and cross-border funding amplify contagion
1998	Russia/LTCM	relative-value convergence can fail under funding and liquidity constraints
2000–2002	dot-com bust	valuation narratives, momentum, and factor crowding reverse sharply
2007–2008	global financial crisis	correlation, credit, liquidity, funding, model, and systemic risks interact
2010	Flash Crash	automated feedback and liquidity withdrawal can compress crisis dynamics into minutes
2011–2012	euro-area sovereign stress	sovereign, bank, collateral, and redenomination risks become coupled
2015	Swiss franc discontinuity	policy-regime change can produce extreme jumps and liquidity gaps
2020	pandemic market shock	cross-asset deleveraging and market-function stress challenge normal-times liquidity assumptions
2022	rate/inflation repricing	duration and leverage risks re-emerge after a long low-rate regime
2022	UK gilt / LDI stress	leveraged hedging can create procyclical collateral and liquidity feedback

D Bibliographical essay

Several literatures that are usually separated are necessary for a genuinely historical understanding of quantitative finance. For early money, debt, and exchange, Davies [98], Goetzmann [155], Graeber [159] provide contrasting broad narratives. For merchant banking and the transition toward organised capital markets, de Roover’s Medici studies, Neal’s account of early international capital markets, Gelderblom and Jonker’s work on Dutch finance, and Petram’s study of Amsterdam are particularly useful [103, 104, 146, 296, 315]. Braudel provides a wider institutional canvas [65].

For the intellectual history of modern finance, the primary papers remain unusually readable: Bachelier, Markowitz, Samuelson, Fama, Black–Scholes, Merton, and the martingale-pricing papers [30, 58, 133, 175, 178, 277, 284, 337]. Bernstein, Mehrling, Derman, MacKenzie, Weatherall, and Patterson provide complementary historical and sociological interpretations [47, 107, 269, 281, 311, 379].

Volatility is a literature large enough to constitute a separate history. The path from Mandelbrot through ARCH/GARCH, stochastic volatility, implied surfaces, realized volatility, and rough volatility can be followed through Andersen et al. [15], Bollerslev [61], Cont [87], Engle [126], Gatheral [143], Gatheral et al. [144], Heston [191], Mandelbrot [276], Taylor [362]. Microstructure and electronic trading are surveyed and systematised by O’Hara, Harris, Hasbrouck, and later algorithmic-trading texts [81, 174, 183, 250, 307]. The empirical HFT literature is represented here by work on latency, market making, price discovery, and the Flash Crash [69, 184, 234, 282].

For crises and regulation, no single chronology suffices. Kindleberger–Aliber and Reinhart–Rogoff

provide long-run perspectives; Brunnermeier, Gorton, and the FCIC address the global financial crisis; Basel and European Union primary documents establish the regulatory framework [40, 71, 129, 137, 158, 233, 328]. Machine learning and AI are represented by both general methodological texts and finance-specific applications [5, 156, 166, 185, 348, 387, 390].

References

- [1] Frédéric Abergel, Marouane Anane, Anirban Chakraborti, Aymen Jedidi, and Ioane Muni Toke. *Limit Order Books*. Cambridge University Press, Cambridge, 2016.
- [2] Carlo Acerbi. Spectral measures of risk: A coherent representation of subjective risk aversion. *Journal of Banking & Finance*, 26(7):1505–1518, 2002.
- [3] Claudio Albanese. Fourth level blas and high performance pricing / oplib cuda lab. Thalesian GPUs in Finance Workshop, 2009. URL <https://thalesians.com/education/schools/first-thalesian-full-day-workshop-gpus-in-finance/>.
- [4] Claudio Albanese and Hongyun Li. Monte carlo pricing using operator methods and measure changes. *SSRN Electronic Journal*, 2009.
- [5] Iñaki Aldasoro, Leonardo Gambacorta, Anton Korinek, Viral Shreeti, and Manuela Stein. Intelligent financial system: How ai is transforming finance. BIS Working Papers 1194, Bank for International Settlements, 2024.
- [6] Irene Aldridge. *High-Frequency Trading: A Practical Guide to Algorithmic Strategies and Trading Systems*. Wiley, 2 edition, 2013.
- [7] Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *The Journal of Risk*, 3(2):5–39, 2001.
- [8] Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *Journal of Risk*, 3(2):5–39, 2001.
- [9] Elisa Alòs. A generalization of the hull and white formula with applications to option pricing approximation. *Finance and Stochastics*, 10:353–365, 2006.
- [10] Elisa Alòs, Jorge A. León, and Josep Vives. On the short-time behavior of the implied volatility for jump-diffusion models with stochastic volatility. *Finance and Stochastics*, 2007. Representative Malliavin/asymptotic work.
- [11] S. M. Amadae. *Rationalizing Capitalist Democracy: The Cold War Origins of Rational Choice Liberalism*. University of Chicago Press, Chicago, 2003.
- [12] Saeed Amen. *Trading Thalesians: What the Ancient World Can Teach Us About Trading Today*. Palgrave Macmillan, 2014.
- [13] Saeed Amen. Takeaways from quantminds 2023, 2023. Practitioner account of the Global Derivatives/QuantMinds transition.
- [14] Leif B. G. Andersen and Vladimir V. Piterbarg. *Interest Rate Modeling*. Atlantic Financial Press, 2010.

- [15] Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. Modeling and forecasting realized volatility. *Econometrica*, 71(2):579–625, 2003.
- [16] Thomas Aquinas. *Political Writings*. Cambridge University Press, Cambridge, 2006. Edited and translated by R. W. Dyson.
- [17] Aristotle. *Politics*. Oxford University Press, Oxford, 1995. Edition associated with Ernest Barker.
- [18] Kenneth J. Arrow. *Social Choice and Individual Values*. Wiley, New York, 1951. Cowles Commission Monograph No. 12.
- [19] Kenneth J. Arrow. The role of securities in the optimal allocation of risk-bearing. *Review of Economic Studies*, 31(2):91–96, 1964.
- [20] Kenneth J. Arrow. The role of securities in the optimal allocation of risk-bearing. *Review of Economic Studies*, 31(2):91–96, 1964. doi: 10.2307/2296188. English publication of work first presented in French in 1953.
- [21] Kenneth J. Arrow and Gerard Debreu. Existence of an equilibrium for a competitive economy. *Econometrica*, 22(3):265–290, 1954. doi: 10.2307/1907353.
- [22] Kenneth J. Arrow and Leonid Hurwicz. Decentralization and computation in resource allocation. In Ralph W. Pfouts, editor, *Essays in Economics and Econometrics*, pages 34–104. University of North Carolina Press, Chapel Hill, 1960.
- [23] Kenneth J. Arrow and Leonid Hurwicz, editors. *Studies in Resource Allocation Processes*. Cambridge University Press, Cambridge, 1977.
- [24] Philippe Artzner, Freddy Delbaen, Jean-Marc Eber, and David Heath. Coherent measures of risk. *Mathematical Finance*, 9(3):203–228, 1999.
- [25] Amine Assouel, Antoine Jacquier, and Alexei Kondratyev. A quantum generative adversarial network for distributions. *arXiv preprint arXiv:2110.02742*, 2021.
- [26] Marco Avellaneda and Jeong-Hyun Lee. Statistical arbitrage in the us equities market. *Quantitative Finance*, 10(7):761–782, 2010.
- [27] Marco Avellaneda and Jeong-Hyun Lee. Statistical arbitrage in the us equities market. *Quantitative Finance*, 10(7):761–782, 2010.
- [28] Marco Avellaneda and Sasha Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008.
- [29] Marco Avellaneda and Sasha Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008.
- [30] Louis Bachelier. Théorie de la spéculation. *Annales Scientifiques de l’École Normale Supérieure*, 17:21–86, 1900.
- [31] Emmanuel Bacry, Iacopo Mastromatteo, and Jean-François Muzy. Hawkes processes in finance. *Market Microstructure and Liquidity*, 1(1), 2015.

- [32] Walter Bagehot. The only game in town. *Financial Analysts Journal*, 27(2):12–14, 22, 1971. Pseudonym used by Jack L. Treynor.
- [33] David H. Bailey, Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu. The probability of backtest overfitting. *Journal of Computational Finance*, 2014.
- [34] Gurdip Bakshi, Charles Cao, and Zhiwu Chen. Empirical performance of alternative option pricing models. *The Journal of Finance*, 52(5):2003–2049, 1997.
- [35] Clifford A. Ball and Antonio Roma. Stochastic volatility option pricing. *Journal of Financial and Quantitative Analysis*, 29(4):589–607, 1994.
- [36] Ole E. Barndorff-Nielsen and Neil Shephard. Econometric analysis of realized volatility and its use in estimating stochastic volatility models. *Journal of the Royal Statistical Society: Series B*, 64(2):253–280, 2002.
- [37] Basel Committee on Banking Supervision. International convergence of capital measurement and capital standards. Technical report, Bank for International Settlements, Basel, 1988.
- [38] Basel Committee on Banking Supervision. International convergence of capital measurement and capital standards: A revised framework. Technical report, Bank for International Settlements, Basel, 2004.
- [39] Basel Committee on Banking Supervision. Basel iii: A global regulatory framework for more resilient banks and banking systems. Technical report, Bank for International Settlements, Basel, 2011.
- [40] Basel Committee on Banking Supervision. Minimum capital requirements for market risk. Technical report, Bank for International Settlements, Basel, 2019.
- [41] Martin Baxter and Andrew Rennie. *Financial Calculus: An Introduction to Derivative Pricing*. Cambridge University Press, Cambridge, 1996.
- [42] Christian Bayer, Peter Friz, and Jim Gatheral. Pricing under rough volatility. *Quantitative Finance*, 16(6):887–904, 2016.
- [43] Tobias Becker, Oskar Mencer, Stephen Weston, and Georgi Gaydadjiev. Maxeler data-flow in computational finance. In *FPGA Based Accelerators for Financial Applications*, pages 243–266. Springer, 2015.
- [44] Richard Bellman. *Dynamic Programming*. Princeton University Press, Princeton, NJ, 1957.
- [45] Lorenzo Bergomi. Smile dynamics ii. *Risk*, 2005.
- [46] Lorenzo Bergomi. *Stochastic Volatility Modeling*. Chapman and Hall/CRC, 2016.
- [47] Peter L. Bernstein. *Capital Ideas: The Improbable Origins of Modern Wall Street*. Free Press, New York, 1992.
- [48] Peter L. Bernstein. *Capital Ideas Evolving*. Wiley, Hoboken, NJ, 2007.
- [49] Dimitri P. Bertsekas. *Dynamic Programming and Optimal Control*. Athena Scientific, 4 edition, 2017.

- [50] Dimitri P. Bertsekas. *Reinforcement Learning and Optimal Control*. Athena Scientific, 2019.
- [51] Dimitris Bertsimas and Andrew W. Lo. Optimal control of execution costs. *Journal of Financial Markets*, 1(1):1–50, 1998.
- [52] Fischer Black. Capital market equilibrium with restricted borrowing. *The Journal of Business*, 45(3):444–455, 1972.
- [53] Fischer Black. Fact and fantasy in the use of options. *Financial Analysts Journal*, 31(4):36–41, 61–72, 1975.
- [54] Fischer Black. Studies of stock price volatility changes. In *Proceedings of the 1976 Meetings of the American Statistical Association, Business and Economic Statistics Section*, pages 177–181, 1976.
- [55] Fischer Black. Noise. *The Journal of Finance*, 41(3):528–543, 1986.
- [56] Fischer Black and Piotr Karasinski. Bond and option pricing when short rates are lognormal. *Financial Analysts Journal*, 47(4):52–59, 1991.
- [57] Fischer Black and Myron Scholes. The valuation of option contracts and a test of market efficiency. *The Journal of Finance*, 27(2):399–417, 1972.
- [58] Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.
- [59] Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973. doi: 10.1086/260062.
- [60] Fischer Black, Emanuel Derman, and William Toy. A one-factor model of interest rates and its application to treasury bond options. *Financial Analysts Journal*, 46(1):33–39, 1990.
- [61] Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- [62] Michael D. Bordo and Barry Eichengreen, editors. *A Retrospective on the Bretton Woods System: Lessons for International Monetary Reform*. University of Chicago Press, Chicago, 1993.
- [63] Jean-Philippe Bouchaud, Yuval Gefen, Marc Potters, and Matthieu Wyart. Fluctuations and response in financial markets: The subtle nature of random price changes. *Quantitative Finance*, 4(2):176–190, 2004.
- [64] Jean-Philippe Bouchaud, J. Doyne Farmer, and Fabrizio Lillo. How markets slowly digest changes in supply and demand. In *Handbook of Financial Markets: Dynamics and Evolution*, pages 57–160. Elsevier, 2009.
- [65] Fernand Braudel. *Civilization and Capitalism, 15th–18th Century, Volume II: The Wheels of Commerce*. Harper & Row, New York, 1982.
- [66] Leo Breiman. Random forests. *Machine Learning*, 45:5–32, 2001.
- [67] Damiano Brigo and Fabio Mercurio. *Interest Rate Models: Theory and Practice*. Springer, 2 edition, 2006.

- [68] Damiano Brigo, Massimo Morini, and Andrea Pallavicini. *Counterparty Credit Risk, Collateral and Funding*. Wiley, 2013.
- [69] Jonathan Brogaard, Terrence Hendershott, and Ryan Riordan. High-frequency trading and price discovery. *The Review of Financial Studies*, 27(8):2267–2306, 2014.
- [70] Xavier Brokmann, David Itkin, Johannes Muhle-Karbe, and Peter Schmidt. Tackling nonlinear price impact with linear strategies. *Mathematical Finance*, 35(2):422–440, 2025.
- [71] Markus K. Brunnermeier. Deciphering the liquidity and credit crunch 2007–2008. *Journal of Economic Perspectives*, 23(1):77–100, 2009.
- [72] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.
- [73] Hans Buehler, Lukas Gonon, Josef Teichmann, and Ben Wood. Deep hedging. *Quantitative Finance*, 19(8):1271–1291, 2019.
- [74] Christoph Burgard and Mats Kjaer. Partial differential equation representations of derivatives with bilateral counterparty risk and funding costs. *The Journal of Credit Risk*, 7(3):1–19, 2011.
- [75] Sarah Butcher. Is a slang programming job at goldman sachs a technology career killer? eFinancialCareers, 2018.
- [76] John Y. Campbell and Ludger Hentschel. No news is good news: An asymmetric model of changing volatility in stock returns. *Journal of Financial Economics*, 31(3):281–318, 1992.
- [77] Bryan Cantrill. A conversation with arthur whitney. *ACM Queue*, 2009.
- [78] Luca Capriotti and Mike Giles. Fast correlation greeks by adjoint algorithmic differentiation. *Risk*, 2012.
- [79] Peter Carr and Roger Lee. Volatility derivatives. *Annual Review of Financial Economics*, 1:319–339, 2009.
- [80] Peter Carr and Dilip Madan. Option valuation using the fast fourier transform. *Journal of Computational Finance*, 2(4):61–73, 1999.
- [81] Álvaro Cartea, Sebastian Jaimungal, and José Penalva. *Algorithmic and High-Frequency Trading*. Cambridge University Press, Cambridge, 2015.
- [82] Charles Castelli. *The Theory of “Options” in Stocks and Shares*. F. C. Mathieson, London, 1877.
- [83] Ernest P. Chan. *Quantitative Trading: How to Build Your Own Algorithmic Trading Business*. Wiley, Hoboken, NJ, 2009.
- [84] Ernest P. Chan. *Algorithmic Trading: Winning Strategies and Their Rationale*. Wiley, Hoboken, NJ, 2013.
- [85] Edward Chancellor. *Devil Take the Hindmost: A History of Financial Speculation*. Farrar, Straus and Giroux, New York, 1999.

- [86] Peter K. Clark. A subordinated stochastic process model with finite variance for speculative prices. *Econometrica*, 41(1):135–155, 1973.
- [87] Rama Cont. Empirical properties of asset returns: Stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.
- [88] Rama Cont and David-Antoine Fournié. Functional itô calculus and stochastic integral representation of martingales. *Annals of Probability*, 41(1):109–133, 2013.
- [89] Rama Cont and Milena Vuletić. Simulation of arbitrage-free implied volatility surfaces. *Applied Mathematical Finance*, 30(2):94–121, 2023.
- [90] John C. Cox, Stephen A. Ross, and Mark Rubinstein. Option pricing: A simplified approach. *Journal of Financial Economics*, 7(3):229–263, 1979.
- [91] John C. Cox, Jonathan E. Ingersoll, and Stephen A. Ross. A theory of the term structure of interest rates. *Econometrica*, 53(2):385–407, 1985.
- [92] CQG and MathWorks. Cqg connects to mathworks matlab to advance quantitative trading analytics, 2010. Press release featuring Steve Wilcockson.
- [93] Stéphane Crépey, Tomasz R. Bielecki, and Damiano Brigo. *Counterparty Risk and Funding: A Tale of Two Puzzles*. Chapman & Hall/CRC, Boca Raton, 2014.
- [94] Jakša Cvitanović and Ioannis Karatzas. Convex duality in constrained portfolio optimization. *Annals of Applied Probability*, 2(4):767–818, 1992.
- [95] Jakša Cvitanović and Fernando Zapatero. *Introduction to the Economics and Mathematics of Financial Markets*. MIT Press, 2004.
- [96] George B. Dantzig. *Linear Programming and Extensions*. Princeton University Press, Princeton, NJ, 1963.
- [97] George B. Dantzig. Linear programming. *Operations Research*, 50(1):42–47, 2002. doi: 10.1287/opre.50.1.42.17798.
- [98] Glyn Davies. *A History of Money: From Ancient Times to the Present Day*. University of Wales Press, Cardiff, 3 edition, 2002.
- [99] Mark H. A. Davis. Piecewise-deterministic markov processes: A general class of non-diffusion stochastic models. *Journal of the Royal Statistical Society. Series B*, 46(3):353–388, 1984.
- [100] Mark H. A. Davis. *Markov Models and Optimization*. Chapman and Hall, 1993.
- [101] Mark H. A. Davis, Vassilios G. Panas, and Thaleia Zariphopoulou. European option pricing with transaction costs. *SIAM Journal on Control and Optimization*, 31(2):470–493, 1993.
- [102] Joseph de la Vega. *Confusion de Confusiones*. [s.n.], Amsterdam, 1688. Early treatise on the Amsterdam securities market.
- [103] Raymond de Roover. The medici bank: Financial and commercial operations. *The Journal of Economic History*, 6(2):153–172, 1946.
- [104] Raymond de Roover. *The Rise and Decline of the Medici Bank, 1397–1494*. Harvard University Press, Cambridge, MA, 1963.

- [105] Gerard Debreu. *Theory of Value: An Axiomatic Analysis of Economic Equilibrium*. Yale University Press, New Haven, 1959. Cowles Foundation Monograph No. 17.
- [106] Alexander Denev and Douglas Mortensen. *Probabilistic Graphical Models: A New Way of Thinking in Financial Modelling*. Risk Books, 2015.
- [107] Emanuel Derman. *My Life as a Quant: Reflections on Physics and Finance*. Wiley, Hoboken, NJ, 2004.
- [108] Emanuel Derman and Iraj Kani. Riding on a smile. *Risk*, 7(2):32–39, 1994.
- [109] Emanuel Derman and Iraj Kani. Riding on a smile. *Risk*, 1994.
- [110] Matthew F. Dixon, Igor Halperin, and Paul Bilokon. *Machine Learning in Finance: From Theory to Practice*. Springer, 2020.
- [111] Joseph L. Doob. *Stochastic Processes*. Wiley, New York, 1953.
- [112] Darrell Duffie. *Dynamic Asset Pricing Theory*. Princeton University Press, 3 edition, 2001.
- [113] Darrell Duffie and Kenneth J. Singleton. Modeling term structures of defaultable bonds. *Review of Financial Studies*, 12(4):687–720, 1999.
- [114] Daniel J. Duffy. *Finite Difference Methods in Financial Engineering*. Wiley, 2006.
- [115] Nicholas Dunbar. *Inventing Money: The Story of Long-Term Capital Management and the Legends Behind It*. Wiley, Chichester, 2000.
- [116] Bruno Dupire. Pricing with a smile. *Risk*, 7(1):18–20, 1994.
- [117] David Easley, Marcos López de Prado, and Maureen O’Hara. Flow toxicity and liquidity in a high-frequency world. *The Review of Financial Studies*, 25(5):1457–1493, 2012.
- [118] eFinancialCareers. Sarah butcher – global editor, 2026. URL <https://www.efinancialcareers.co.uk/news/author/sarahbutcher-en>. Accessed 2026-08-14.
- [119] Albert Einstein. über die von der molekularkinetischen theorie der wärme geforderte bewegung von in ruhenden flüssigkeiten suspendierten teilchen. *Annalen der Physik*, 322(8):549–560, 1905.
- [120] Omar El Euch and Mathieu Rosenbaum. The characteristic function of rough heston models. *Mathematical Finance*, 29(1):3–38, 2019.
- [121] Nicole El Karoui, Shige Peng, and Marie-Claire Quenez. Backward stochastic differential equations in finance. *Mathematical Finance*, 7(1):1–71, 1997.
- [122] Michael Ellman. Optimal planning. *Soviet Studies*, 20(2):226–248, 1968.
- [123] Michael Ellman. *Planning Problems in the USSR: The Contribution of Mathematical Methods to Their Solution, 1960–1971*. Cambridge University Press, Cambridge, 1973.
- [124] Youssef Elouerkhaoui. Pricing and hedging in a dynamic credit model. *International Journal of Theoretical and Applied Finance*, 2007.

- [125] Youssef Elouerkhaoui. *Credit Correlation: Theory and Practice*. Palgrave Macmillan, 2017. Representative credit-correlation work.
- [126] Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of united kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- [127] Richard E. Ericson. System for optimal functioning of the (socialist) economy, 2018. URL <https://www.aeaweb.org/conference/2018/preliminary/powerpoint/eGQKDKHd>. Conference presentation on the history and structure of SOFE.
- [128] European Parliament and Council. Directive 2004/39/ec on markets in financial instruments, 2004. Official Journal of the European Union.
- [129] European Parliament and Council. Directive 2014/65/eu on markets in financial instruments (mifid ii), 2014. Official Journal of the European Union.
- [130] European Parliament and Council. Regulation (eu) no 600/2014 on markets in financial instruments (mifir), 2014. Official Journal of the European Union.
- [131] Frank J. Fabozzi, Petter N. Kolm, Dessislava A. Pachamanova, and Sergio M. Focardi, editors. *Robust Portfolio Optimization and Management*. Wiley, 2007.
- [132] Eugene F. Fama. The behavior of stock-market prices. *The Journal of Business*, 38(1):34–105, 1965.
- [133] Eugene F. Fama. Efficient capital markets: A review of theory and empirical work. *The Journal of Finance*, 25(2):383–417, 1970.
- [134] J. Doyne Farmer, Paolo Patelli, and Ilija I. Zovko. The predictive power of zero intelligence in financial markets. *Proceedings of the National Academy of Sciences*, 102(6):2254–2259, 2005.
- [135] Nikolai P. Fedorenko. *Optimal Functioning System for a Socialist Economy*. Progress Publishers, Moscow, 1974. English-language work associated with the System for Optimal Functioning of the Socialist Economy (SOFE).
- [136] Niall Ferguson. *The Ascent of Money: A Financial History of the World*. Penguin Press, New York, 2008.
- [137] Financial Crisis Inquiry Commission. The financial crisis inquiry report. Technical report, U.S. Government Printing Office, Washington, DC, 2011.
- [138] Thomas Fischer and Christopher Krauss. Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research*, 270(2):654–669, 2018.
- [139] Filipe Fontanela, Antoine Jacquier, and Mugad Oumgari. A quantum algorithm for linear pdes arising in finance. *arXiv preprint arXiv:1912.02753*, 2019.
- [140] Martin Forde and Antoine Jacquier. Small-time asymptotics for implied volatility under the heston model. *International Journal of Theoretical and Applied Finance*, 12, 2009.
- [141] Saul I. Gass. The life and times of the father of linear programming. *OR/MS Today*, 32(2), 2005.

- [142] Evan Gatev, William N. Goetzmann, and K. Geert Rouwenhorst. Pairs trading: Performance of a relative-value arbitrage rule. *The Review of Financial Studies*, 19(3):797–827, 2006.
- [143] Jim Gatheral. *The Volatility Surface: A Practitioner’s Guide*. Wiley, Hoboken, NJ, 2006.
- [144] Jim Gatheral, Thibault Jaisson, and Mathieu Rosenbaum. Volatility is rough. *Quantitative Finance*, 18(6):933–949, 2018.
- [145] Weiqing Ge, Parsa Lalbakhsh, Lucas Isai, Artem Lenskiy, and Hanna Suominen. Neural network-based financial volatility forecasting: A systematic review. *ACM Computing Surveys*, 55(1):1–30, 2022.
- [146] Oscar Gelderblom and Joost Jonker. Completing a financial revolution: The finance of the dutch east india trade and the rise of the amsterdam capital market, 1595–1612. *The Journal of Economic History*, 64(3):641–672, 2004.
- [147] Hélyette Geman. *Commodities and Commodity Derivatives*. Wiley, 2005.
- [148] Hélyette Geman, Nicole El Karoui, and Jean-Charles Rochet. Changes of numeraire, changes of probability measure and option pricing. *Journal of Applied Probability*, 32:443–458, 1995.
- [149] Slava Gerovitch. Internyet: Why the soviet union did not build a nationwide computer network. *History and Technology*, 24(4):335–350, 2008. doi: 10.1080/07341510802044736.
- [150] Michael B. Giles. Multilevel monte carlo path simulation. *Operations Research*, 56(3):607–617, 2008.
- [151] Michael B. Giles and Łukasz Szpruch. Antithetic multilevel monte carlo estimation for multi-dimensional sdes without lévy area simulation. *Annals of Applied Probability*, 24(4):1585–1620, 2014.
- [152] I. V. Girsanov. On transforming a certain class of stochastic processes by absolutely continuous substitution of measures. *Theory of Probability and Its Applications*, 5(3):285–301, 1960. doi: 10.1137/1105027.
- [153] Paul Glasserman. *Monte Carlo Methods in Financial Engineering*, volume 53 of *Stochastic Modelling and Applied Probability*. Springer, New York, 2004. doi: 10.1007/978-0-387-21617-1.
- [154] Lawrence R. Glosten and Paul R. Milgrom. Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1):71–100, 1985.
- [155] William N. Goetzmann. *Money Changes Everything: How Finance Made Civilization Possible*. Princeton University Press, Princeton, 2016.
- [156] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, Cambridge, MA, 2016.
- [157] Neil J. Gordon, David J. Salmond, and Adrian F. M. Smith. Novel approach to nonlinear/non-gaussian bayesian state estimation. *IEE Proceedings F – Radar and Signal Processing*, 140(2):107–113, 1993.
- [158] Gary B. Gorton. *Slapped by the Invisible Hand: The Panic of 2007*. Oxford University Press, New York, 2010.

- [159] David Graeber. *Debt: The First 5,000 Years*. Melville House, Brooklyn, 2011.
- [160] Jon Gregory. *The xVA Challenge: Counterparty Credit Risk, Funding, Collateral and Capital*. Wiley, 3 edition, 2015.
- [161] Jon Gregory. *The xVA Challenge: Counterparty Risk, Funding, Collateral, Capital and Initial Margin*. Wiley, Chichester, 4 edition, 2020.
- [162] Richard C. Grinold and Ronald N. Kahn. *Active Portfolio Management*. McGraw-Hill, New York, 2 edition, 2000.
- [163] Sanford J. Grossman and Merton H. Miller. Liquidity and market structure. *The Journal of Finance*, 43(3):617–633, 1988.
- [164] Sanford J. Grossman and Joseph E. Stiglitz. On the impossibility of informationally efficient markets. *American Economic Review*, 70(3):393–408, 1980.
- [165] Sanford J. Grossman and Joseph E. Stiglitz. On the impossibility of informationally efficient markets. *American Economic Review*, 70(3):393–408, 1980.
- [166] Shihao Gu, Bryan Kelly, and Dacheng Xiu. Empirical asset pricing via machine learning. *The Review of Financial Studies*, 33(5):2223–2273, 2020.
- [167] Lewis Gudgeon, Sam M. Werner, Daniel Perez, and William J. Knottenbelt. DeFi protocols for loanable funds: Interest rates, liquidity and market efficiency. In *Proceedings of the 2nd ACM Conference on Advances in Financial Technologies*, pages 92–112. ACM, 2020. doi: 10.1145/3419614.3423254.
- [168] Olivier Guéant. *The Financial Mathematics of Market Liquidity: From Optimal Execution to Market Making*. Chapman and Hall/CRC, 2017.
- [169] Patrick S. Hagan, Deep Kumar, Andrew S. Lesniewski, and Diana E. Woodward. Managing smile risk. *Wilmott Magazine*, pages 84–108, 2002.
- [170] Igor Halperin. Qlbs: Q-learner in the black-scholes(-merton) worlds. *arXiv preprint arXiv:1712.04609*, 2017.
- [171] Igor Halperin. The qlbs q-learner goes nuclear: Fitted q iteration, inverse rl, and option portfolios. *Quantitative Finance*, 19(9):1543–1553, 2019.
- [172] Ben Hambly, Renyuan Xu, and Huining Yang. Recent advances in reinforcement learning in finance. *Mathematical Finance*, 33(3):437–503, 2023.
- [173] Peter R. Hansen and Asger Lunde. A forecast comparison of volatility models: Does anything beat a garch(1,1)? *Journal of Applied Econometrics*, 20(7):873–889, 2005.
- [174] Larry Harris. *Trading and Exchanges: Market Microstructure for Practitioners*. Oxford University Press, New York, 2003.
- [175] J. Michael Harrison and David M. Kreps. Martingales and arbitrage in multiperiod securities markets. *Journal of Economic Theory*, 20(3):381–408, 1979.
- [176] J. Michael Harrison and David M. Kreps. Martingales and arbitrage in multiperiod securities markets. *Journal of Economic Theory*, 20(3):381–408, 1979.

- [177] J. Michael Harrison and David M. Kreps. Martingales and arbitrage in multiperiod securities markets. *Journal of Economic Theory*, 20(3):381–408, 1979. doi: 10.1016/0022-0531(79)90043-7.
- [178] J. Michael Harrison and Stanley R. Pliska. Martingales and stochastic integrals in the theory of continuous trading. *Stochastic Processes and their Applications*, 11(3):215–260, 1981.
- [179] J. Michael Harrison and Stanley R. Pliska. Martingales and stochastic integrals in the theory of continuous trading. *Stochastic Processes and their Applications*, 11(3):215–260, 1981. doi: 10.1016/0304-4149(81)90026-0.
- [180] J. Michael Harrison and Stanley R. Pliska. A stochastic calculus model of continuous trading: Complete markets. *Stochastic Processes and their Applications*, 15(3):313–316, 1983. doi: 10.1016/0304-4149(83)90038-8.
- [181] Campbell R. Harvey, Ashwin Ramachandran, and Joey Santoro. *DeFi and the Future of Finance*. Wiley, Hoboken, NJ, 2021.
- [182] Dominik Harz and William J. Knottenbelt. Towards safer smart contracts: A survey of languages and verification methods, 2018.
- [183] Joel Hasbrouck. *Empirical Market Microstructure: The Institutions, Economics, and Econometrics of Securities Trading*. Oxford University Press, New York, 2007.
- [184] Joel Hasbrouck and Gideon Saar. Low-latency trading. *Journal of Financial Markets*, 16(4): 646–679, 2013.
- [185] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning*. Springer, New York, 2 edition, 2009.
- [186] Espen Gaarder Haug. *The Complete Guide to Option Pricing Formulas*. McGraw-Hill, 2 edition, 2007.
- [187] Espen Gaarder Haug and Nassim Nicholas Taleb. Why we have never used the black-scholes-merton option pricing formula. *Wilmott Magazine*, 2008.
- [188] Friedrich A. Hayek. The use of knowledge in society. *American Economic Review*, 35(4): 519–530, 1945.
- [189] David Heath, Robert Jarrow, and Andrew Morton. Bond pricing and the term structure of interest rates: A new methodology for contingent claims valuation. *Econometrica*, 60(1): 77–105, 1992.
- [190] David Heath, Robert Jarrow, and Andrew Morton. Bond pricing and the term structure of interest rates: A new methodology for contingent claims valuation. *Econometrica*, 60(1): 77–105, 1992.
- [191] Steven L. Heston. A closed-form solution for options with stochastic volatility with applications to bond and currency options. *The Review of Financial Studies*, 6(2):327–343, 1993.
- [192] Yves Hilpisch. *Python for Finance: Mastering Data-Driven Finance*. O’Reilly, 2 edition, 2018.
- [193] Yves Hilpisch. *Artificial Intelligence in Finance*. O’Reilly, 2020.

- [194] Thomas S. Y. Ho and Sang-Bin Lee. Term structure movements and pricing interest rate contingent claims. *The Journal of Finance*, 41(5):1011–1029, 1986.
- [195] David Hochfelder. *The Telegraph in America, 1832–1920*. Johns Hopkins University Press, Baltimore, 2012.
- [196] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997.
- [197] Blanka Horvath, Aitor Muguruza, and Mehdi Tomas. Deep learning in rough volatility. *Applied Mathematical Finance*, 2021. Representative work.
- [198] Blanka Horvath, Aitor Muguruza, and Mehdi Tomas. Deep learning volatility: A deep neural network perspective on pricing and calibration in rough volatility models. *Quantitative Finance*, 21(1):11–27, 2021.
- [199] Sanja Hukovic. Model risk, quantitative risk standards, and ai governance, 2026. Career and conference record including LSEG model-risk leadership.
- [200] John Hull and Alan White. The pricing of options on assets with stochastic volatilities. *The Journal of Finance*, 42(2):281–300, 1987.
- [201] John Hull and Alan White. Pricing interest-rate-derivative securities. *The Review of Financial Studies*, 3(4):573–592, 1990.
- [202] John Hull and Alan White. Cva and wrong-way risk. *Financial Analysts Journal*, 2012. Representative practitioner research.
- [203] John C. Hull. *Options, Futures, and Other Derivatives*. Pearson, 10 edition, 2018.
- [204] John C. Hull. *Options, Futures, and Other Derivatives*. Pearson, 11 edition, 2022.
- [205] Leonid Hurwicz. Optimality and informational efficiency in resource allocation processes. In Kenneth J. Arrow, Samuel Karlin, and Patrick Suppes, editors, *Mathematical Methods in the Social Sciences, 1959*, pages 27–46. Stanford University Press, Stanford, CA, 1960.
- [206] Leonid Hurwicz. On informationally decentralized systems. In C. B. McGuire and Roy Radner, editors, *Decision and Organization*, pages 297–336. North-Holland, Amsterdam, 1972.
- [207] Dragos I. Ilie, Sam M. Werner, Iain D. Stewart, and William J. Knottenbelt. Unstable throughput: When the difficulty algorithm breaks. In *2021 IEEE International Conference on Blockchain and Cryptocurrency (ICBC)*, pages 1–5. IEEE, 2021. doi: 10.1109/ICBC51069.2021.9461086.
- [208] Imperial College London. Tom hadfield: Senior research fellow in quantitative finance, 2026. URL <https://profiles.imperial.ac.uk/t.hadfield>. Accessed 2026-08-14.
- [209] Imperial College London. Cristopher salvi – research: Stochastic analysis, deep learning, mathematical ai, formal verification, lean 4, 2026. URL <https://profiles.imperial.ac.uk/c.salvi/grants>. Accessed 2026-08-14.
- [210] Informa Connect. Quantminds international, 2026. URL <https://informaconnect.com/quantminds-international/>. Conference history and programme.

- [211] Jonathan E. Ingersoll. *Theory of Financial Decision Making*. Rowman and Littlefield, 1987.
- [212] David Itkin. Stochastic portfolio theory with price impact. *arXiv preprint arXiv:2506.07993*, 2025.
- [213] Kiyosi Itô. Stochastic integral. *Proceedings of the Imperial Academy*, 20(8):519–524, 1944.
- [214] Antoine Jacquier and Oleksiy Kondratyev. *Quantum Machine Learning and Optimisation in Finance*. Packt, 2024.
- [215] Antoine Jacquier, Oleksiy Kondratyev, Gordon Lee, and Mugad Oumgari. Quantum computing for financial mathematics. *arXiv preprint arXiv:2311.06621*, 2023.
- [216] Eric Jacquier, Nicholas G. Polson, and Peter E. Rossi. Bayesian analysis of stochastic volatility models. *Journal of Business & Economic Statistics*, 12(4):371–389, 1994.
- [217] Peter Jaeckel. *Monte Carlo Methods in Finance*. Wiley, 2002.
- [218] Peter Jaeckel. Let’s be rational. *Wilmott*, 2015. Implied-volatility inversion algorithm.
- [219] Robert A. Jarrow and Stuart M. Turnbull. Pricing derivatives on financial securities subject to credit risk. *Journal of Finance*, 50(1):53–85, 1995.
- [220] Elliot Jones and William Knottenbelt. IsabeLLM: Automated theorem proving applied to formally verifying consensus. *Distributed Ledger Technologies: Research and Practice*, 2026. doi: 10.1145/3822513.
- [221] Elliot Jones, Toshiko Matsui, and William Knottenbelt. Stablecoins as dry powder: A copula-based risk analysis of cryptocurrency markets. In *2026 IEEE International Conference on Blockchain and Cryptocurrency (ICBC)*. IEEE, 2026. doi: 10.1109/ICBC67748.2026.11575537.
- [222] Philippe Jorion. *Value at Risk: The New Benchmark for Managing Financial Risk*. McGraw-Hill, New York, 3 edition, 2007.
- [223] Mark S. Joshi. *C++ Design Patterns and Derivatives Pricing*. Cambridge University Press, 2 edition, 2008.
- [224] J.P. Morgan/Reuters. Riskmetrics technical document. Technical report, J.P. Morgan/Reuters, New York, 1996.
- [225] JPMorgan Chase. Report of the review committee on the chief investment office, 2013. Internal review published after the 2012 CIO losses.
- [226] Leonid V. Kantorovich. Mathematical methods of organizing and planning production. *Management Science*, 6(4):366–422, 1960. English translation of work originating in 1939.
- [227] Leonid V. Kantorovich. *The Best Use of Economic Resources*. Harvard University Press, Cambridge, MA, 1965. English translation of *Ekonomicheskii raschet nailuchshego ispolzovaniya resursov*, first published in Russian in 1959.
- [228] Leonid V. Kantorovich. Mathematics in economics: Achievements, difficulties, perspectives, 1975. URL <https://www.nobelprize.org/prizes/economic-sciences/1975/kantorovich/lecture/>. Prize lecture, Sveriges Riksbank Prize in Economic Sciences in Memory of Alfred Nobel.

- [229] J. L. Kelly. A new interpretation of information rate. *Bell System Technical Journal*, 35(4): 917–926, 1956.
- [230] Leonid G. Khachiyan. A polynomial algorithm in linear programming. *Soviet Mathematics Doklady*, 20:191–194, 1979. Translation of Doklady Akademii Nauk SSSR 244, 1093–1096.
- [231] Ha Young Kim and Chang Hyun Won. Forecasting the volatility of stock price index: A hybrid model integrating lstm with multiple garch-type models. *Expert Systems with Applications*, 103:25–37, 2018.
- [232] Sangjoon Kim, Neil Shephard, and Siddhartha Chib. Stochastic volatility: Likelihood inference and comparison with arch models. *The Review of Economic Studies*, 65(3):361–393, 1998.
- [233] Charles P. Kindleberger and Robert Z. Aliber. *Manias, Panics, and Crashes: A History of Financial Crises*. Palgrave Macmillan, Basingstoke, 6 edition, 2011.
- [234] Andrei Kirilenko, Albert S. Kyle, Mehrdad Samadi, and Tugkan Tuzun. The flash crash: High-frequency trading in an electronic market. *The Journal of Finance*, 72(3):967–998, 2017.
- [235] Olessia Kirtchik. The soviet scientific programme on ai: If a machine cannot think, can it control? *BJHS Themes*, 8:111–125, 2023. doi: 10.1017/bjt.2023.4. Includes discussion of Soviet cybernetics and automated economic management.
- [236] Olessia Kirtchik. Soviet mathematical economics. *Journal of Economic Perspectives*, 2024. URL <https://www.aeaweb.org/articles/materials/22058>. Materials and historical survey published by the American Economic Association.
- [237] Petter N. Kolm and Gordon Ritter. Dynamic replication and hedging: A reinforcement learning approach. *Journal of Financial Data Science*, 1(1):159–171, 2019.
- [238] Petter N. Kolm and Gordon Ritter. Modern perspectives on reinforcement learning in finance. *Journal of Machine Learning in Finance*, 1(1), 2020.
- [239] A. N. Kolmogorov. *Foundations of the Theory of Probability*. Chelsea Publishing Company, New York, 1950. English translation of the 1933 *Grundbegriffe der Wahrscheinlichkeitsrechnung*.
- [240] Andrey N. Kolmogorov. *Grundbegriffe der Wahrscheinlichkeitsrechnung*. Springer, Berlin, 1933.
- [241] Tjalling C. Koopmans, editor. *Activity Analysis of Production and Allocation*. Wiley, New York, 1951. Cowles Commission Monograph No. 13.
- [242] Tjalling C. Koopmans. *Three Essays on the State of Economic Science*. McGraw-Hill, New York, 1957.
- [243] KX Systems. kdb+ and q documentation, 2026. URL <https://code.kx.com/q/>. Accessed 2026-08-14.
- [244] Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.
- [245] Oskar Lange. On the economic theory of socialism: Part one. *Review of Economic Studies*, 4(1):53–71, 1936. doi: 10.2307/2967660.

- [246] Oskar Lange. On the economic theory of socialism: Part two. *Review of Economic Studies*, 4 (2):123–142, 1937.
- [247] Anthony Lee, Christopher Yau, Michael B. Giles, Arnaud Doucet, and Christopher C. Holmes. On the utility of graphics cards to perform massively parallel simulation of advanced monte carlo methods. *Journal of Computational and Graphical Statistics*, 2010. Preprint arXiv:0905.2441.
- [248] Adam Leeds. Dreams in cybernetic fugue: Cold war technoscience, the intelligentsia, and the birth of soviet mathematical economics. *Historical Studies in the Natural Sciences*, 46(5): 633–668, 2016. doi: 10.1525/hsns.2016.46.5.633.
- [249] Charles-Albert Lehalle and Sophie Laruelle. *Market Microstructure in Practice*. World Scientific, 2013.
- [250] Charles-Albert Lehalle and Sophie Laruelle. *Market Microstructure in Practice*. World Scientific, 2 edition, 2018.
- [251] Charles-Albert Lehalle and Eyal Neuman. Incorporating signals into optimal trading. *Finance and Stochastics*, 23:275–311, 2019.
- [252] Wassily W. Leontief. *The Structure of American Economy, 1919–1929: An Empirical Application of Equilibrium Analysis*. Harvard University Press, Cambridge, MA, 1941.
- [253] Wassily W. Leontief. *Input-Output Economics*. Oxford University Press, New York, 1966.
- [254] Abba P. Lerner. *The Economics of Control: Principles of Welfare Economics*. Macmillan, New York, 1944.
- [255] Jonathan Levy. *Freaks of Fortune: The Emerging World of Capitalism and Risk in America*. Harvard University Press, Cambridge, MA, 2012.
- [256] John Lintner. The valuation of risk assets and the selection of risky investments in stock portfolios and capital budgets. *The Review of Economics and Statistics*, 47(1):13–37, 1965.
- [257] Alexander Lipton. *Mathematical Methods for Foreign Exchange: A Financial Engineer’s Approach*. World Scientific, 2001.
- [258] Alexander Lipton and Andrew Rennie, editors. *The Oxford Handbook of Credit Derivatives*. Oxford University Press, 2013.
- [259] Chun Liu and John M. Maheu. Are there structural breaks in realized volatility? *Journal of Financial Econometrics*, 6(3):326–360, 2008.
- [260] Yang Liu. Novel volatility forecasting using deep learning–long short term memory recurrent neural networks. *Expert Systems with Applications*, 132:99–109, 2019.
- [261] Andrew W. Lo. The adaptive markets hypothesis. *The Journal of Portfolio Management*, 30 (5):15–29, 2004.
- [262] Andrew W. Lo. *Adaptive Markets: Financial Evolution at the Speed of Thought*. Princeton University Press, Princeton, 2017.
- [263] Logos Research and Imperial College London. Logos research launches to make ai reasoning trustworthy in high-stakes environments, 2026. URL <https://www.logosresearch.ai/news/imperial-ai-reasoning-trustworthy/>.

- [264] Marcos López de Prado. *Advances in Financial Machine Learning*. Wiley, Hoboken, NJ, 2018.
- [265] Marcos López de Prado. *Advances in Financial Machine Learning*. Wiley, 2018.
- [266] Alejandro López-Lira and Yuehua Tang. Can chatgpt forecast stock price movements? return predictability and large language models. Technical Report 4412788, SSRN, 2023.
- [267] Roger Lowenstein. *When Genius Failed: The Rise and Fall of Long-Term Capital Management*. Random House, New York, 2000.
- [268] James MacDonald. *A Free Nation Deep in Debt: The Financial Roots of Democracy*. Farrar, Straus and Giroux, New York, 2006.
- [269] Donald MacKenzie. *An Engine, Not a Camera: How Financial Models Shape Markets*. MIT Press, Cambridge, MA, 2006.
- [270] Donald MacKenzie. *Trading at the Speed of Light: How Ultrafast Algorithms Are Transforming Financial Markets*. Princeton University Press, Princeton, 2021.
- [271] Donald MacKenzie, Fabian Muniesa, and Lucia Siu, editors. *Do Economists Make Markets? On the Performativity of Economics*. Princeton University Press, Princeton, 2007.
- [272] Dilip B. Madan and Eugene Seneta. The variance gamma model for share market returns. *Journal of Business*, 63(4):511–524, 1990.
- [273] Dilip B. Madan, Peter P. Carr, and Eric C. Chang. The variance gamma process and option pricing. *European Finance Review*, 2:79–105, 1998.
- [274] Ananth Madhavan. Market microstructure: A survey. *Journal of Financial Markets*, 3(3): 205–258, 2000.
- [275] Burton G. Malkiel. The efficient market hypothesis and its critics. *Journal of Economic Perspectives*, 17(1):59–82, 2003.
- [276] Benoit B. Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4):394–419, 1963.
- [277] Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952.
- [278] Harry M. Markowitz. *Portfolio Selection: Efficient Diversification of Investments*. Wiley, New York, 1959. Cowles Foundation Monograph No. 16.
- [279] Dan McGinn, David Birch, David Akroyd, Miguel Molina-Solana, Yike Guo, and William J. Knottenbelt. Visualizing dynamic bitcoin transaction patterns. *Big Data*, 4(2):109–119, 2016. doi: 10.1089/big.2015.0056.
- [280] Richard A. Meese and Kenneth Rogoff. Empirical exchange rate models of the seventies: Do they fit out of sample? *Journal of International Economics*, 14(1–2):3–24, 1983.
- [281] Perry Mehrling. *Fischer Black and the Revolutionary Idea of Finance*. Wiley, Hoboken, NJ, 2005.
- [282] Albert J. Menkveld. High frequency trading and the new-market makers. *Journal of Financial Markets*, 16(4):712–740, 2013.

- [283] Robert C. Merton. Theory of rational option pricing. *Bell Journal of Economics and Management Science*, 4(1):141–183, 1973. doi: 10.2307/3003143.
- [284] Robert C. Merton. Theory of rational option pricing. *The Bell Journal of Economics and Management Science*, 4(1):141–183, 1973.
- [285] Robert C. Merton. Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1–2):125–144, 1976.
- [286] Robert C. Merton. A functional perspective of financial intermediation. *Financial Management*, 24(2):23–41, 1995.
- [287] Attilio Meucci. *Risk and Asset Allocation*. Springer, 2005.
- [288] Attilio Meucci. Fully flexible views: Theory and practice. *Risk*, 2008.
- [289] Philip Mirowski. *Machine Dreams: Economics Becomes a Cyborg Science*. Cambridge University Press, Cambridge, 2002.
- [290] Ludwig von Mises. Die wirtschaftrrechnung im sozialistischen gemeinwesen. *Archiv für Sozialwissenschaft und Sozialpolitik*, 47:86–121, 1920. English translation commonly published as *Economic Calculation in the Socialist Commonwealth*.
- [291] Franco Modigliani and Merton H. Miller. The cost of capital, corporation finance and the theory of investment. *American Economic Review*, 48(3):261–297, 1958.
- [292] Jan Mossin. Equilibrium in a capital asset market. *Econometrica*, 34(4):768–783, 1966.
- [293] Marek Musiela. Stochastic pdes and term structure models. *Journées Internationales de Finance*, 1993. Working paper/lecture notes.
- [294] Marek Musiela and Marek Rutkowski. *Martingale Methods in Financial Modelling*. Springer, 2 edition, 2005.
- [295] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system, 2008. White paper.
- [296] Larry Neal. *The Rise of Financial Capitalism: International Capital Markets in the Age of Reason*. Cambridge University Press, Cambridge, 1990.
- [297] Salih N. Neftci. *An Introduction to the Mathematics of Financial Derivatives*. Academic Press, 2 edition, 2000.
- [298] Salih N. Neftci. *Principles of Financial Engineering*. Academic Press, 2 edition, 2008.
- [299] Daniel B. Nelson. Conditional heteroskedasticity in asset returns: A new approach. *Econometrica*, 59(2):347–370, 1991.
- [300] Vasily S. Nemchinov. *The Use of Mathematics in Economics*. MIT Press, Cambridge, MA, 1964. English edition edited with an introduction by Alec Nove.
- [301] Eyal Neuman and Moritz Voß. Trading with the crowd. *Mathematical Finance*, 33(3):548–617, 2023.
- [302] Victor Niederhoffer. *The Education of a Speculator*. Wiley, 1997.

- [303] Nobel Prize Outreach. The sveriges riksbank prize in economic sciences in memory of alfred nobel 1975: Leonid vitaliyevich kantorovich and tjalling c. koopmans, 1975. URL <https://www.nobelprize.org/prizes/economic-sciences/1975/summary/>.
- [304] Jan Novotny, Paul A. Bilokon, Aris Galiotos, and Frédéric Déléze. *Machine Learning and Big Data with kdb+/q*. Wiley, 2019. doi: 10.1002/9781119404729.
- [305] Viktor V. Novozhilov. *Problems of Cost-Benefit Analysis in Optimal Planning*. International Arts and Sciences Press, White Plains, NY, 1970. English translation of Novozhilov’s work on scarcity, opportunity costs, and optimal planning.
- [306] Anna A. Obizhaeva and Jiang Wang. Optimal trading strategy and supply/demand dynamics. *Journal of Financial Markets*, 16(1):1–32, 2013.
- [307] Maureen O’Hara. *Market Microstructure Theory*. Blackwell, Cambridge, MA, 1995.
- [308] Bernt Øksendal. *Stochastic Differential Equations: An Introduction with Applications*. Springer, 6 edition, 2003.
- [309] Emiliano S. Pagnotta and Thomas Philippon. Competing on speed. *Econometrica*, 90(3): 1067–1115, 2022.
- [310] Panos Parpas and Berç Rustem. Stochastic and robust optimisation methods in finance. *Optimization and Engineering*, 2015. Representative strand; see author’s collected work.
- [311] Scott Patterson. *The Quants: How a New Breed of Math Whizzes Conquered Wall Street and Nearly Destroyed It*. Crown Business, New York, 2010.
- [312] Imanol Pérez Arribas, Cristopher Salvi, and Łukasz Szpruch. Sig-sdes model for quantitative finance. *arXiv preprint arXiv:2006.00218*, 2020.
- [313] Joseph Persky. Retrospectives: Lange and von mises, large-scale enterprises, and the economic case for socialism. *Journal of Economic Perspectives*, 5(4):229–236, 1991. doi: 10.1257/jep.5.4.229.
- [314] Benjamin Peters. *How Not to Network a Nation: The Uneasy History of the Soviet Internet*. MIT Press, Cambridge, MA, 2016.
- [315] Lodewijk Petram. *The World’s First Stock Exchange*. Columbia University Press, New York, 2014.
- [316] Huyên Pham. *Continuous-Time Stochastic Control and Optimization with Financial Applications*. Springer, 2009.
- [317] Vladimir Piterbarg. Funding beyond discounting: Collateral agreements and derivatives pricing. *Risk*, 2010.
- [318] Michael K. Pitt and Neil Shephard. Filtering via simulation: Auxiliary particle filters. *Journal of the American Statistical Association*, 94(446):590–599, 1999.
- [319] Karl Polanyi. *The Great Transformation*. Farrar & Rinehart, New York, 1944.
- [320] L. S. Pontryagin, V. G. Boltyanskii, R. V. Gamkrelidze, and E. F. Mishchenko. *The Mathematical Theory of Optimal Processes*. Interscience, New York, 1962.

- [321] Warren B. Powell. *Approximate Dynamic Programming: Solving the Curses of Dimensionality*. Wiley, 2 edition, 2011.
- [322] Warren B. Powell. *Reinforcement Learning and Stochastic Optimization: A Unified Framework for Sequential Decisions*. Wiley, 2022.
- [323] Alex Preda. *Framing Finance: The Boundaries of Markets and Modern Capitalism*. University of Chicago Press, Chicago, 2009.
- [324] Quantitative Developers Club. Neil fowler: Founder of wbs training, 2026. URL <https://qdc.org.uk/team/>. Public biography; WBS founded in 2000.
- [325] Fintan Quill. Time series analytics for streaming big fast data. Carnegie Mellon Database Seminar, 2017. URL <https://db.cs.cmu.edu/events/time-series-database-lectures-fintan-quill-kx-systems/>.
- [326] William Quinn and John D. Turner. *Boom and Bust: A Global History of Financial Bubbles*. Cambridge University Press, Cambridge, 2020.
- [327] Roy Radner. Existence of equilibrium of plans, prices, and price expectations in a sequence of markets. *Econometrica*, 40(2):289–303, 1972. doi: 10.2307/1909407.
- [328] Carmen M. Reinhart and Kenneth S. Rogoff. *This Time Is Different: Eight Centuries of Financial Folly*. Princeton University Press, Princeton, 2009.
- [329] Eglė Rindzevičiūtė. Soviet policy sciences and earth system governmentality. *Modern Intellectual History*, 17(1):179–208, 2020. doi: 10.1017/S1479244318000161. Discusses cybernetic control and the intellectual context of Soviet policy sciences.
- [330] Risk.net. Rising star in quant finance: Milena vuletić, 2024. URL <https://www.risk.net/awards/7960374/rising-star-in-quant-finance-milena-vuletic>. Risk Awards 2025; published 2024-11-26.
- [331] Risk.net. Rising star in quant finance: David itkin, 2025. URL <https://www.risk.net/awards/7962609/rising-star-in-quant-finance-david-itkin>. Risk Awards 2026; published 2025-11-25.
- [332] Richard Roll. A simple implicit measure of the effective bid-ask spread in an efficient market. *The Journal of Finance*, 39(4):1127–1139, 1984.
- [333] Stephen A. Ross. The arbitrage theory of capital asset pricing. *Journal of Economic Theory*, 13(3):341–360, 1976.
- [334] Stephen A. Ross. The arbitrage theory of capital asset pricing. *Journal of Economic Theory*, 13(3):341–360, 1976.
- [335] Mark Rubinstein. Implied binomial trees. *The Journal of Finance*, 49(3):771–818, 1994.
- [336] Berç Rustem. *Algorithms for Nonlinear Programming and Multiple-Objective Decisions*. Wiley, 1998.
- [337] Paul A. Samuelson. Proof that properly anticipated prices fluctuate randomly. *Industrial Management Review*, 6(2):41–49, 1965.

- [338] Gleb Sandmann and Siem Jan Koopman. Estimation of stochastic volatility models via monte carlo maximum likelihood. *Journal of Econometrics*, 87(2):271–301, 1998.
- [339] Antonio Serra. *A Short Treatise on the Wealth and Poverty of Nations*. [s.n.], 1613. Originally published in Italian as Breve trattato delle cause che possono far abbondare li regni d’oro e argento.
- [340] William F. Sharpe. Capital asset prices: A theory of market equilibrium under conditions of risk. *The Journal of Finance*, 19(3):425–442, 1964.
- [341] Neil Shephard and Torben G. Andersen. Stochastic volatility: Origins and overview. In Torben G. Andersen, Richard A. Davis, Jens-Peter Kreiß, and Thomas Mikosch, editors, *Handbook of Financial Time Series*, pages 233–254. Springer, Berlin, 2009.
- [342] Robert J. Shiller. Do stock prices move too much to be justified by subsequent changes in dividends? *American Economic Review*, 71(3):421–436, 1981.
- [343] Albert N. Shiryaev. *Essentials of Stochastic Finance: Facts, Models, Theory*. World Scientific, 1999.
- [344] Albert N. Shiryaev. *Optimal Stopping Rules*. Springer, 2008.
- [345] Andrei Shleifer and Robert W. Vishny. The limits of arbitrage. *The Journal of Finance*, 52(1):35–55, 1997.
- [346] Steven E. Shreve. *Stochastic Calculus for Finance I: The Binomial Asset Pricing Model*. Springer, 2004.
- [347] Steven E. Shreve. *Stochastic Calculus for Finance II: Continuous-Time Models*. Springer, 2004.
- [348] Justin Sirignano and Rama Cont. Universal features of price formation in financial markets: Perspectives from deep learning. *Quantitative Finance*, 19(9):1449–1459, 2019.
- [349] Adam Smith. *An Inquiry into the Nature and Causes of the Wealth of Nations*. W. Strahan and T. Cadell, London, 1776.
- [350] Société Générale. Report on the events relating to the unauthorized trading activities of Jérôme Kerviel. Technical report, Société Générale, 2008.
- [351] Alexander Sokol. *Long-Term Portfolio Simulation: For XVA, Limits, Liquidity and Regulatory Capital*. Risk Books, 2014.
- [352] Didier Sornette. *Why Stock Markets Crash: Critical Events in Complex Financial Systems*. Princeton University Press, 2003.
- [353] Tom Standage. *The Victorian Internet: The Remarkable Story of the Telegraph and the Nineteenth Century’s On-Line Pioneers*. Walker & Company, New York, 2 edition, 2007.
- [354] Benn Steil. *The Battle of Bretton Woods: John Maynard Keynes, Harry Dexter White, and the Making of a New World Order*. Princeton University Press, Princeton, 2013.
- [355] Harvey Stein. Selected work on model risk and quantitative finance, 2012. Practitioner research and conference literature.

- [356] James Steuart. *An Inquiry into the Principles of Political Oeconomy*. A. Millar and T. Cadell, London, 1767.
- [357] Iain Stewart, Dragos Ilie, Alexei Zamyatin, Sam Werner, M. F. Torshizi, and William J. Knottenbelt. Committing to quantum resistance: A slow defence for bitcoin against a fast quantum computing attack. *Royal Society Open Science*, 5(6):180410, 2018. doi: 10.1098/rsos.180410.
- [358] Nassim Nicholas Taleb. *Dynamic Hedging: Managing Vanilla and Exotic Options*. Wiley, 1997.
- [359] Nassim Nicholas Taleb. *The Black Swan: The Impact of the Highly Improbable*. Random House, New York, 2007.
- [360] Stephen J. Taylor. *Modelling Financial Time Series*. Wiley, Chichester, 1986.
- [361] Stephen J. Taylor. Forecasting the volatility of currency exchange rates. *International Journal of Forecasting*, 3(1):159–170, 1987.
- [362] Stephen J. Taylor. Modeling stochastic volatility: A review and comparative study. *Mathematical Finance*, 4(2):183–204, 1994.
- [363] The Thalesians. The first thalesian full-day workshop: Gpus in finance, 2009. URL <https://thalesians.com/education/schools/first-thalesian-full-day-workshop-gpus-in-finance/>.
- [364] The Thalesians. The thalesians’ history, 2026. URL <https://thalesians.com/about/organization/>. Accessed 2026-08-14.
- [365] The Thalesians. The thalesian seminar series, 2026. URL <https://thalesians.com/seminars/>. Accessed 2026-08-14.
- [366] Edward O. Thorp. *Beat the Dealer: A Winning Strategy for the Game of Twenty-One*. Random House, New York, 1966.
- [367] Edward O. Thorp. The kelly criterion in blackjack, sports betting, and the stock market. *Handbook of Asset and Liability Management*, 2007. Revised exposition of earlier Kelly-criterion applications.
- [368] Edward O. Thorp and Sheen T. Kassouf. *Beat the Market: A Scientific Stock Market System*. Random House, New York, 1967.
- [369] James Tobin. Liquidity preference as behavior towards risk. *Review of Economic Studies*, 25(2):65–86, 1958. doi: 10.2307/2296205.
- [370] Jack L. Treynor. How to rate management of investment funds. *Harvard Business Review*, 43(1):63–75, 1965.
- [371] Ruey S. Tsay. *Analysis of Financial Time Series*. Wiley, Hoboken, NJ, 3 edition, 2010.
- [372] Igor Tulchinsky, editor. *Finding Alphas: A Quantitative Approach to Building Trading Strategies*. Wiley, 2015.

- [373] United States Congress. Banking act of 1933, 1933. Public Law 73–66; provisions commonly known as the Glass–Steagall Act.
- [374] United States Congress. Dodd–frank wall street reform and consumer protection act, 2010. Public Law 111–203.
- [375] U.S. Securities and Exchange Commission. Concept release on equity market structure. Technical report, U.S. Securities and Exchange Commission, 2010.
- [376] U.S. Securities and Exchange Commission and Commodity Futures Trading Commission. Findings regarding the market events of may 6, 2010. Technical report, SEC and CFTC, 2010.
- [377] U.S. Senate Permanent Subcommittee on Investigations. Jpmorgan chase whale trades: A case history of derivatives risks and abuses. Technical report, U.S. Senate, 2013.
- [378] Oldrich Vasicek. An equilibrium characterization of the term structure. *Journal of Financial Economics*, 5(2):177–188, 1977.
- [379] James Owen Weatherall. *The Physics of Wall Street: A Brief History of Predicting the Unpredictable*. Houghton Mifflin Harcourt, Boston, 2013.
- [380] Sam M. Werner, Daniel Perez, Lewis Gudgeon, Aariah Klages-Mundt, Dominik Harz, and William J. Knottenbelt. SoK: Decentralized finance (DeFi). In *Proceedings of the 4th ACM Conference on Advances in Financial Technologies*, pages 30–46. ACM, 2022. doi: 10.1145/3558535.3559780.
- [381] Norbert Wiener. Differential-space. *Journal of Mathematics and Physics*, 2:131–174, 1923.
- [382] Steve Wilcockson. Financial services and quantitative technology: Practitioner work around matlab and research-to-production workflows, 2024. Practitioner and industry material.
- [383] Paul Wilmott. *Derivatives: The Theory and Practice of Financial Engineering*. Wiley, Chichester, 1998.
- [384] Paul Wilmott. *Paul Wilmott on Quantitative Finance*. Wiley, 2 edition, 2006.
- [385] Paul Wilmott, Sam Howison, and Jeff Dewynne. *The Mathematics of Financial Derivatives*. Cambridge University Press, 1995.
- [386] Gavin Wood. Ethereum: A secure decentralised generalised transaction ledger, 2014. Ethereum Yellow Paper.
- [387] Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David Rosenberg, and Gideon Mann. Bloomberggpt: A large language model for finance. *arXiv preprint arXiv:2303.17564*, 2023.
- [388] Xenophon. *Oeconomicus: A Social and Historical Commentary*. Clarendon Press, Oxford, 1994. Edition and commentary by Sarah B. Pomeroy.
- [389] Xihan Xiong, Zhipeng Wang, Tianxiang Cui, William Knottenbelt, and Michael Huth. Market misconduct in decentralized finance (DeFi): Analysis, regulatory challenges and policy implications, 2023.

- [390] Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. Fingpt: Open-source financial large language models. *arXiv preprint arXiv:2306.06031*, 2023.
- [391] Marc Yor. On some exponential functionals of brownian motion. *Advances in Applied Probability*, 24:509–531, 1992.
- [392] Marc Yor, editor. *Exponential Functionals of Brownian Motion and Related Processes*. Springer, 2001.
- [393] Alexei Zamyatin, Dominik Harz, Joshua Lind, Panayiotis Panayiotou, Arthur Gervais, and William J. Knottenbelt. XCLAIM: Trustless, interoperable, cryptocurrency-backed assets. In *2019 IEEE Symposium on Security and Privacy (SP)*, pages 193–210. IEEE, 2019. doi: 10.1109/SP.2019.00085.
- [394] Alexei Zamyatin, Mustafa Al-Bassam, Dionysis Zindros, Eleftherios Kokoris-Kogias, Pedro Moreno-Sanchez, Aggelos Kiayias, and William J. Knottenbelt. SoK: Communication across distributed ledgers. In *Financial Cryptography and Data Security*, volume 12675 of *Lecture Notes in Computer Science*, pages 3–36. Springer, 2021. doi: 10.1007/978-3-662-64331-0_1.
- [395] Philip Ziegler. *The Sixth Great Power: A History of One of the Greatest of All Banking Families, the House of Barings, 1762–1929*. Alfred A. Knopf, New York, 1988.
- [396] Gregory Zuckerman. *The Man Who Solved the Market: How Jim Simons Launched the Quant Revolution*. Portfolio, New York, 2019.